Parallel Processing in Python

DS 5110/CS 5501: Big Data Systems

Spring 2024

Lecture 3

Yue Cheng

Some material taken/derived from:

- Wisconsin CS 320 by Tyler Caraza-Harter.

Learning objectives

- Describe the execution model of
  - process-level parallelism
  - thread-level parallelism
  - task-level parallelism
- Know how to measure the speedup metric
- Understand the difference of strong scaling vs. weak scaling

Outline

- Motivation
- Three parallel execution models
- Measuring speedup metric
- Task parallelism in Dask
- Demo

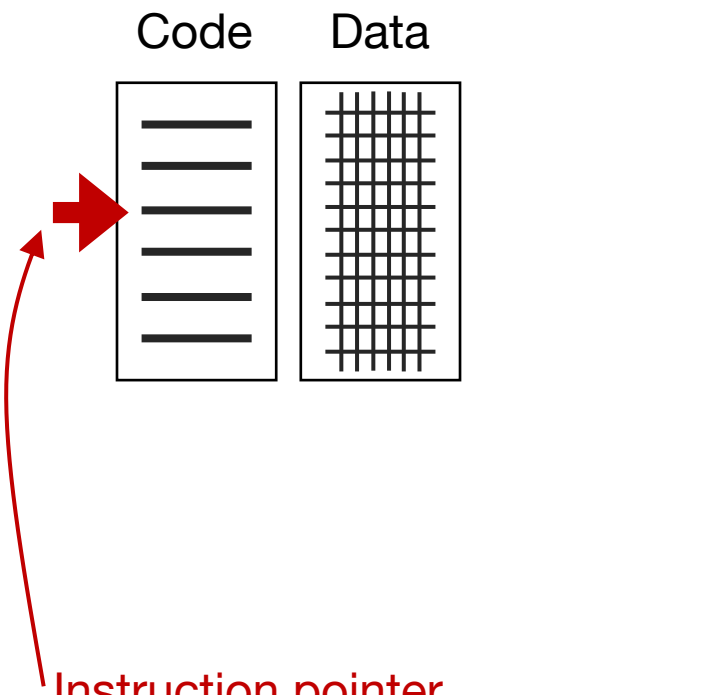
Code
Data
Instruction pointer
(also called “program counter”)

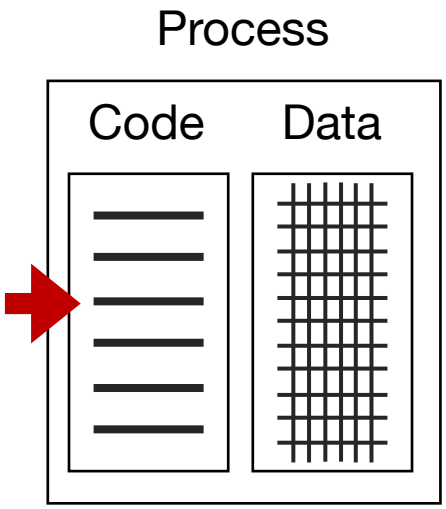

Instruction pointer belongs to a thread within the process

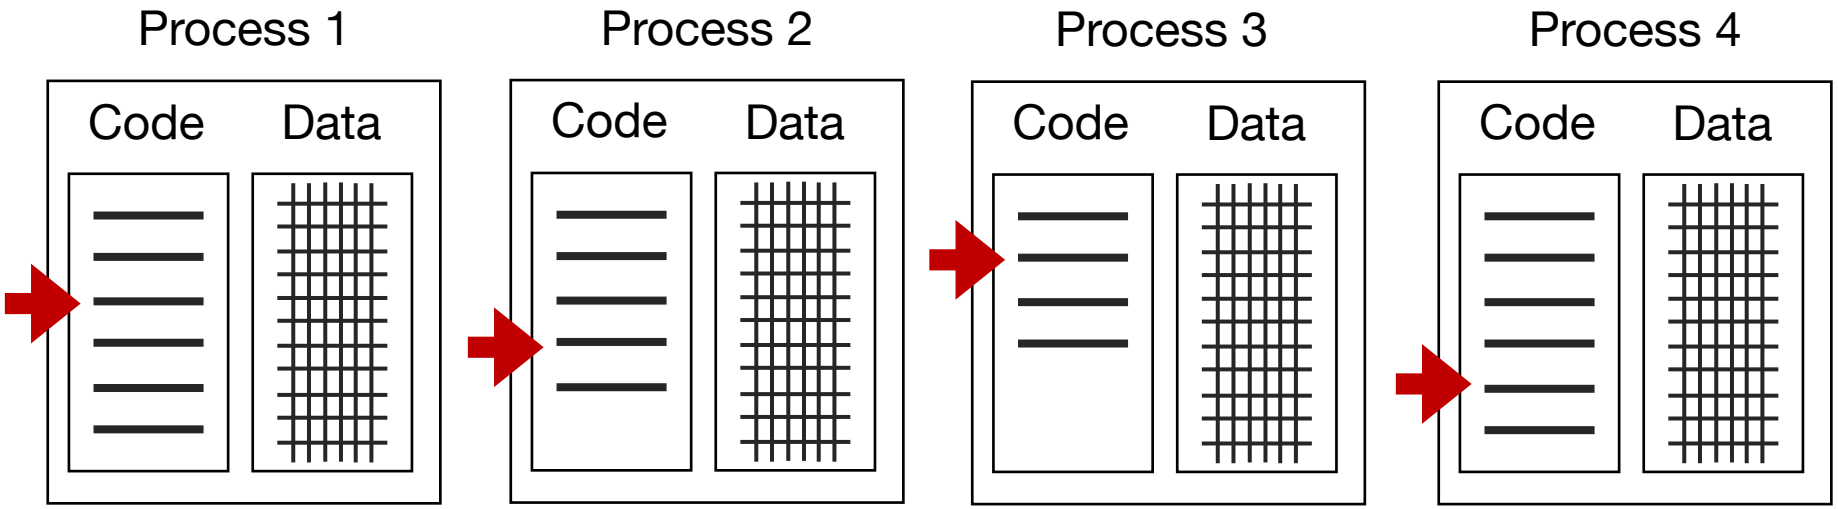

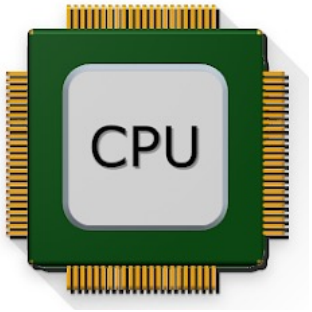

Multi-core processor (CPU)

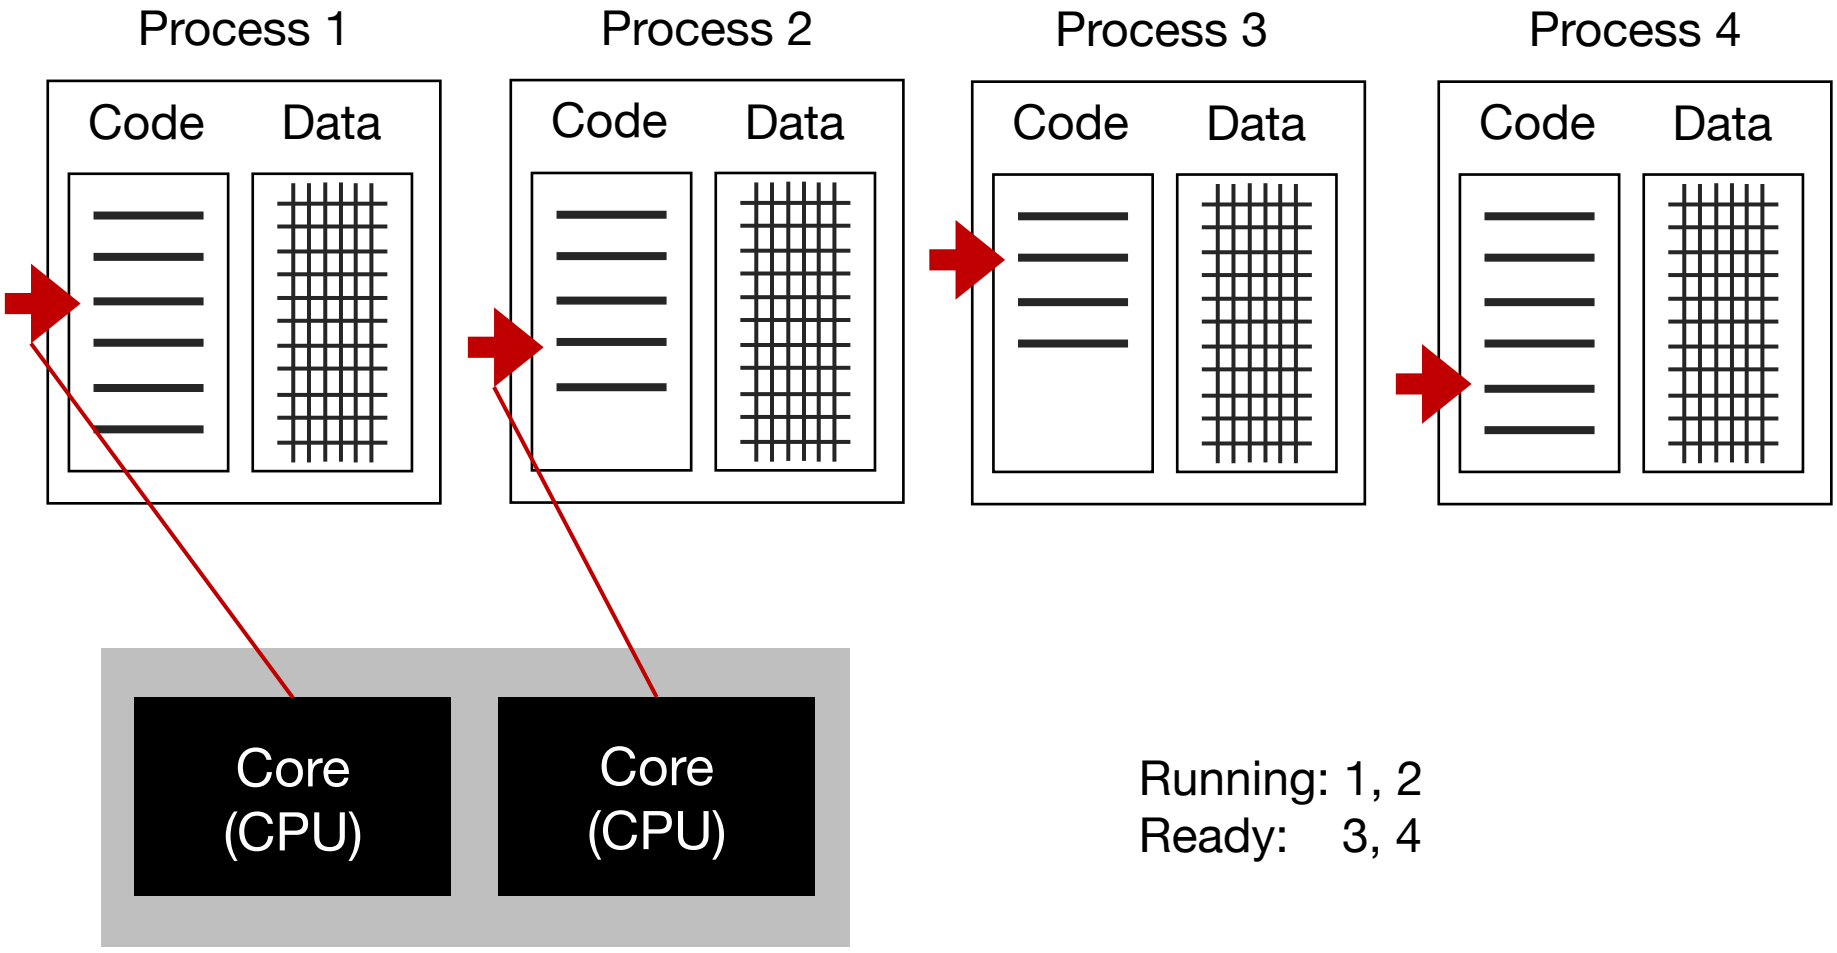
Process 1
Process 2
Process 3
Process 4
Code
Data
Code
Data
Code
Data
Code
Data
Core
(CPU)
Core
(CPU)
Running: 1, 2
Ready: 3, 4
Multi-core processor (CPU)

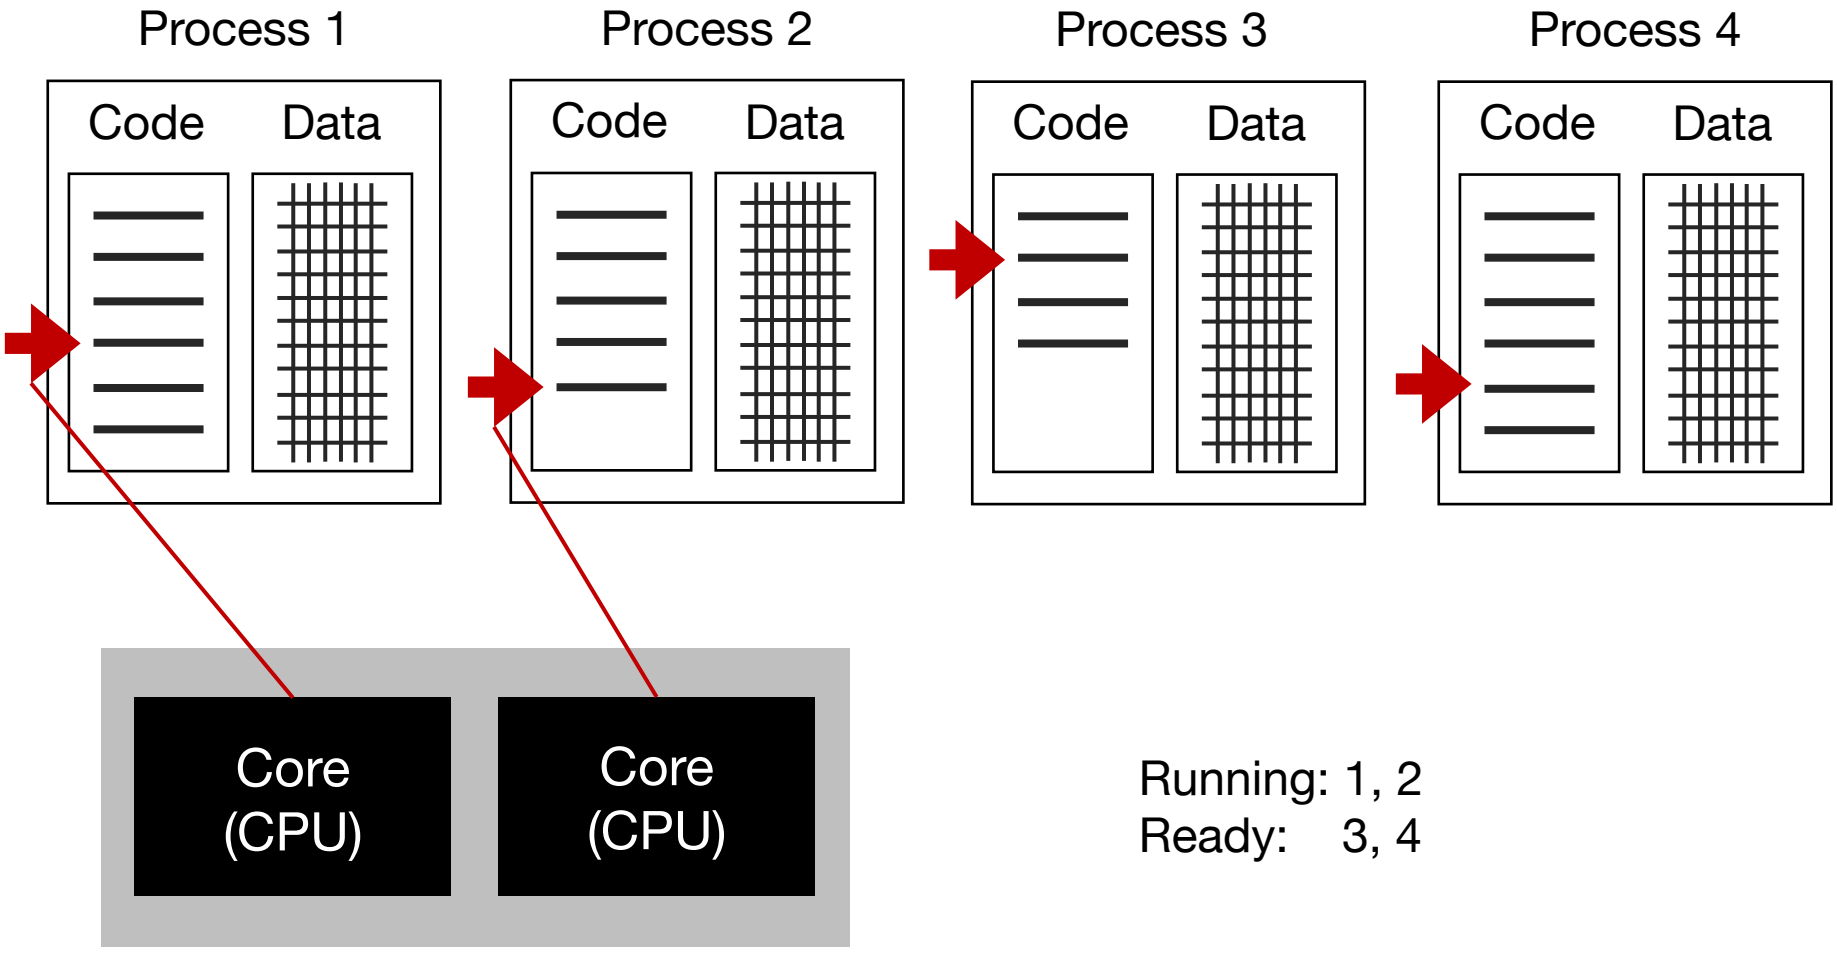
Process 1
Process 2
Process 3
Process 4
Code
Data
Code
Data
Code
Data
Code
Data
Core
(CPU)
Core
(CPU)
Running: 1, 2
Ready: 3, 4
Multi-core processor (CPU)

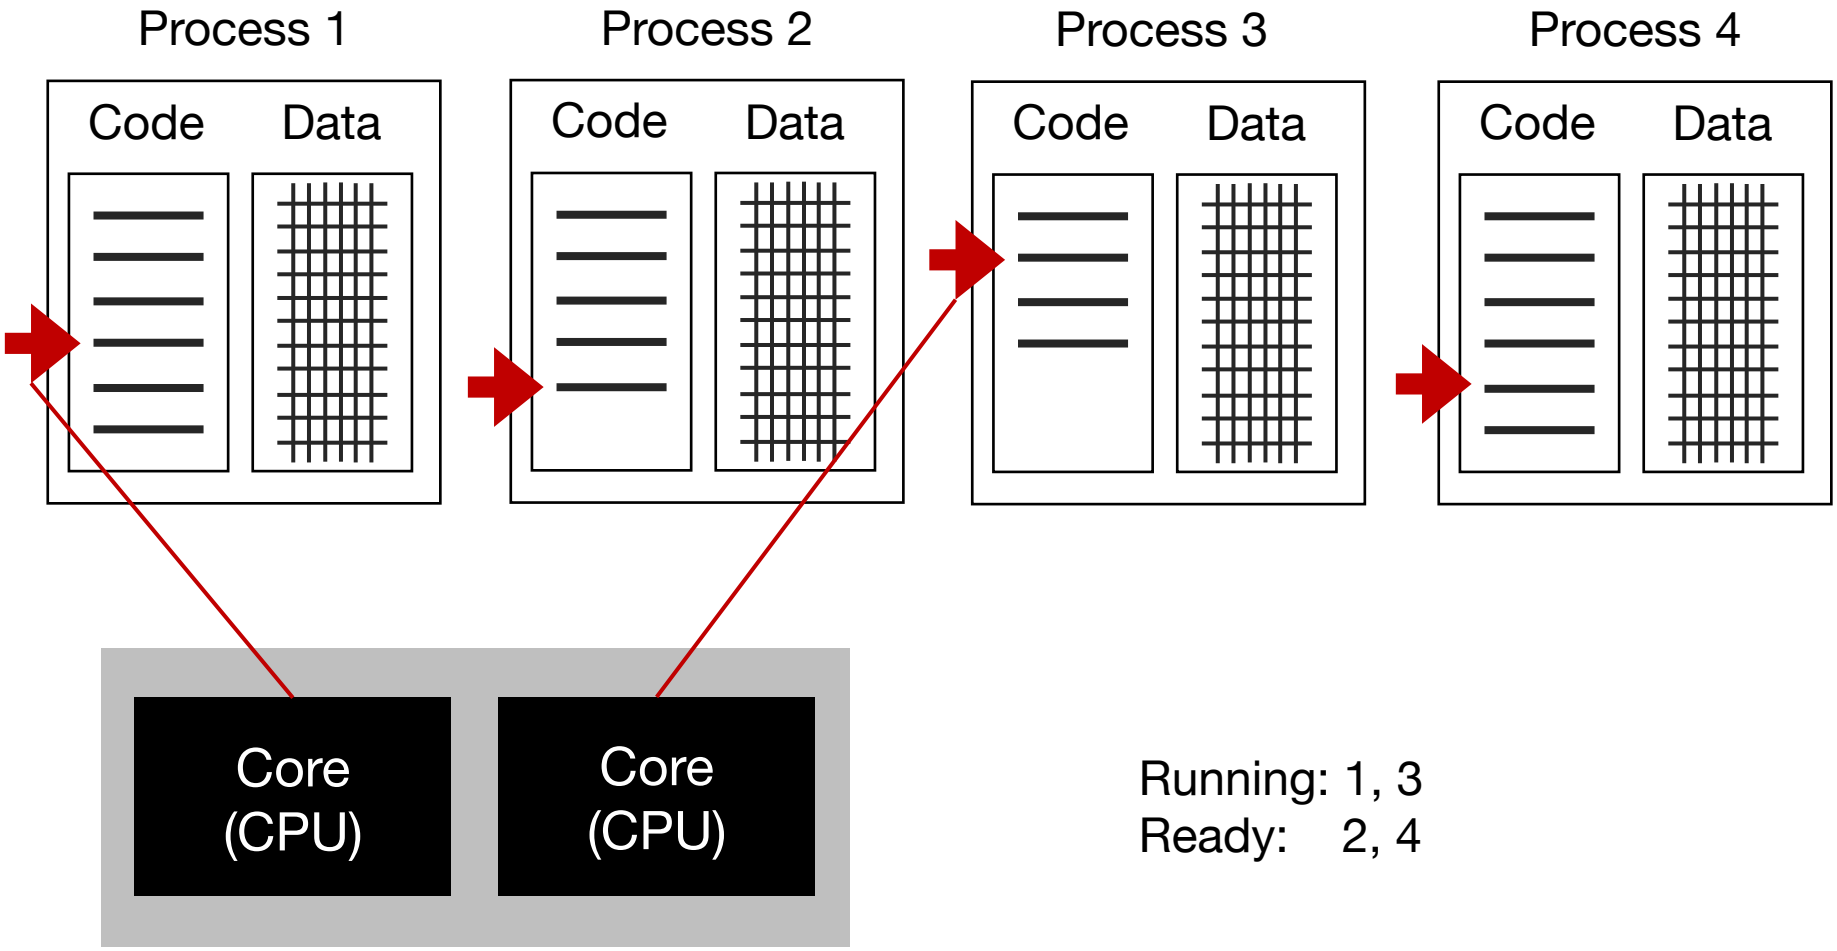
Process 1
Process 2
Process 3
Process 4
Code
Data
Code
Data
Code
Data
Code
Data
Core
(CPU)
Core
(CPU)
Running: 1, 3
Ready: 2, 4
Multi-core processor (CPU)

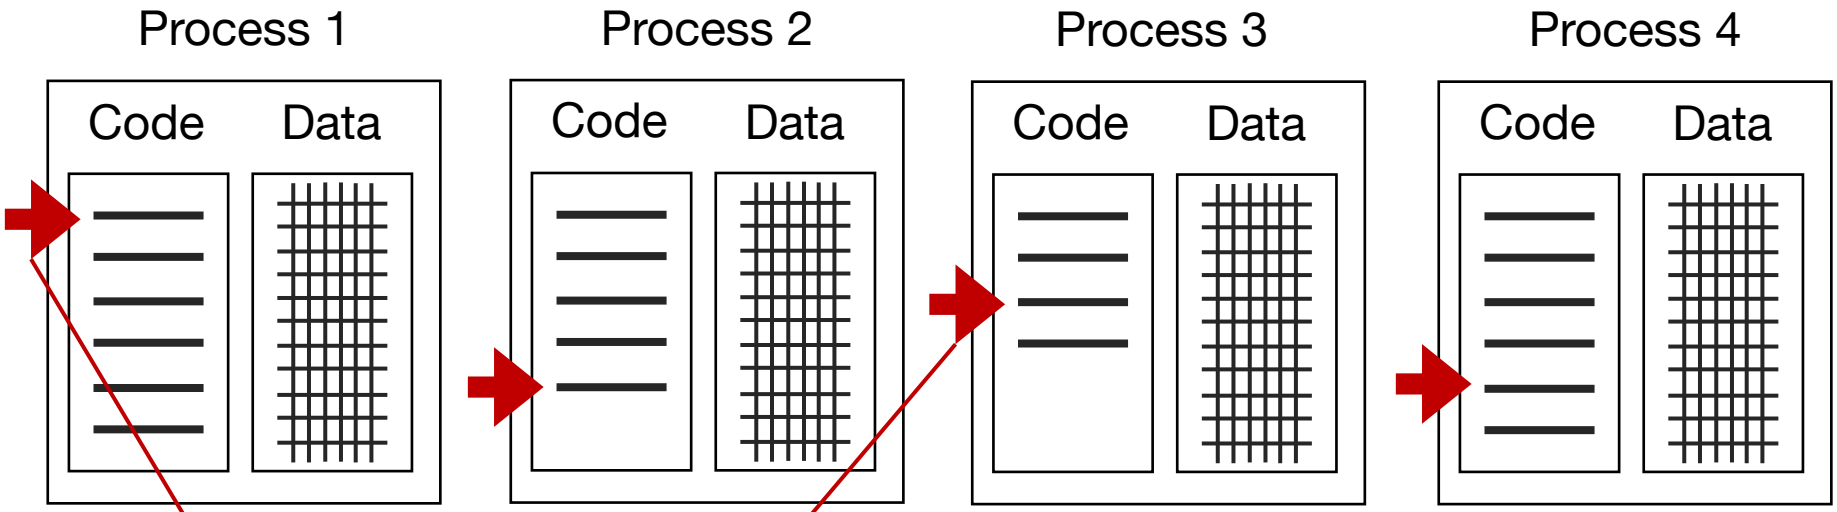

Multi-core processor (CPU)

Running: 1, 3
Ready: 2, 4

The more cores we have, the more tasks we can run simultaneously

Parallel execution models

- Process-level parallelism
- Thread-level parallelism
- Task-level parallelism

Parallel execution models

- **Process-level parallelism**
- Thread-level parallelism
- Task-level parallelism

Process-level parallelism

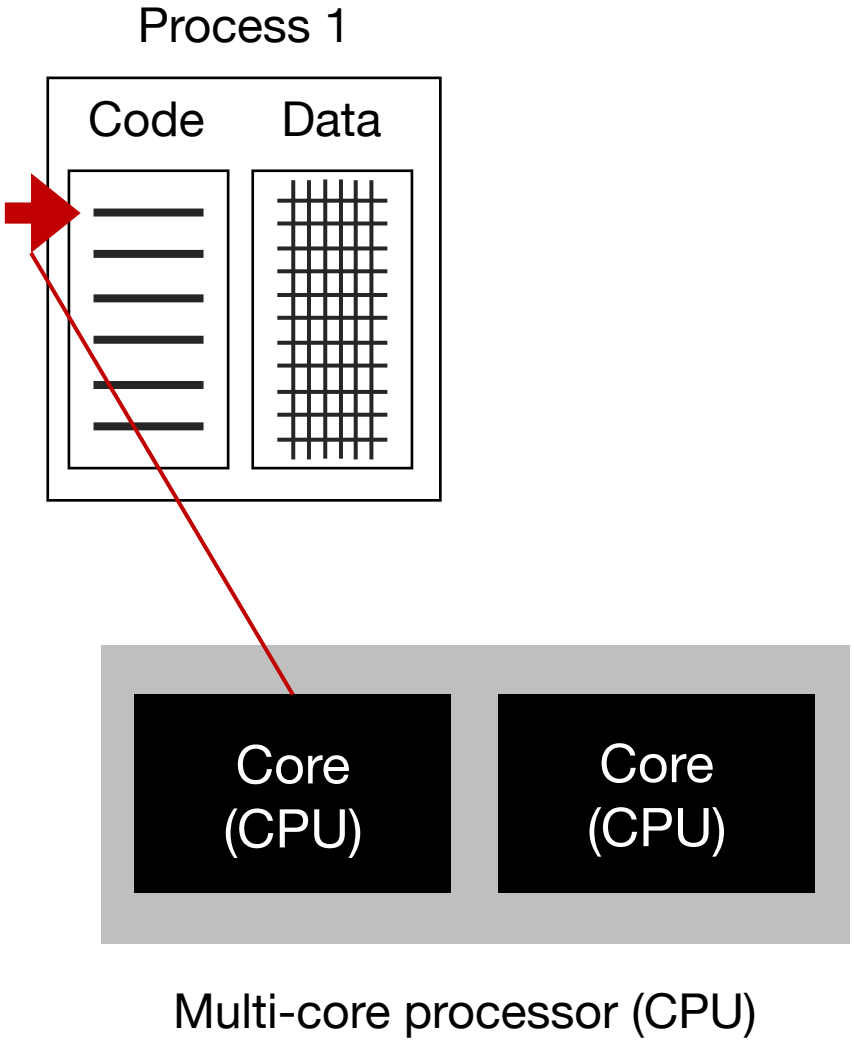

Process-level parallelism

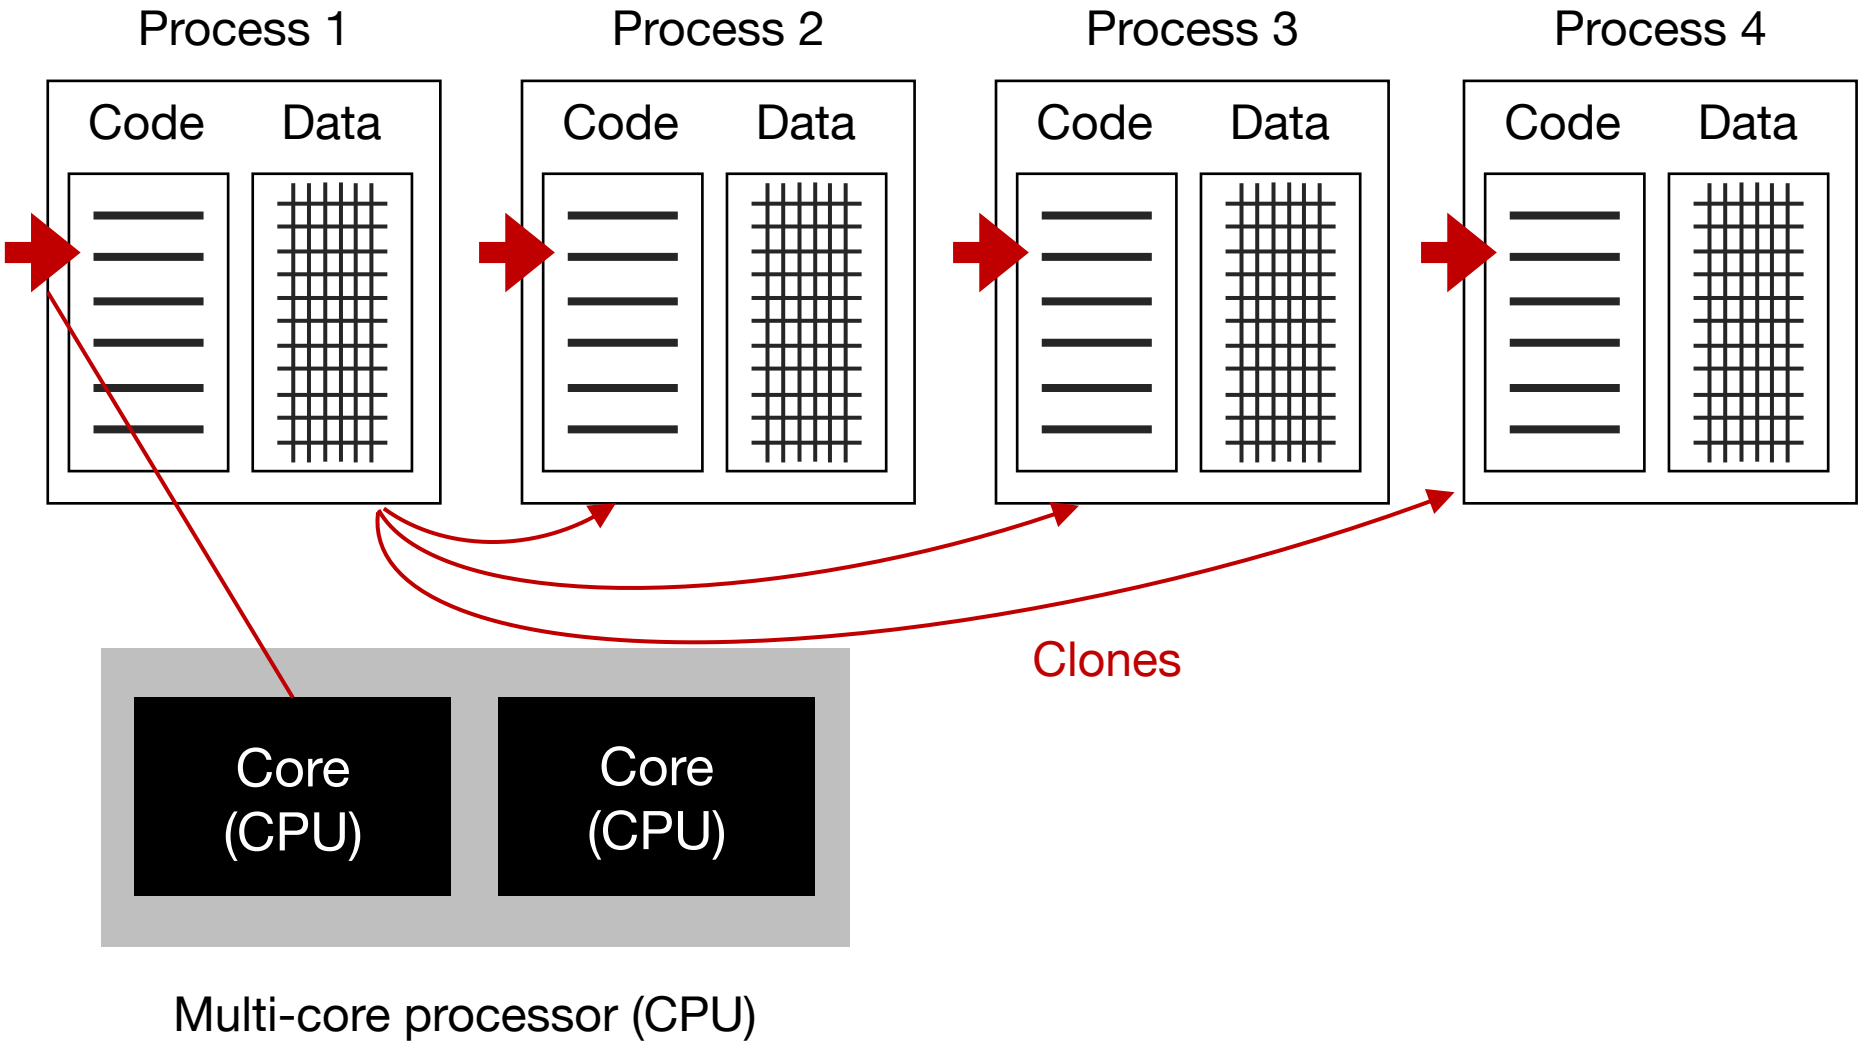

Process-level parallelism

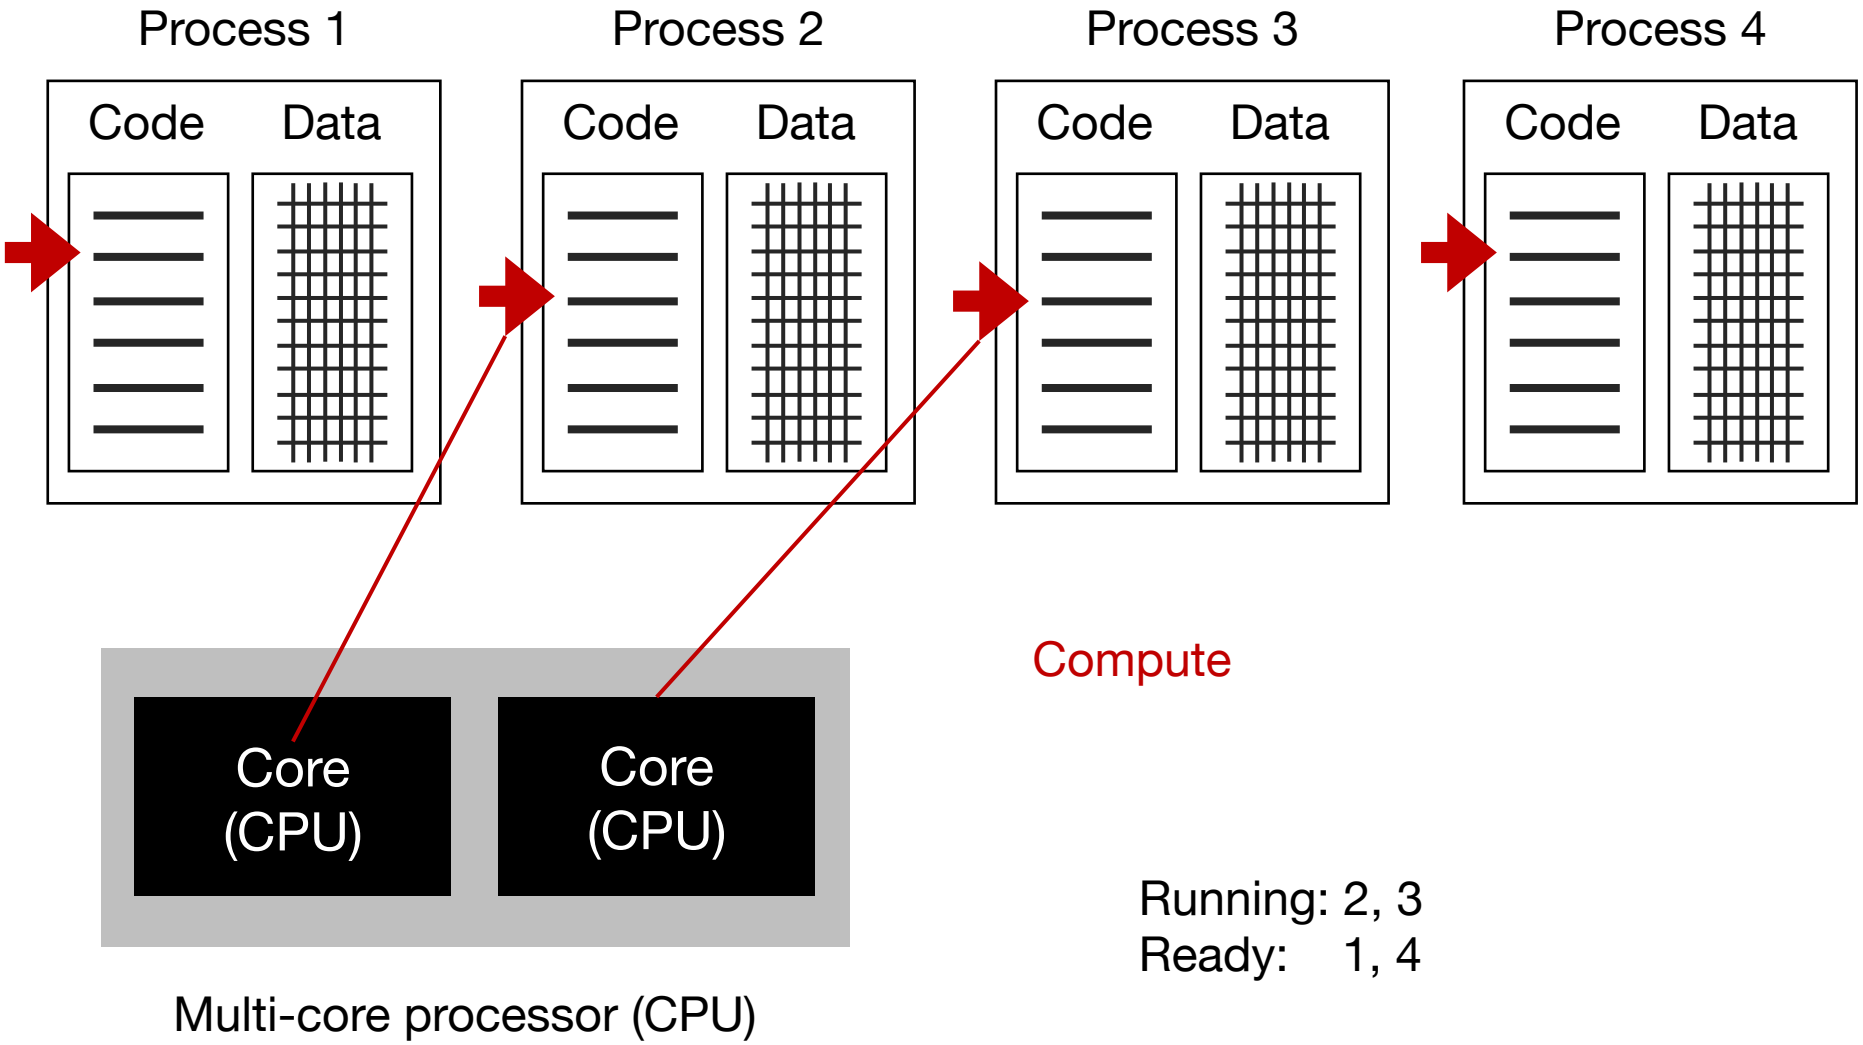

Process-level parallelism

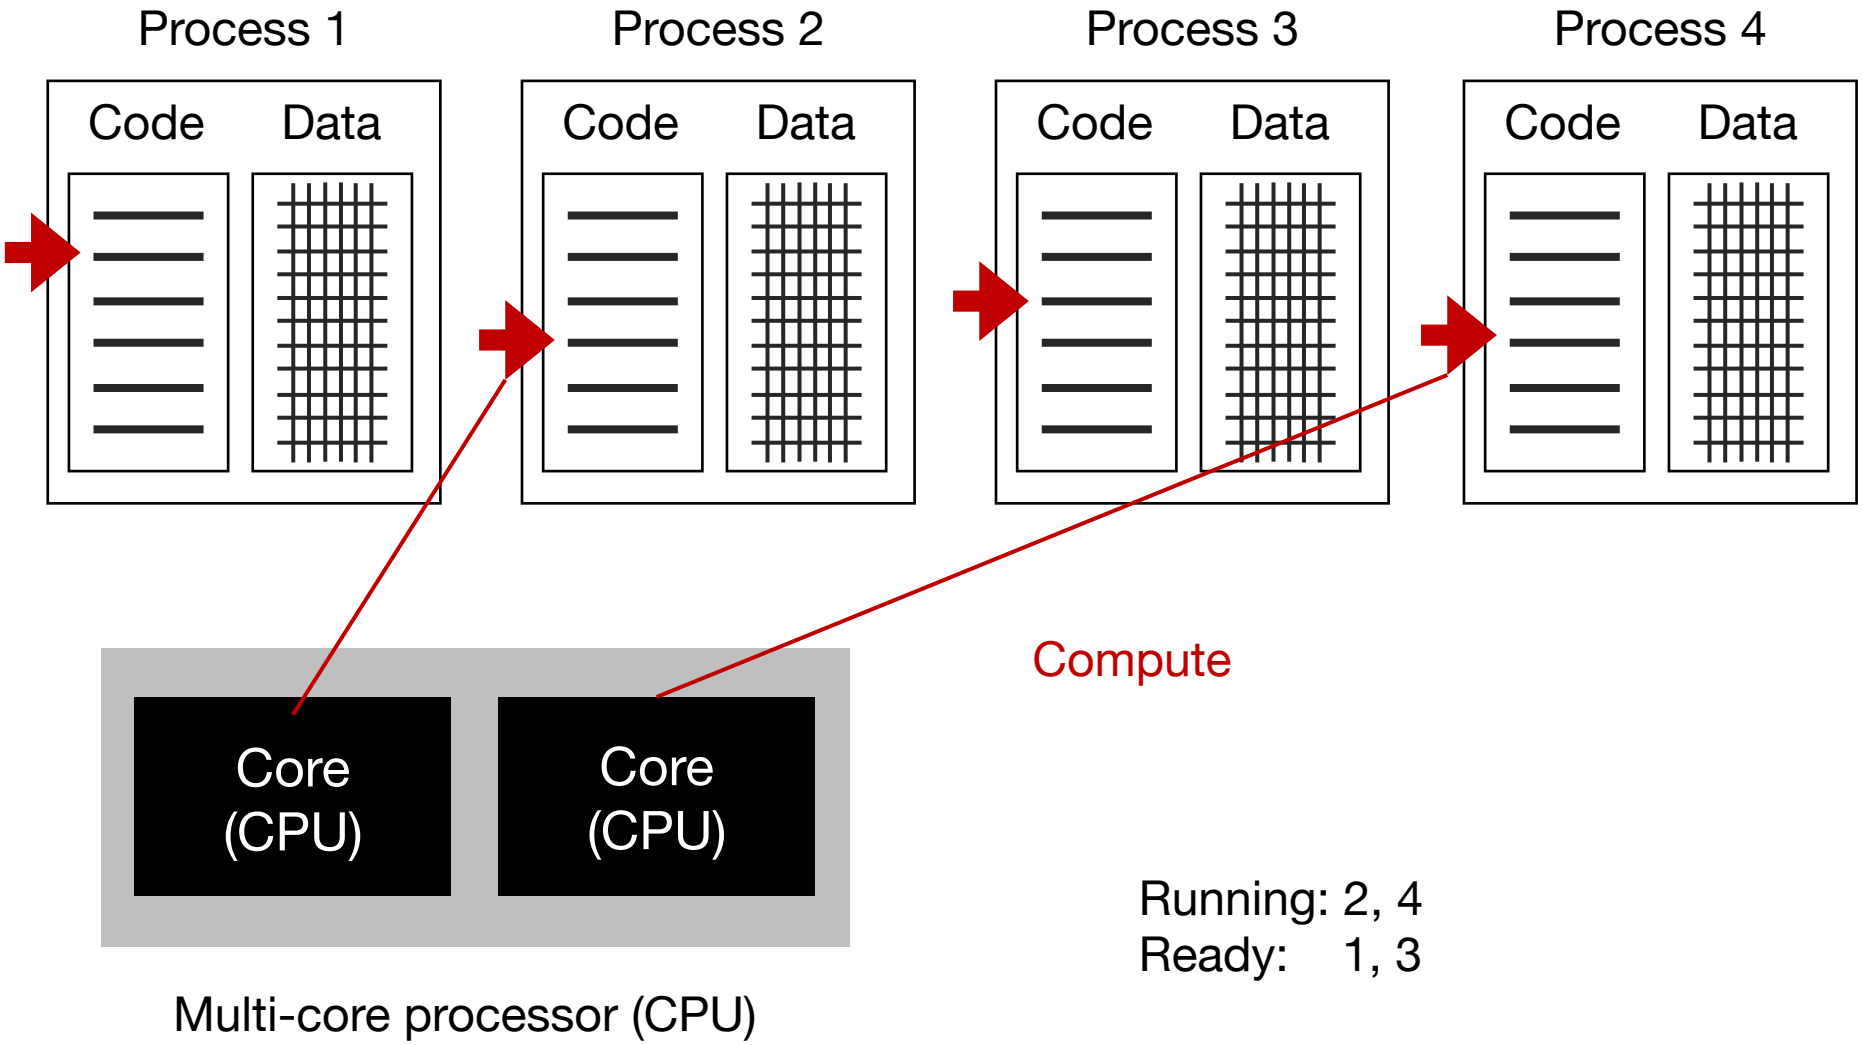

Process-level parallelism

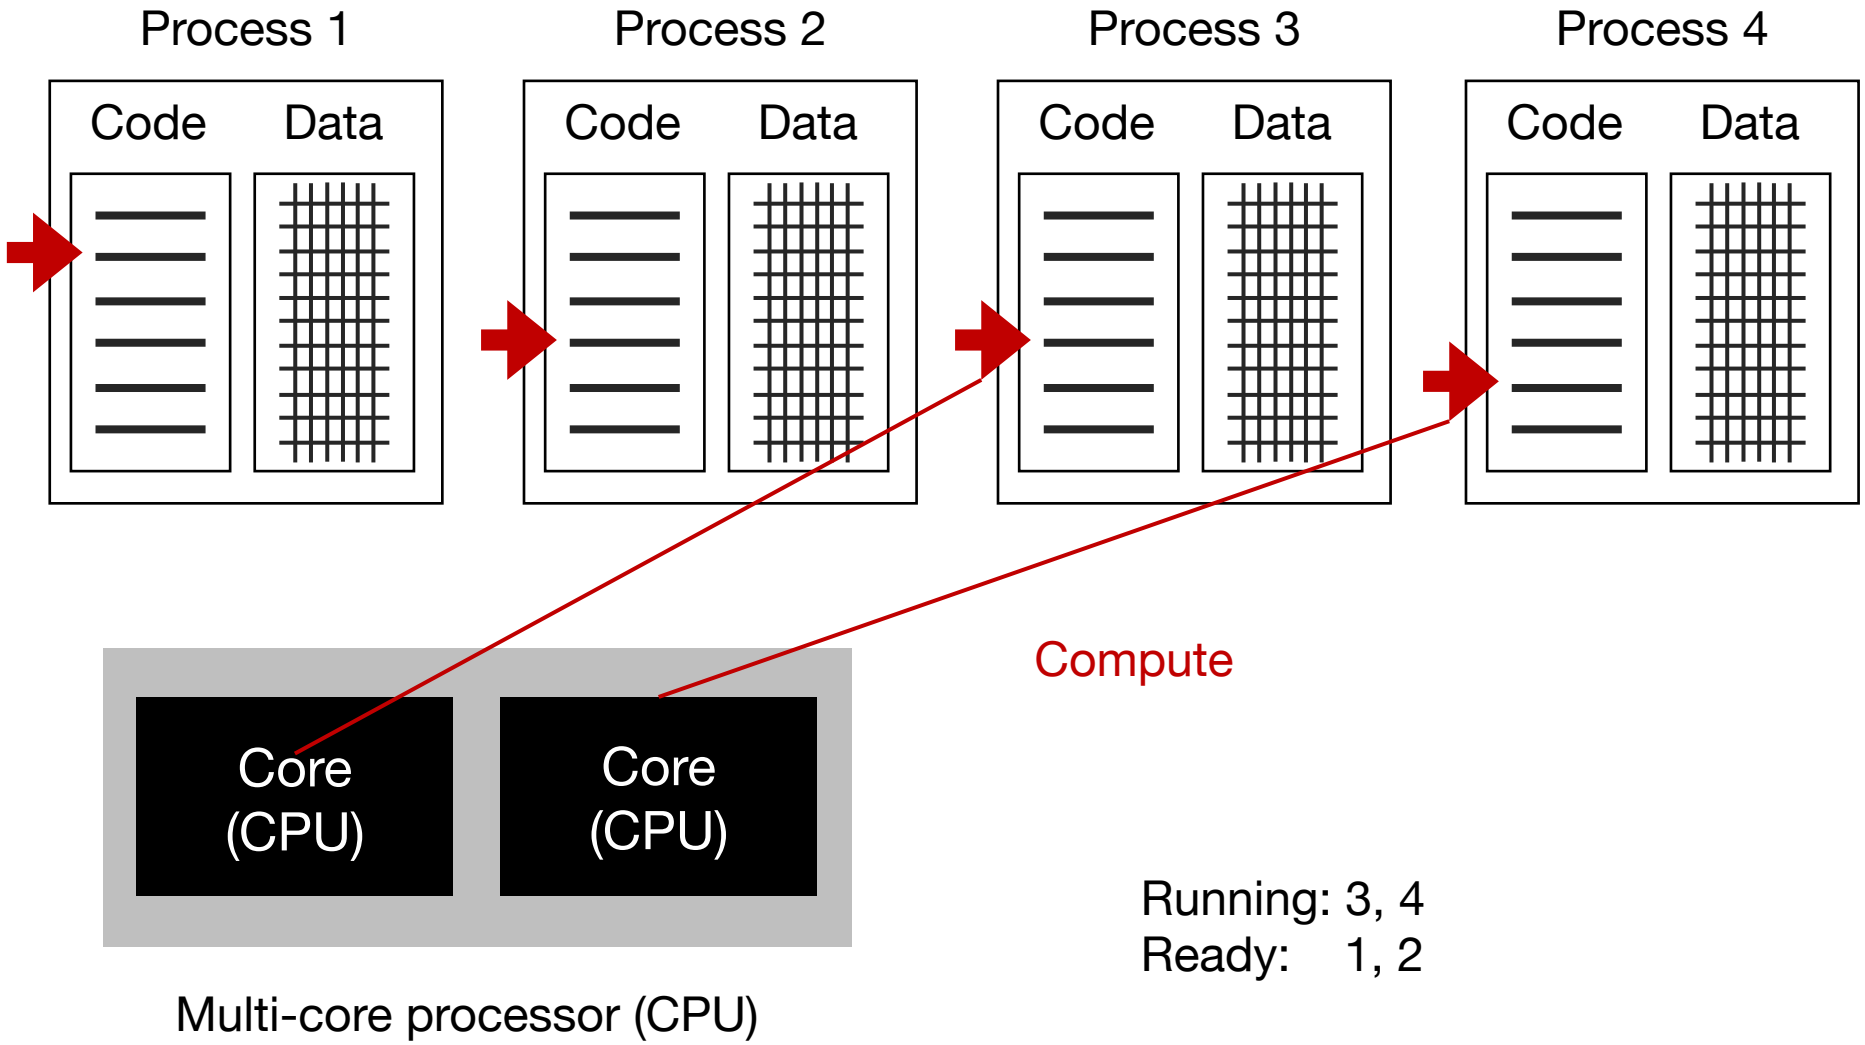

Process-level parallelism

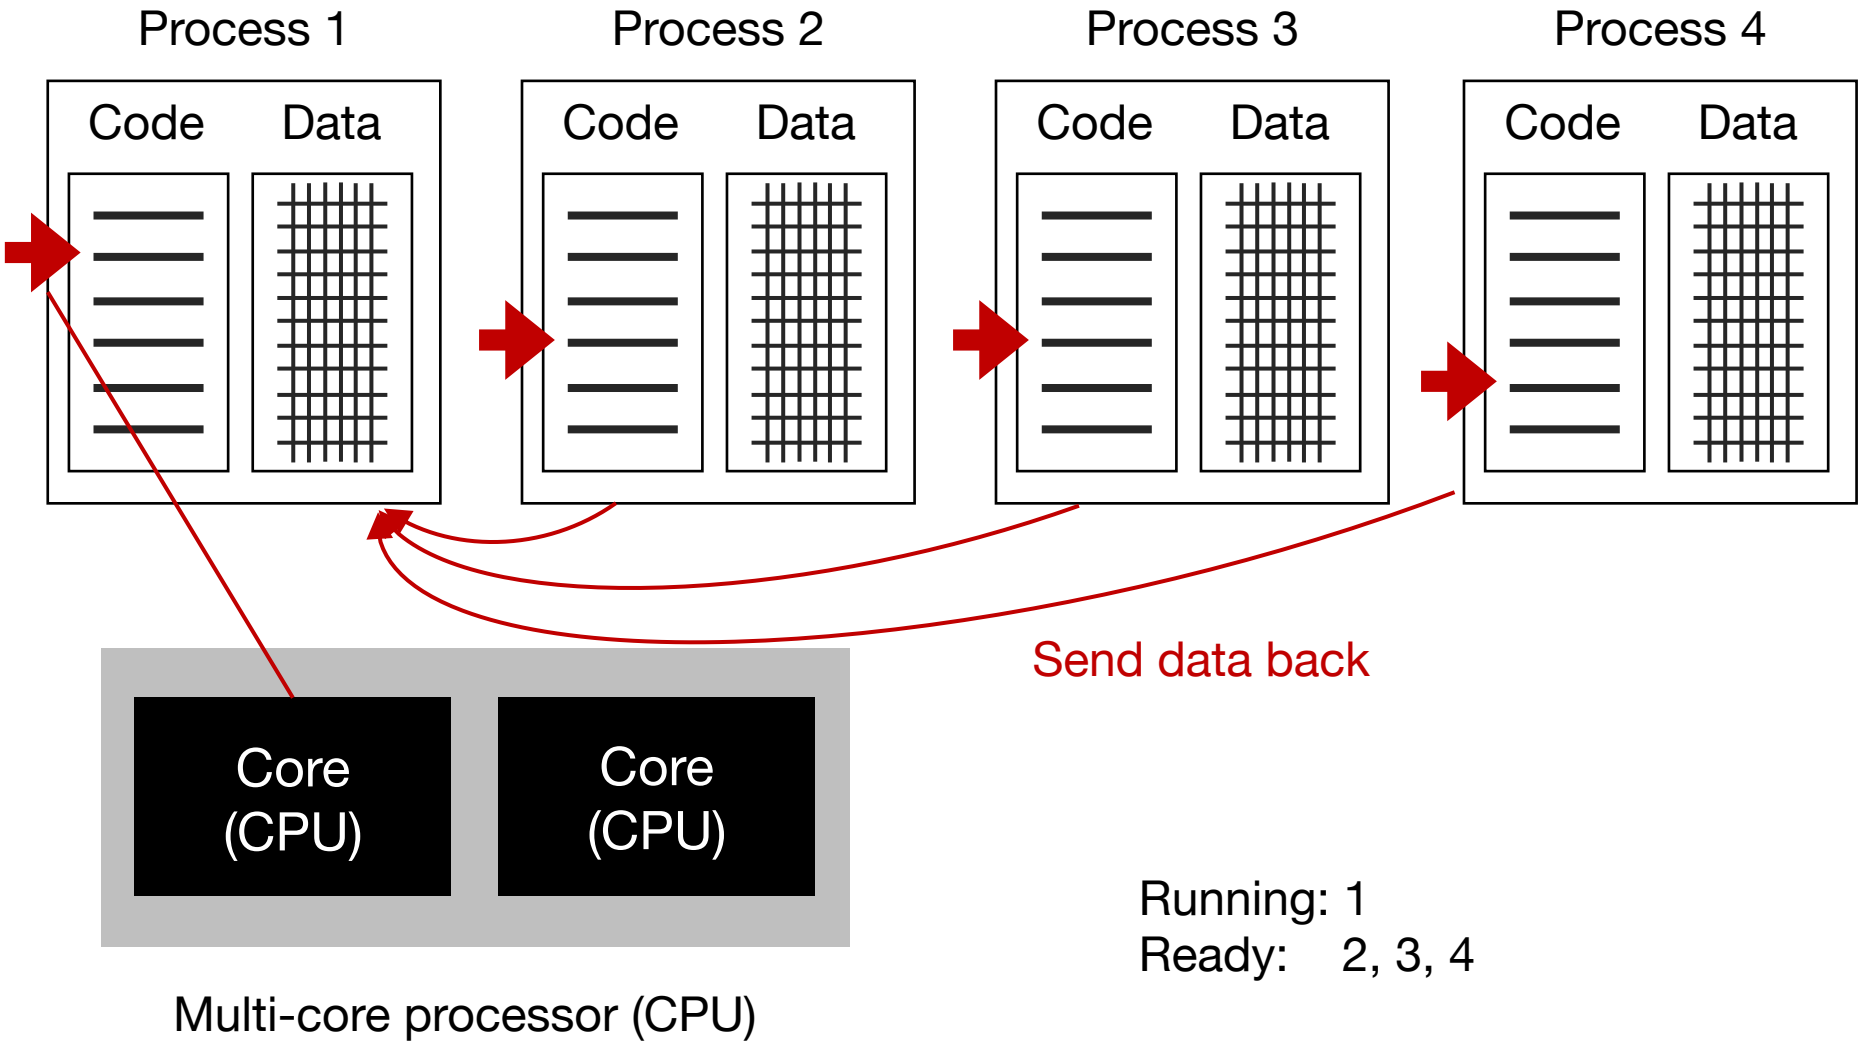

Process-level parallelism

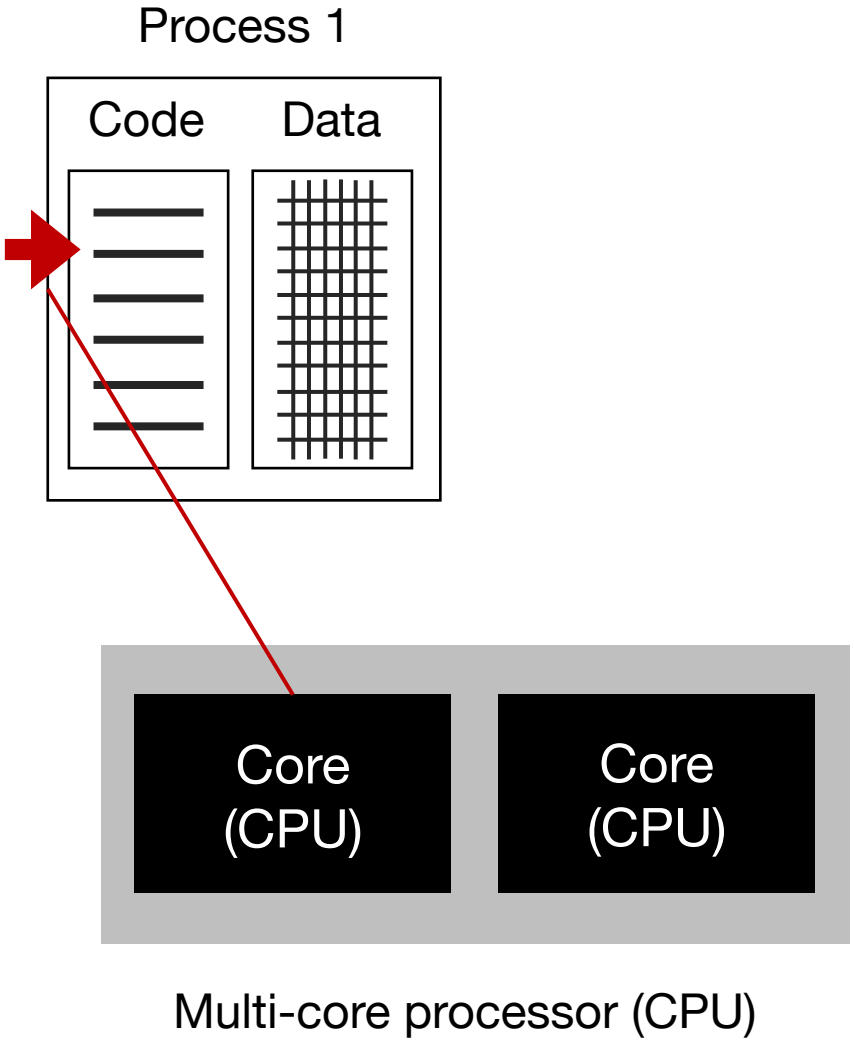

Process-level parallelism in Python

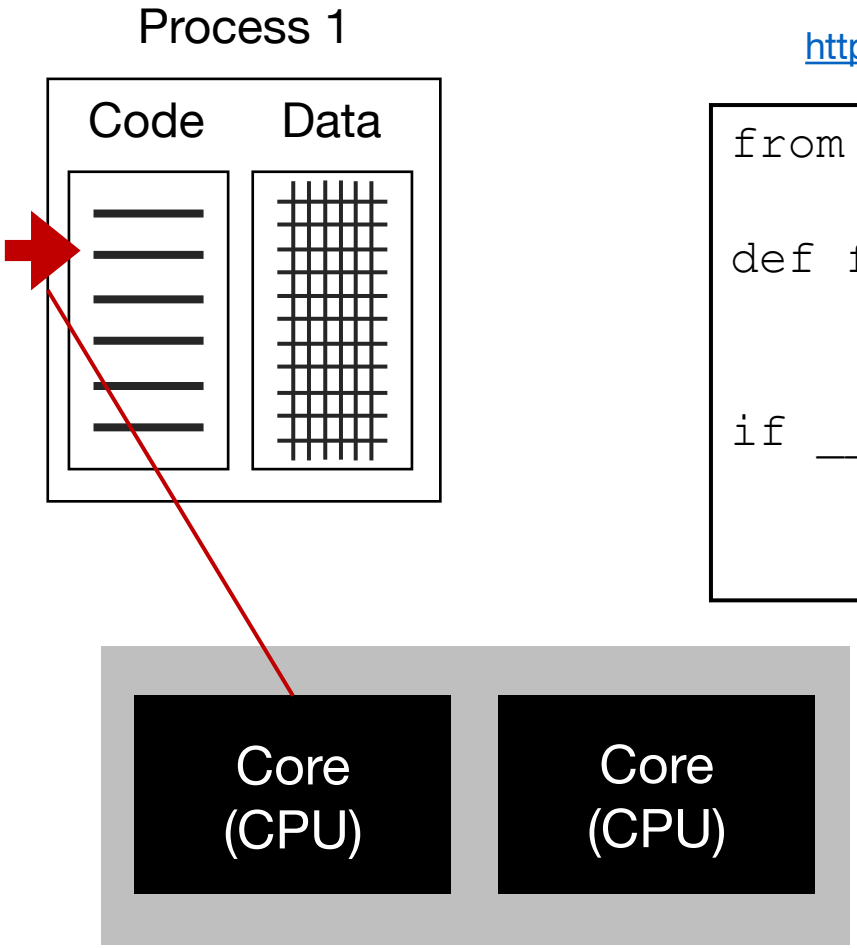

https://docs.python.org/3/library/multiprocessing.html

```
from multiprocessing import Pool

def f(x):
     return x*x

if __name__ == '__main__':
     with Pool(4) as p:
          print(p.map(f, [1,2,3]))
```

Parallel execution models

- Process-level parallelism
- **Thread-level parallelism**
- Task-level parallelism

Thread-level parallelism

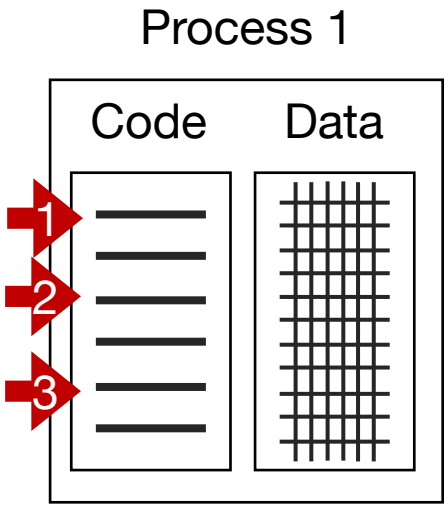

Threads give us multiple instruction pointers in a process, allowing us to execute multiple parts of the code at the same time!

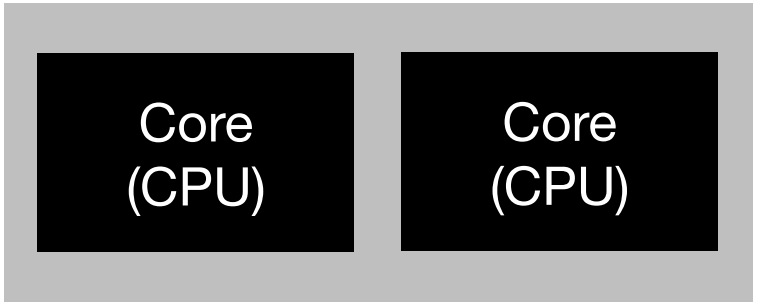

Multi-core processor (CPU)

Thread-level parallelism

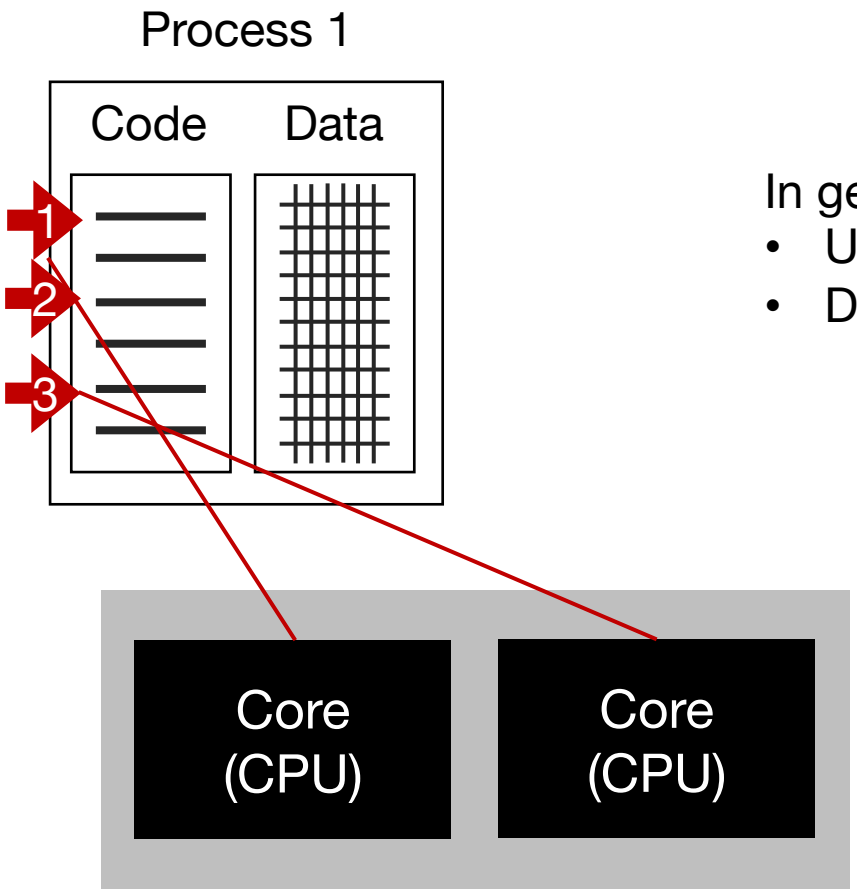

In general, threads help:

- Use multiple cores
- Do useful work when threads are blocking

Running: 1, 3
Ready: 2

Thread-level parallelism in Python

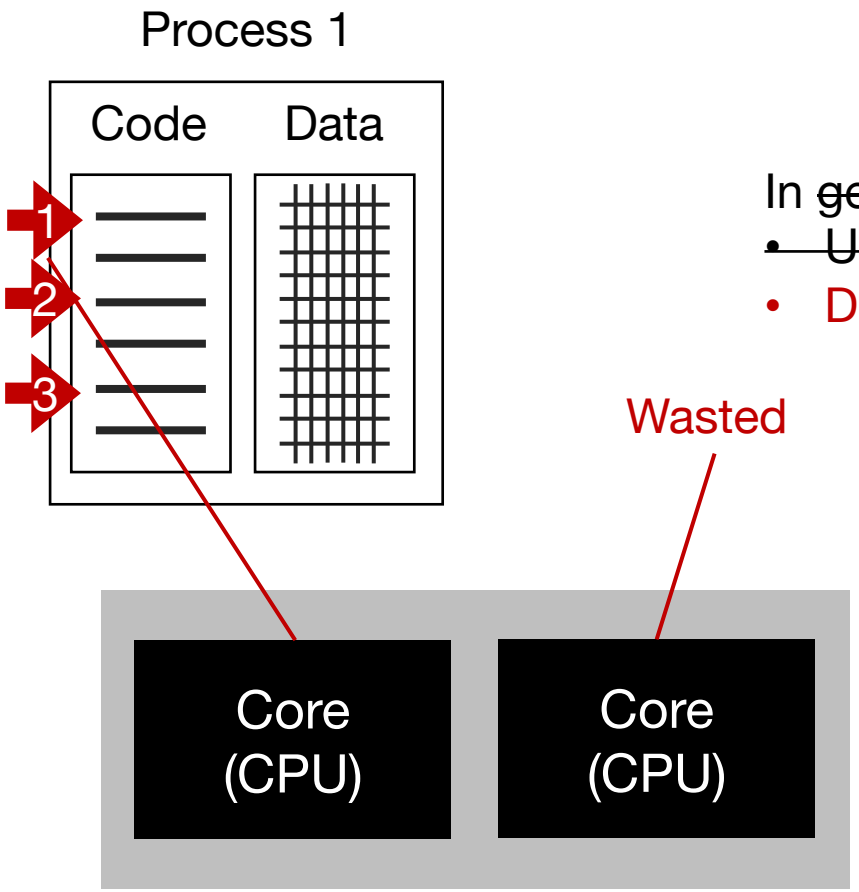

Multi-core processor (CPU)

In ~~general~~ Python, threads help:

- ~~Use multiple cores~~ (b/c of the GIL)
- Do useful work when threads are blocking

https://wiki.python.org/moin/GlobalInterpreterLock

Running: 1
Ready: 3
Blocked: 2

Thread-level parallelism in Python

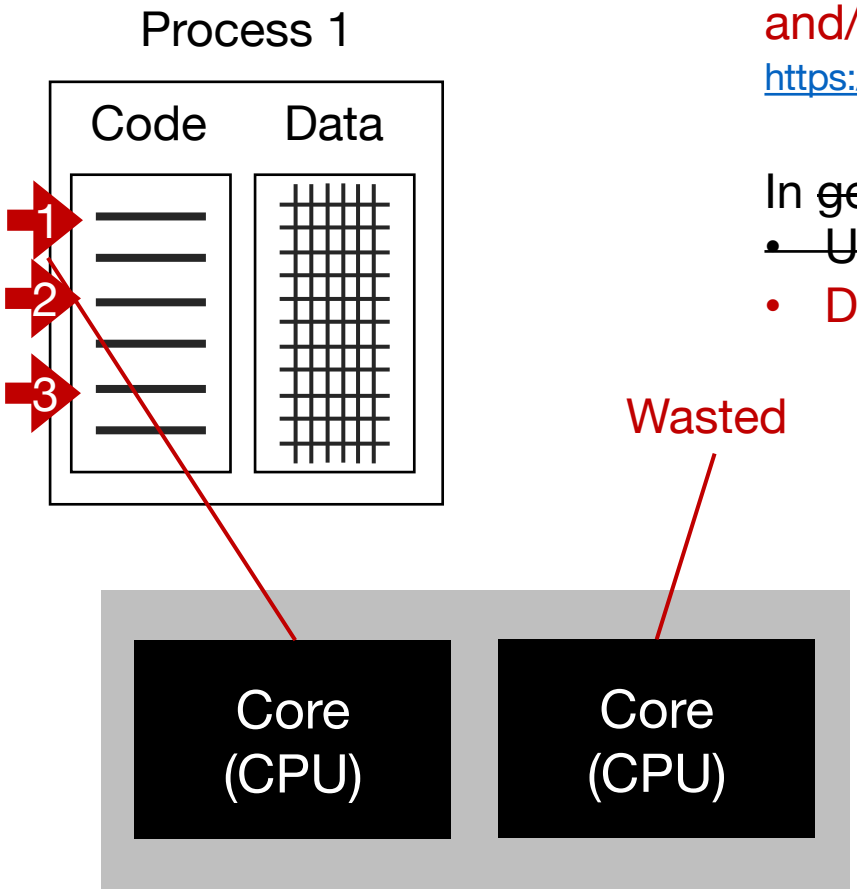

Recommendation: Don't use threads unless you learn a lot on asynchronous processing and/or coroutines
https://docs.python.org/3/library/asyncio-task.html

In ~~general~~ Python, threads help:

- ~~Use multiple cores~~ (b/c of the GIL)
- Do useful work when threads are blocking

https://wiki.python.org/moin/GlobalInterpreterLock

Running: 1
Ready: 3
Blocked: 2

Demo ...

Parallel execution models

- Process-level parallelism
- Thread-level parallelism
- **Task-level parallelism**

Task-level parallelism

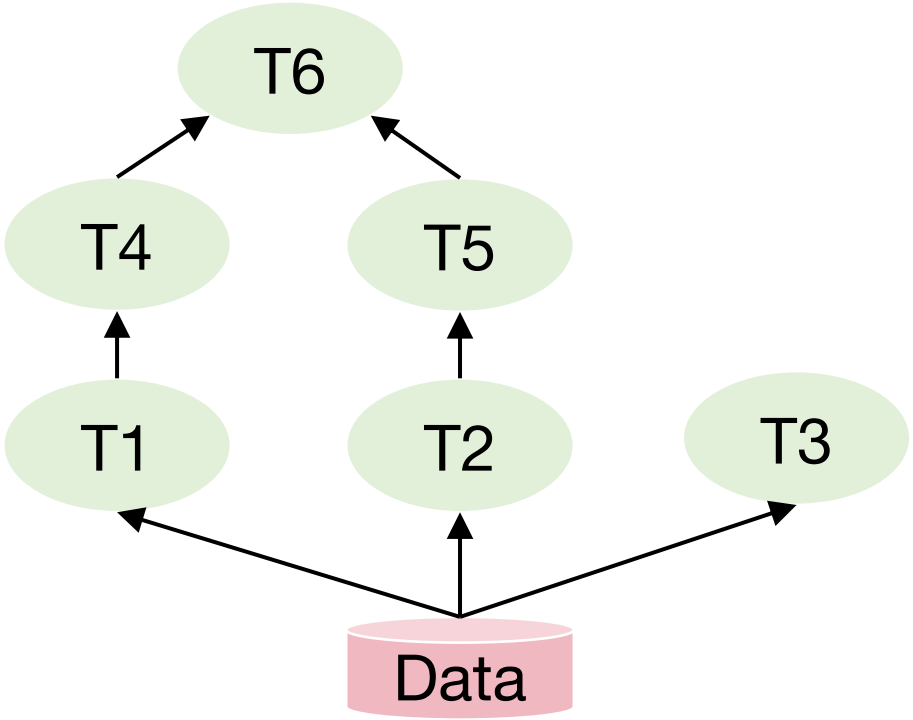

Task DAG
(Directed Acyclic Graph)

Task-level parallelism

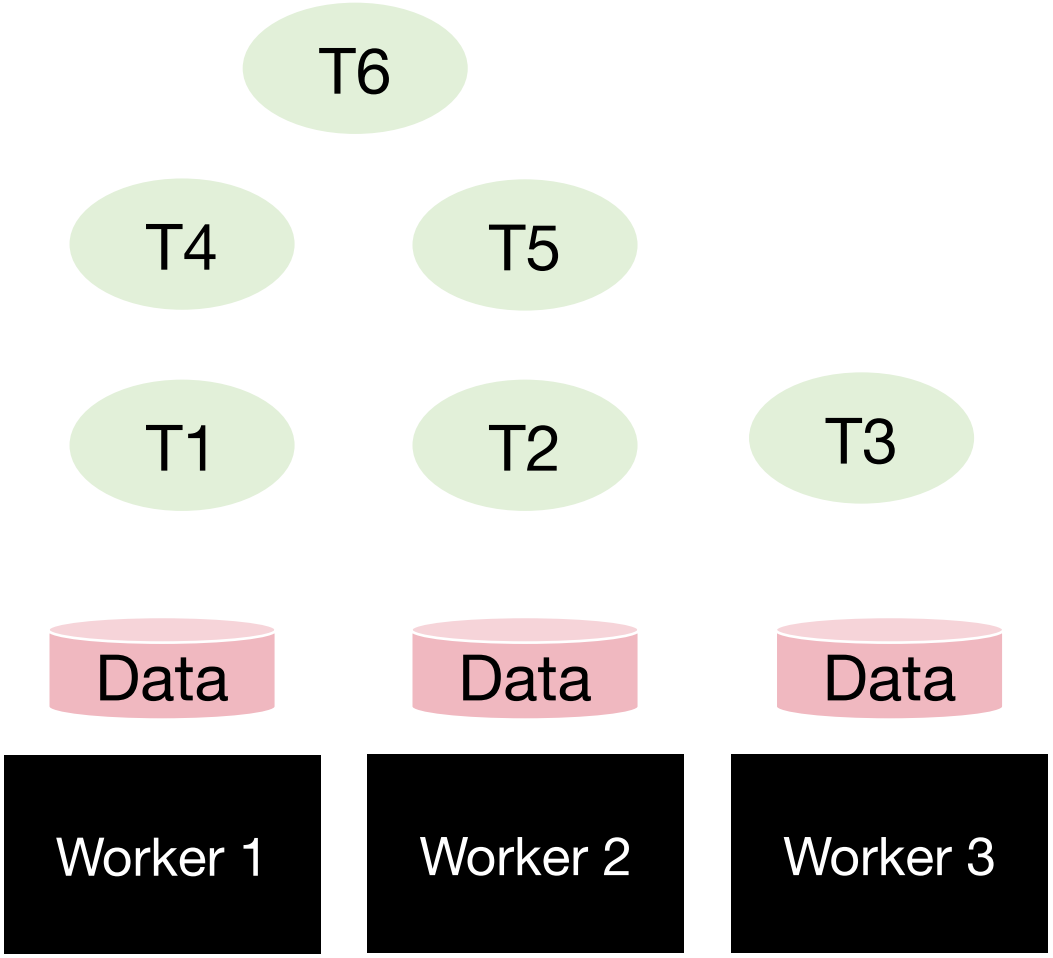

S1: Copy whole dataset to all workers

Task-level parallelism

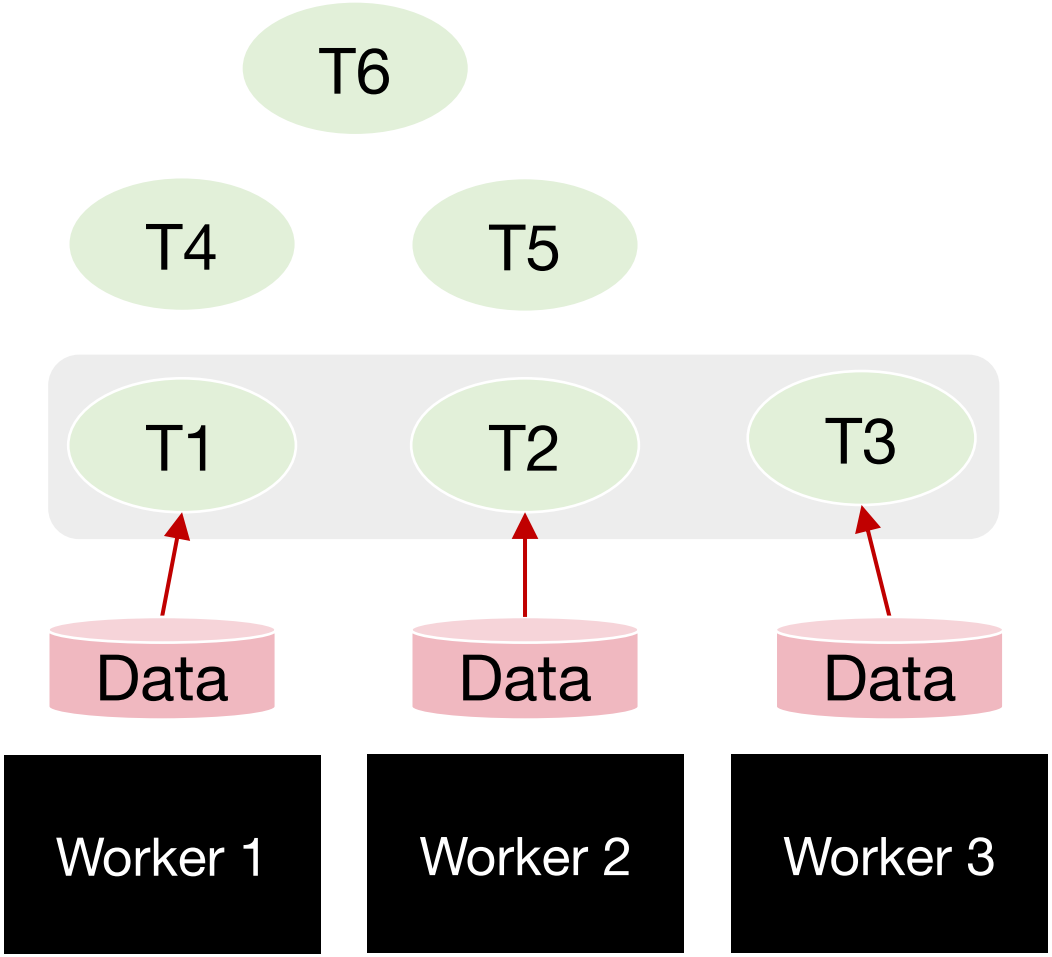

S2: Schedule T1 to W1, T2 to W2, T3 to W3

S1: Copy whole dataset to all workers

Task-level parallelism

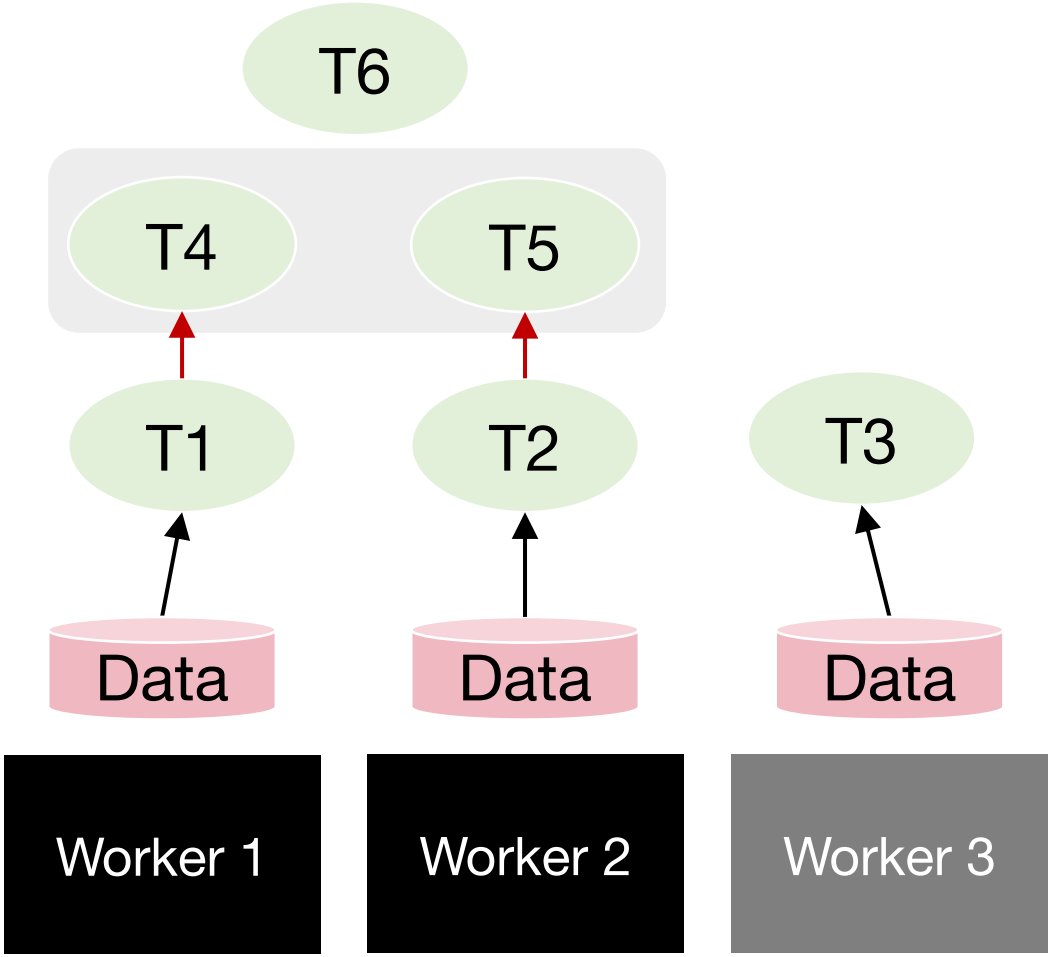

S3: Run T4 after T1 on W1, run T5 after T2 on W2; after T3, W3 is idle

S2: Schedule T1 to W1, T2 to W2, T3 to W3

S1: Copy whole dataset to all workers

Task-level parallelism

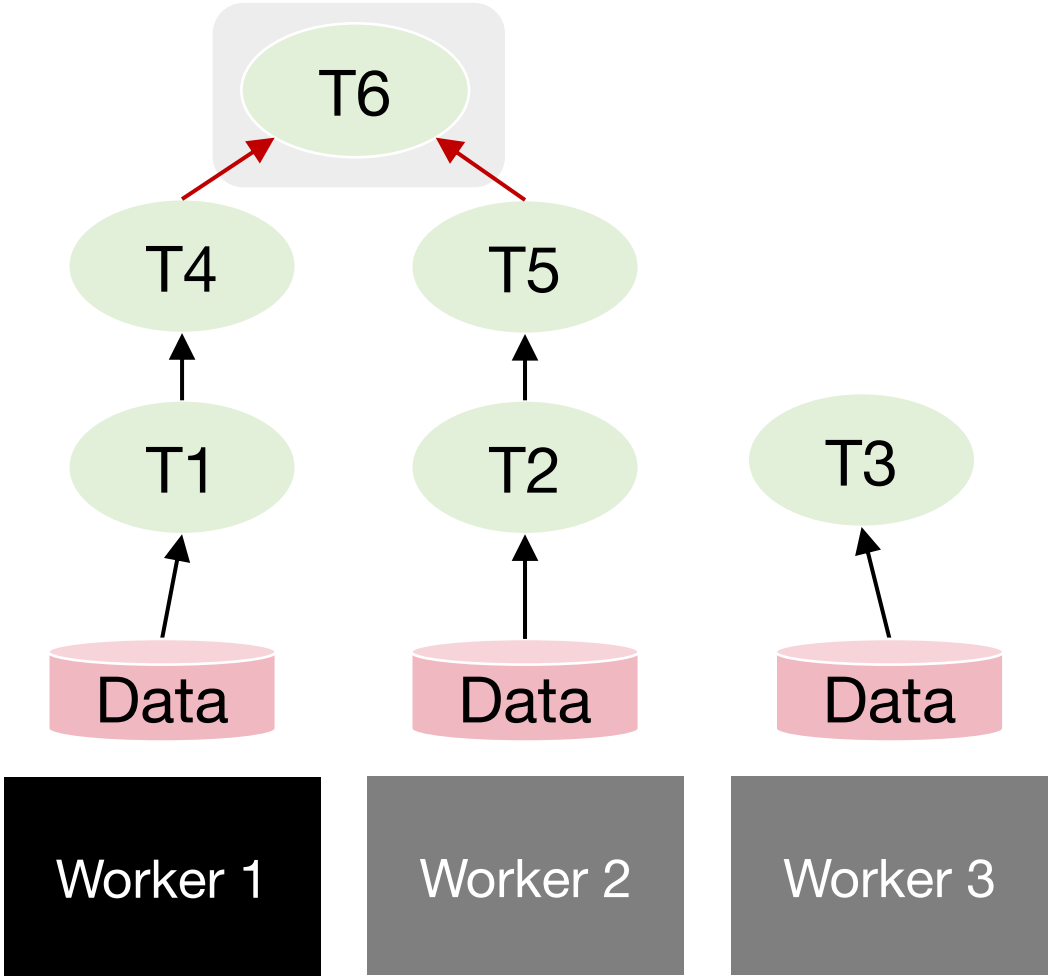

S4: After T4 and T5 ends, run T6 on W1; after T5, W2 is idle

S3: Run T4 after T1 on W1, run T5 after T2 on W2; after T3, W3 is idle

S2: Schedule T1 to W1, T2 to W2, T3 to W3

S1: Copy whole dataset to all workers

Task-level parallelism

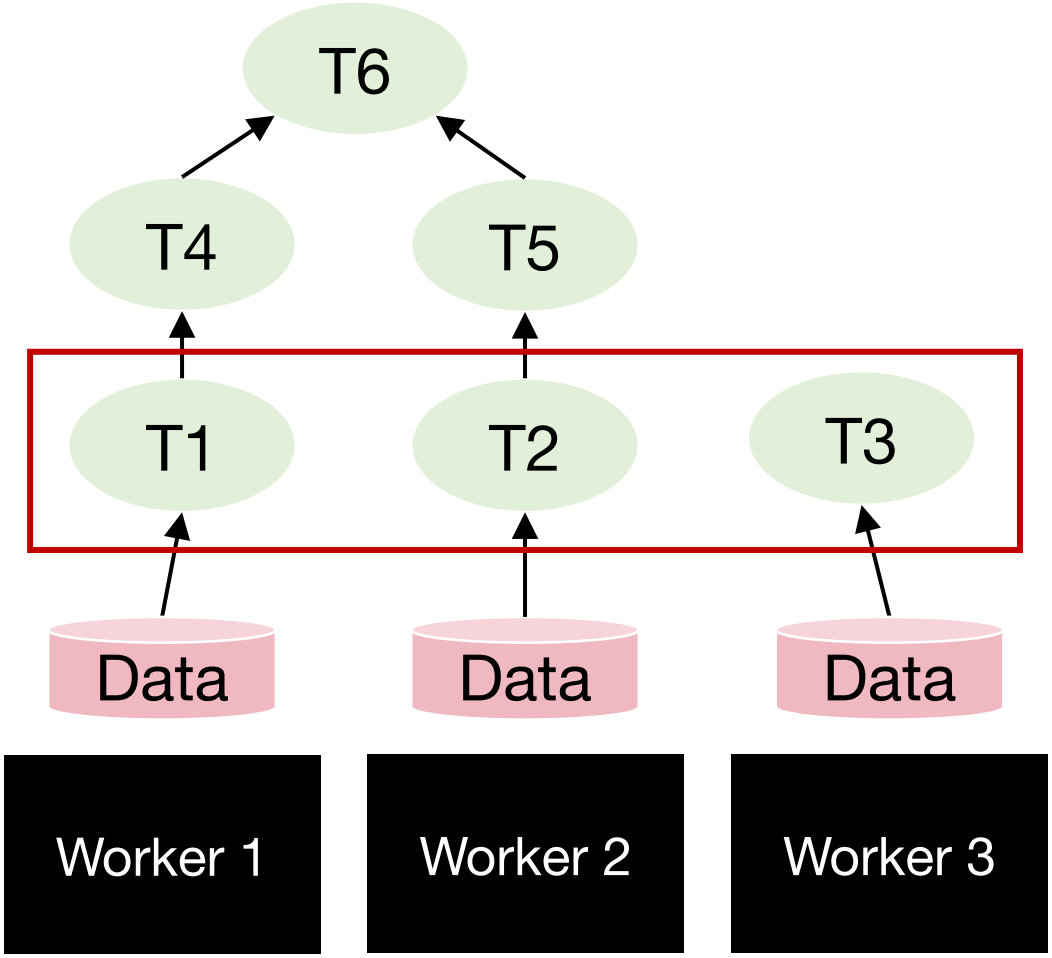

Degree of parallelism is the largest amount of parallelism possible in the DAG:

- How many tasks can be run in parallel at most

Task-level parallelism

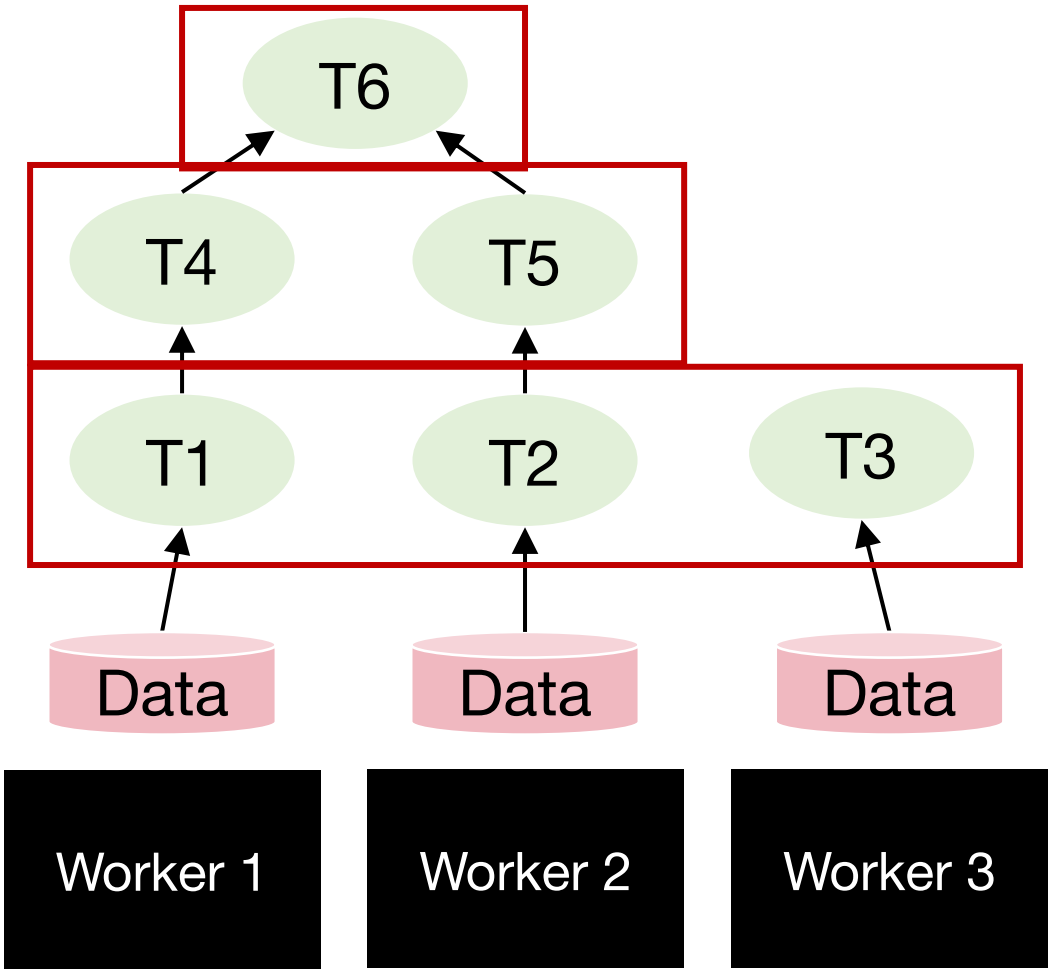

Observations:

Resource wastage on idle workers

Overtime degree of parallelism drops!

Degree of parallelism is the largest amount of parallelism possible in the DAG:

- How many tasks can be run in parallel at most

Quantify benefit of parallelism: Speedup

$$\textbf{Speedup} = \frac{\text{Completion time given 1 worker}}{\text{Completion time given N worker}}$$

Quantify benefit of parallelism: Speedup

$$\textbf{Speedup} = \frac{\text{Completion time given 1 worker}}{\text{Completion time given N worker}}$$

Q: Given N workers, can we get a speedup of N?

Quantify speedup

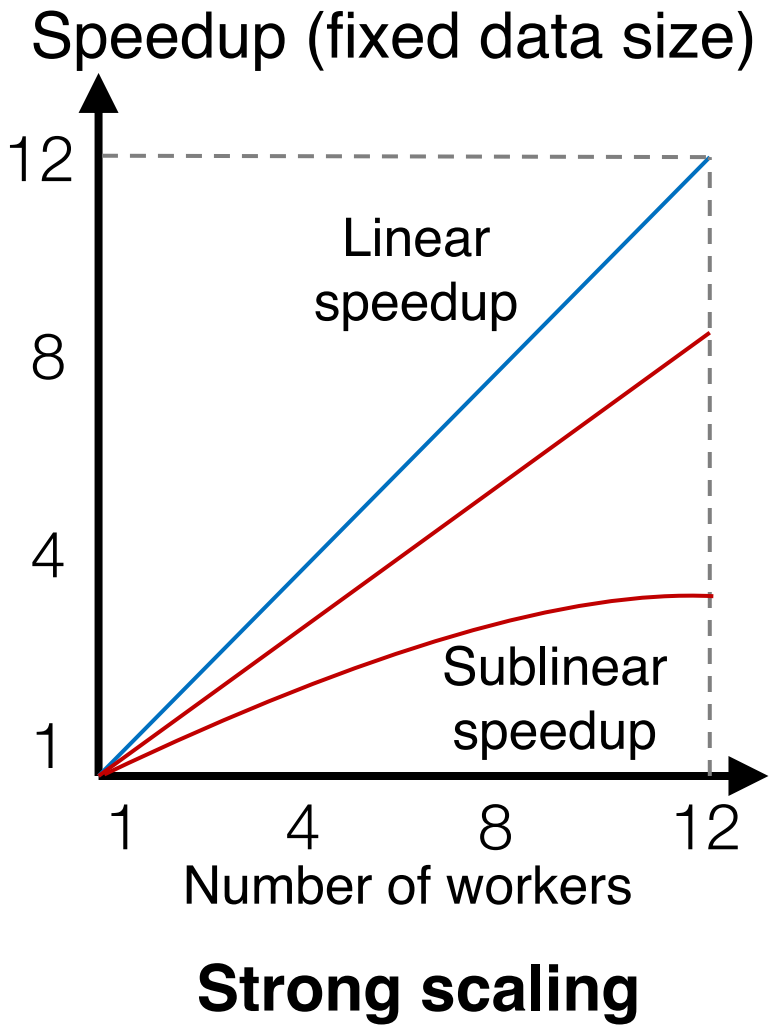

Strong scaling

Quantify speedup

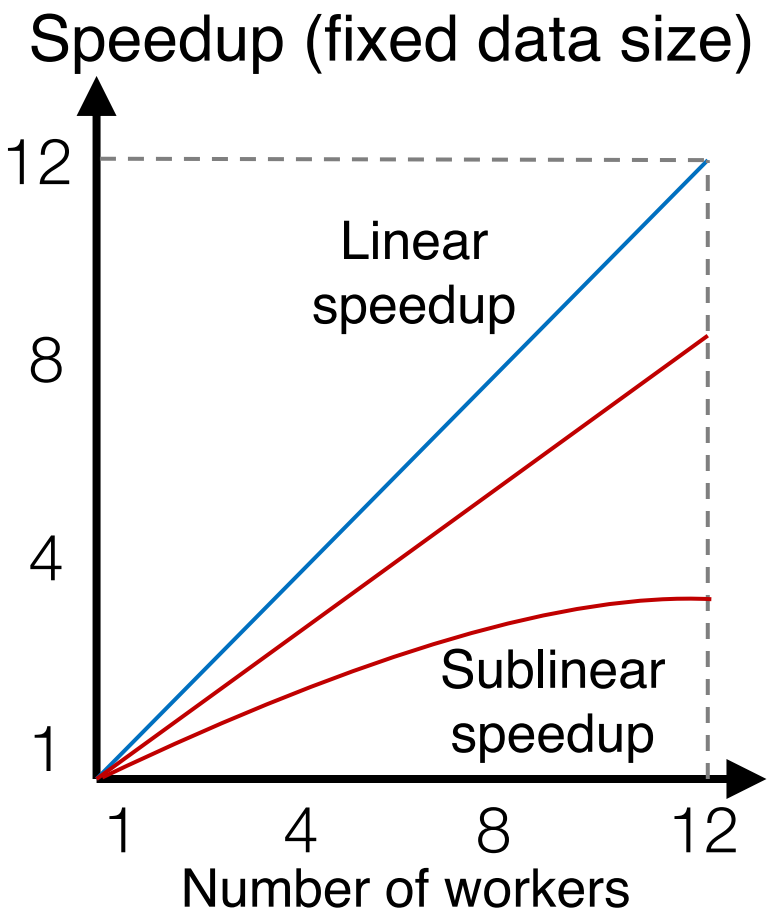

Strong scaling

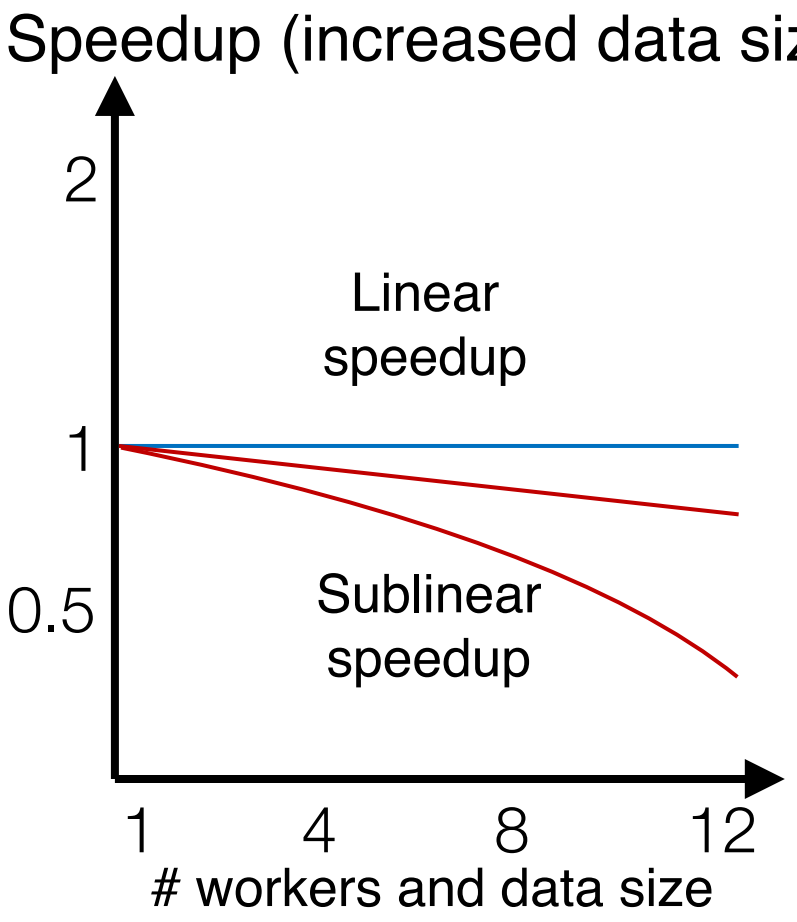

Weak scaling

Idle resources in task-level parallelism

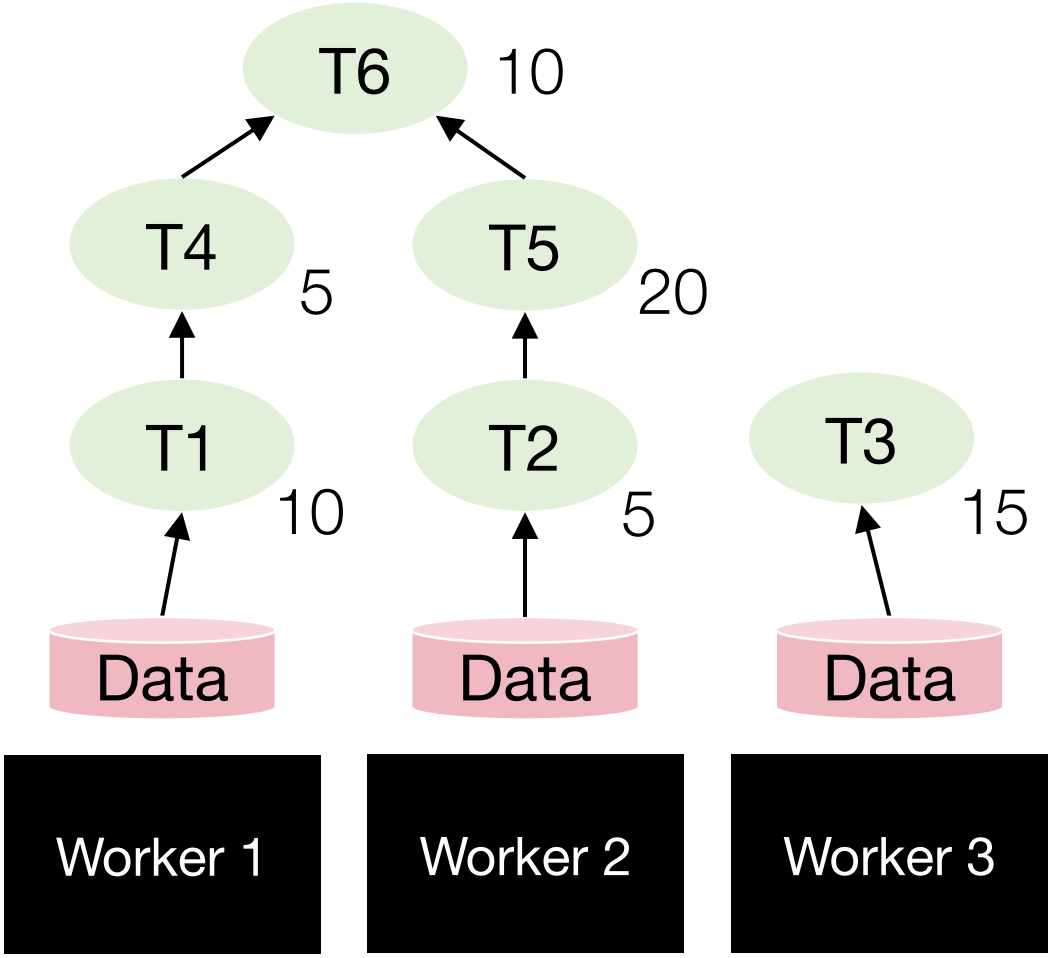

Task completion time varies

Idle resources in task-level parallelism

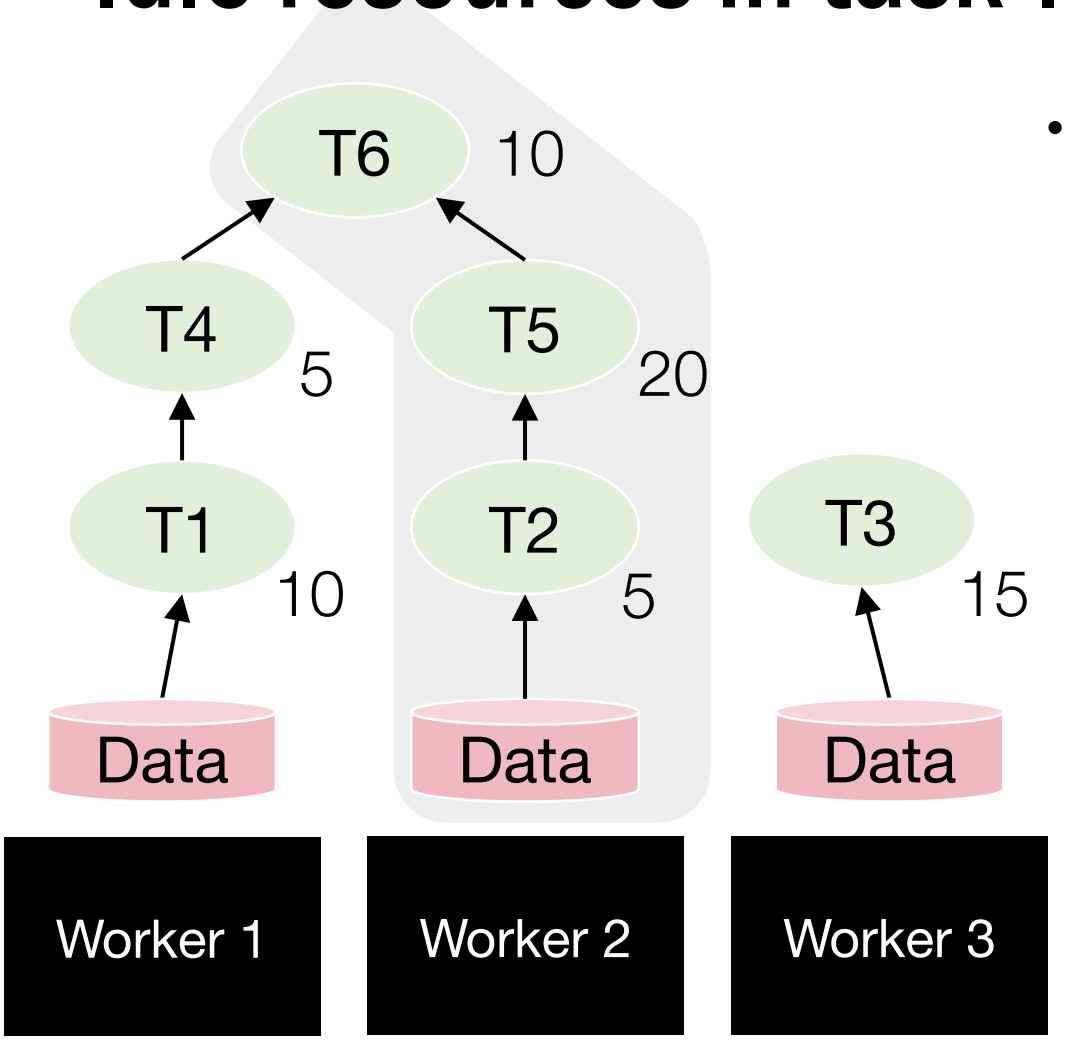

- Job completion time is always bounded by the **longest path** in the DAG

Task completion time varies

Idle resources in task-level parallelism

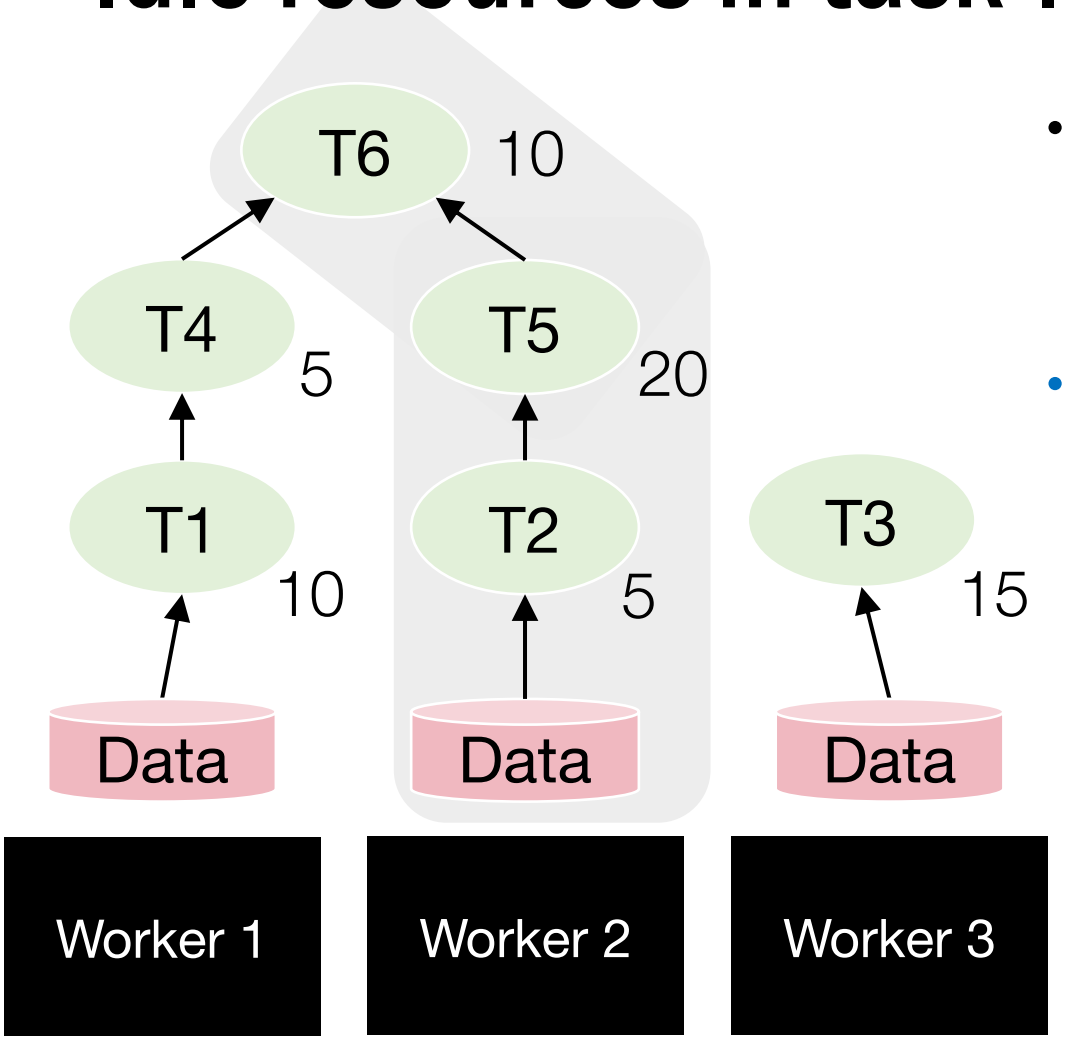

- Job completion time is always bounded by the **longest path** in the DAG
- **Potential optimization:** The scheduler can elastically release a worker if it knows the worker will be idle till the end
  - Can **save $ cost** in cloud

Task completion time varies

Idle resources in task-level parallelism

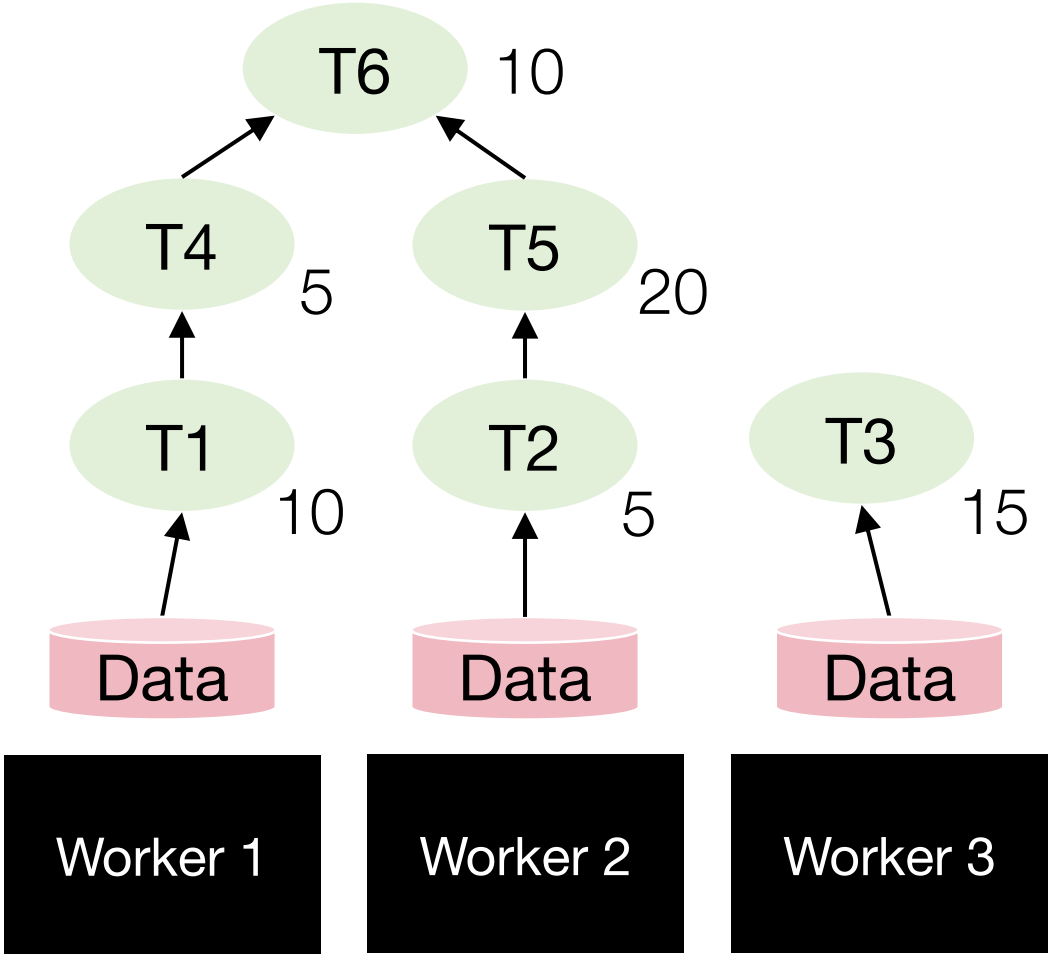

Q: What's the job completion time with 1 worker?

Q: What's the job completion time with 3 worker?

Q: What's the speedup?

Task parallelism in Dask

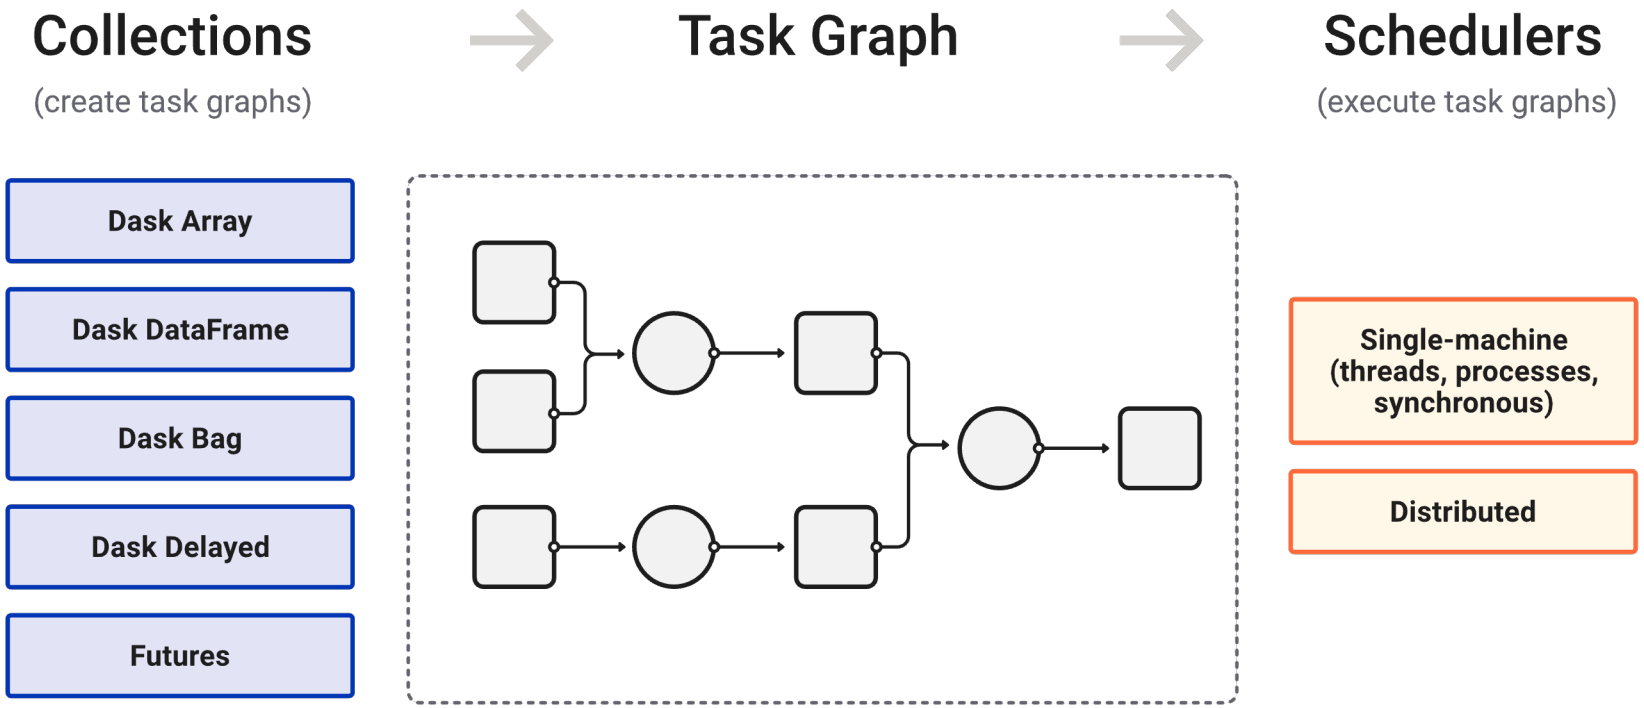

* https://docs.dask.org/en/stable/
* https://docs.dask.org/en/stable/scheduling.html

Dask's task graph and workflow

```
import dask
import dask.array as da
x = da.random.normal(size=1_000_000, chunks=100_000)
```

Dask's task graph and workflow

```
import dask
import dask.array as da
x = da.random.normal(size=1_000_000, chunks=100_000)
```

```
data = x.compute()
```

Lazy evaluation: Dask computation can be triggered manually, e.g., `.compute()`

- only when the result is needed

Dask's task graph and workflow

```
import dask
import dask.array as da
x = da.random.normal(size=1_000_000, chunks=100_000)
```

```
data = x.compute()
```

Lazy evaluation: Dask computation can be triggered manually, e.g., `.compute()`

- only when the result is needed

```
dask.visualize(x)
```

Draw the task graph using `.visualize()`

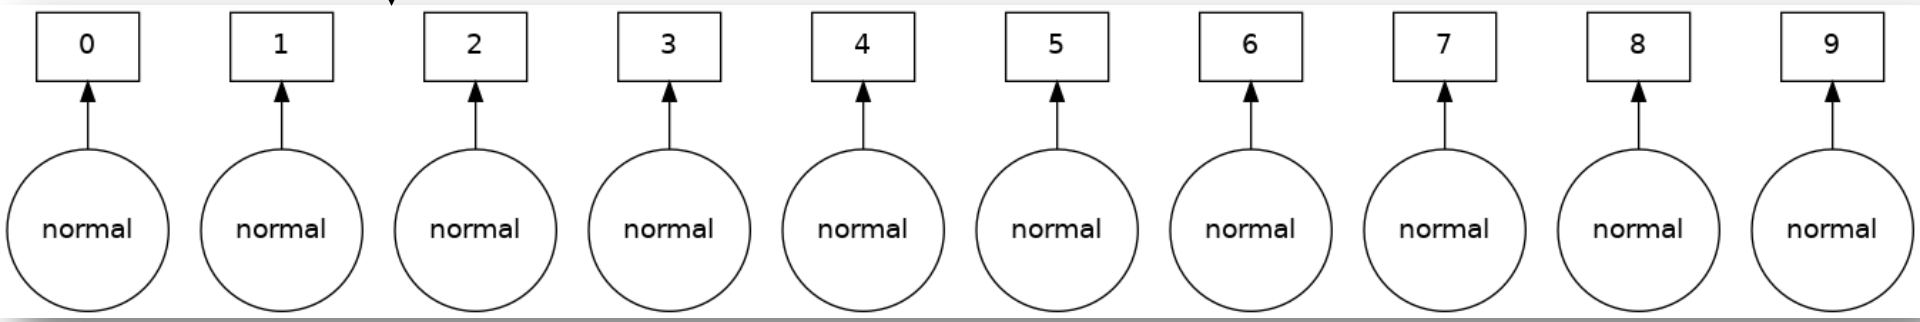

Dask task graph

Demo ...