

Ray Core Internals

DS 5110/CS 5501: Big Data Systems

Spring 2024

Lecture 6b

Yue Cheng

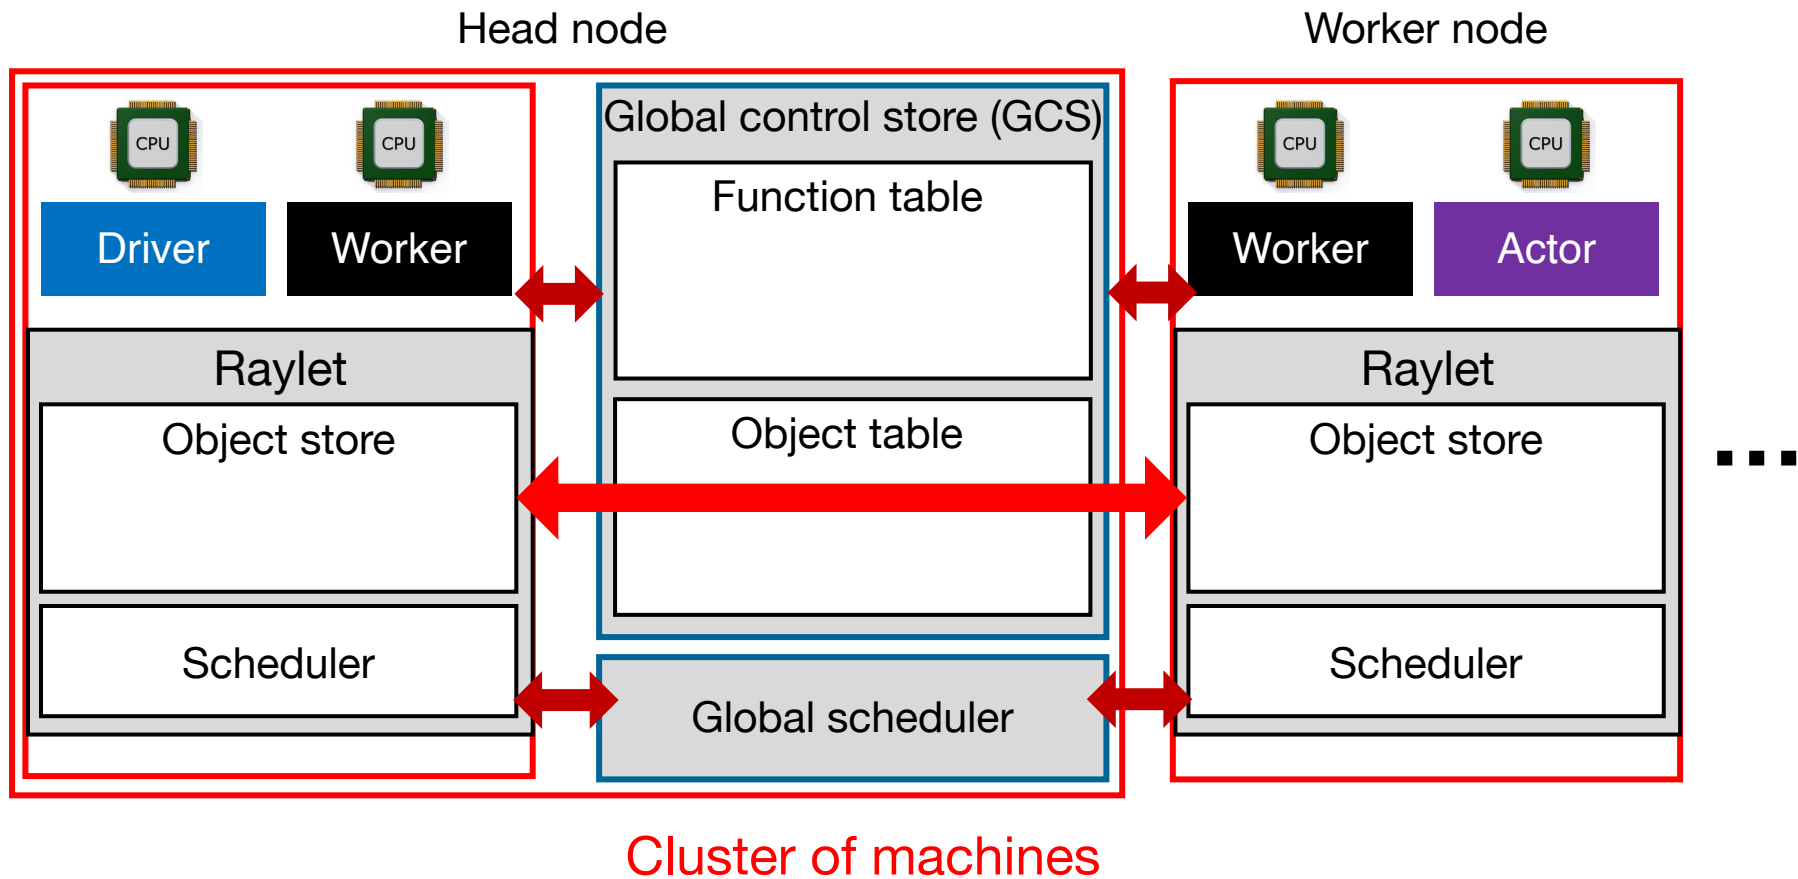


UNIVERSITY
of VIRGINIA

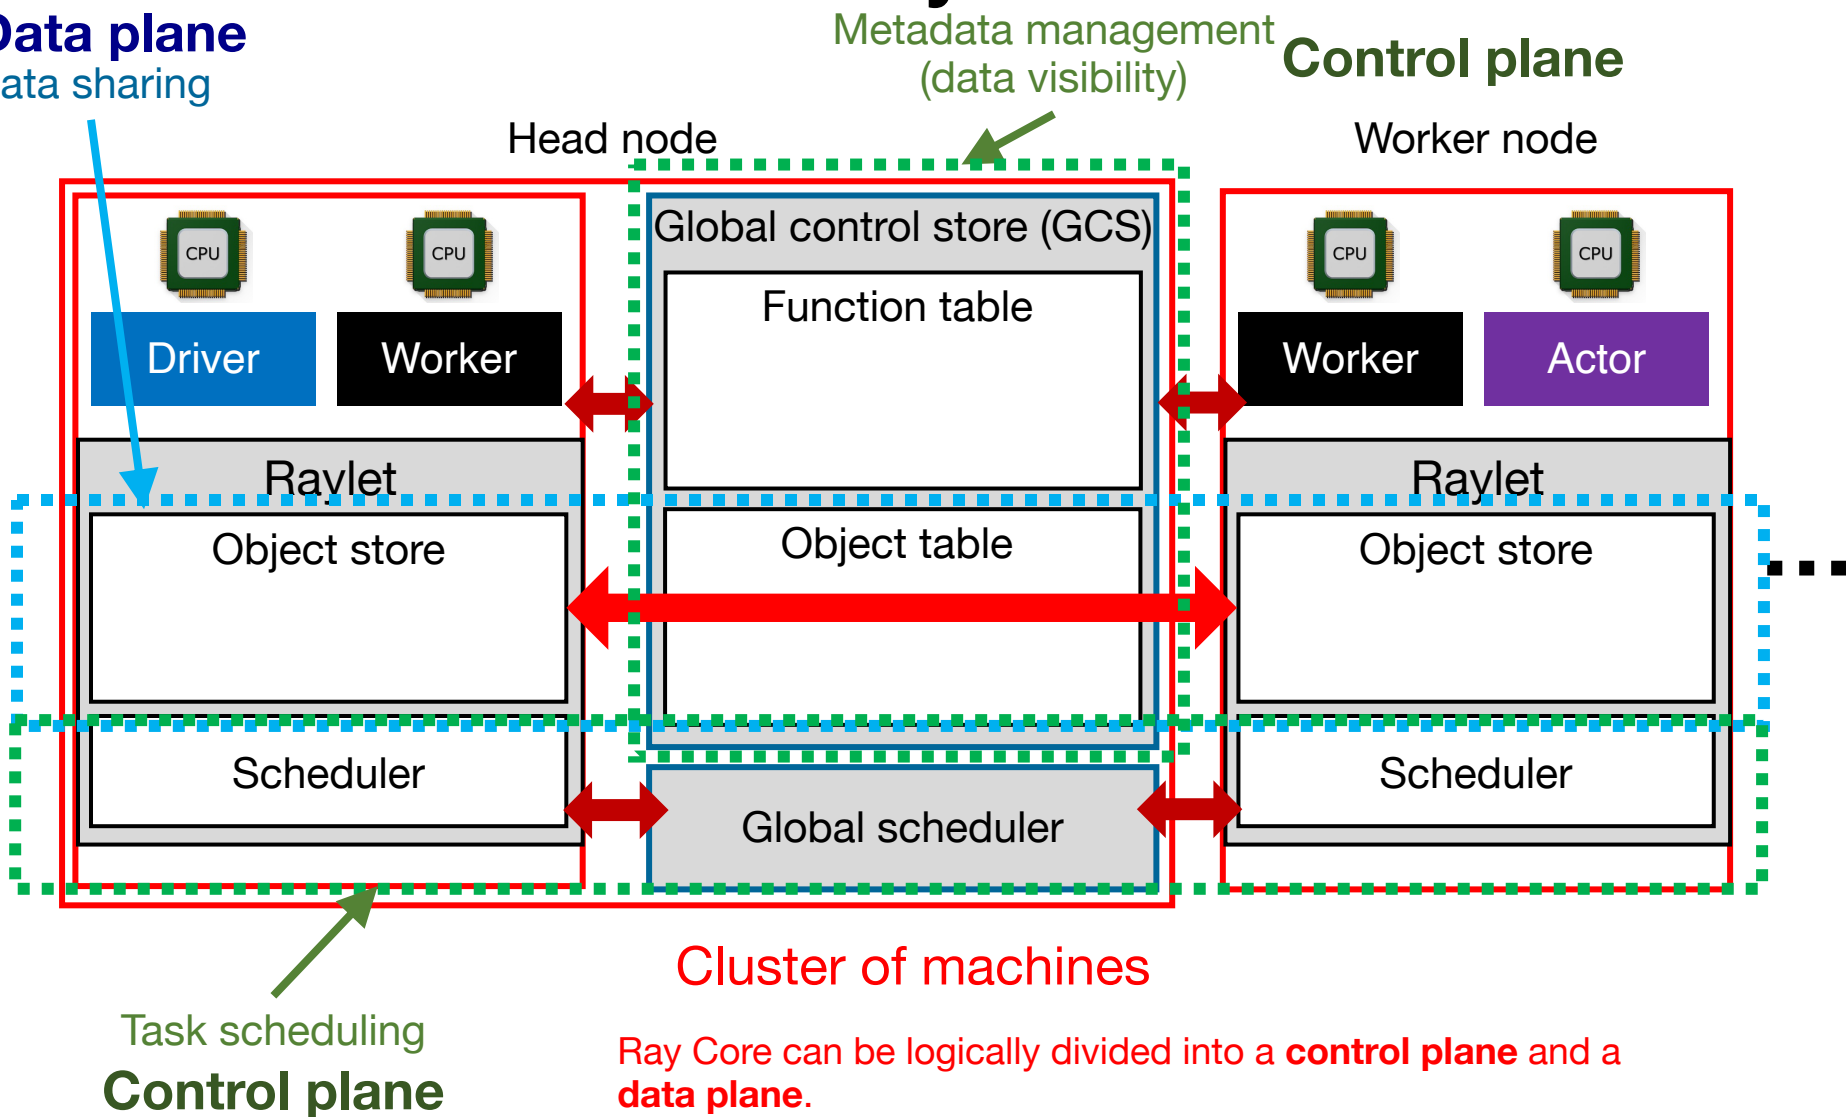
Learning objectives

- Understand how Ray tasks and actors are managed under the hood
- Know the difference between a control plane and a data plane

Under the hood: Ray Core architecture



Under the hood: Ray Core architecture



Life cycle of a remote task

Executing a task remotely

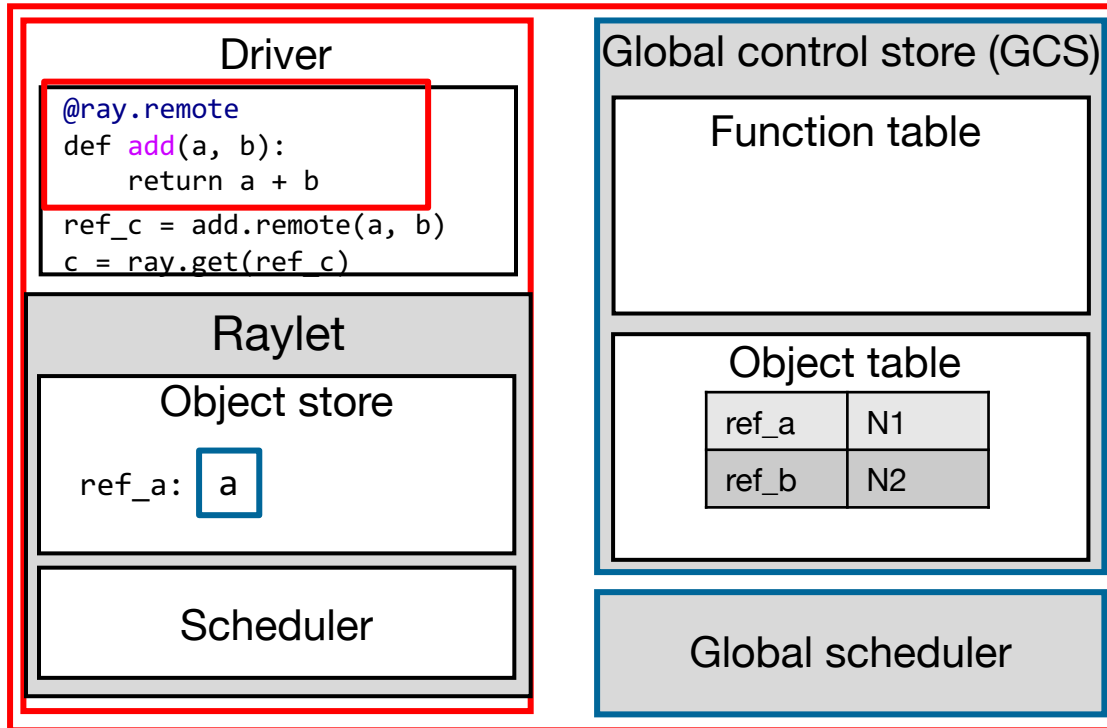
Executing a task remotely →

```
@ray.remote
```

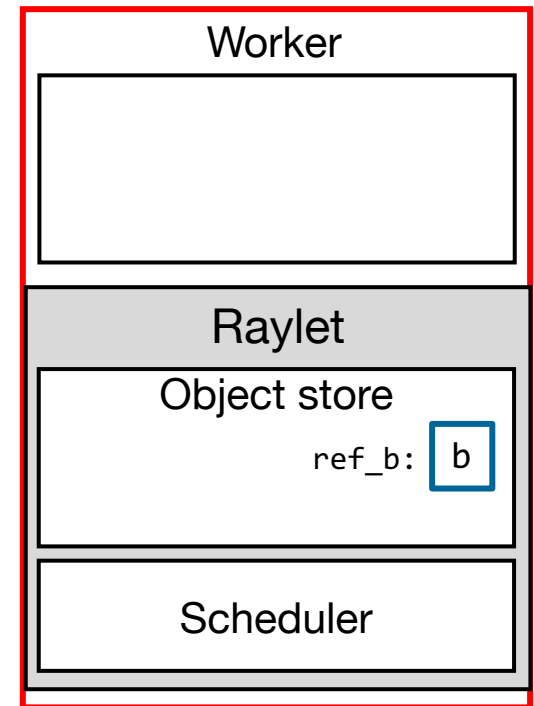
```
def add(a, b):  
    return a + b
```

```
ref_c = add.remote(a, b)  
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



Cluster of machines

Now, the remote task function `add` is initialized...

Executing a task remotely →

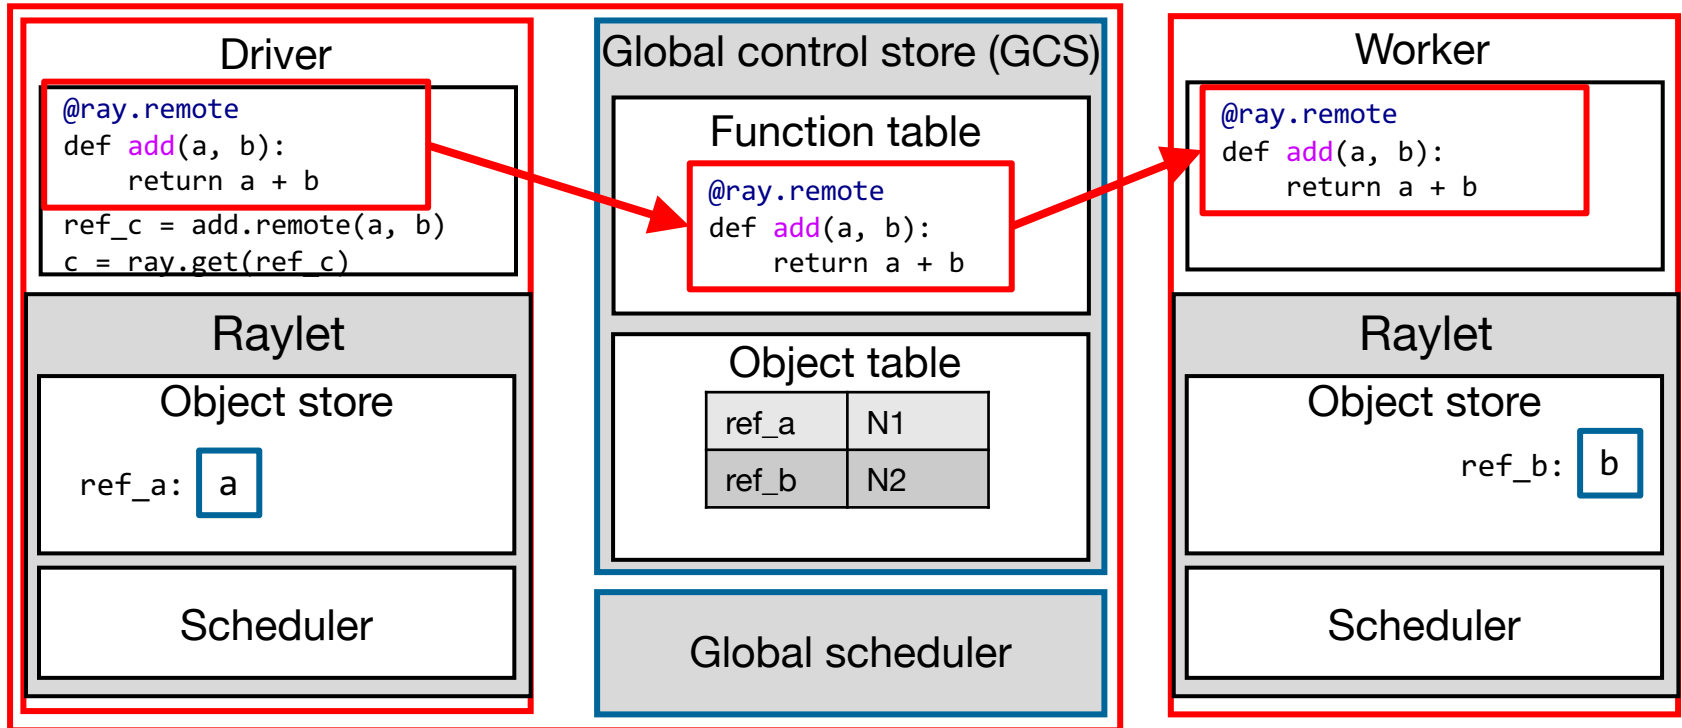
```
@ray.remote
```

```
def add(a, b):  
    return a + b
```

```
ref_c = add.remote(a, b)  
c = ray.get(ref_c)
```

Head node (N1)

Worker node (N2)



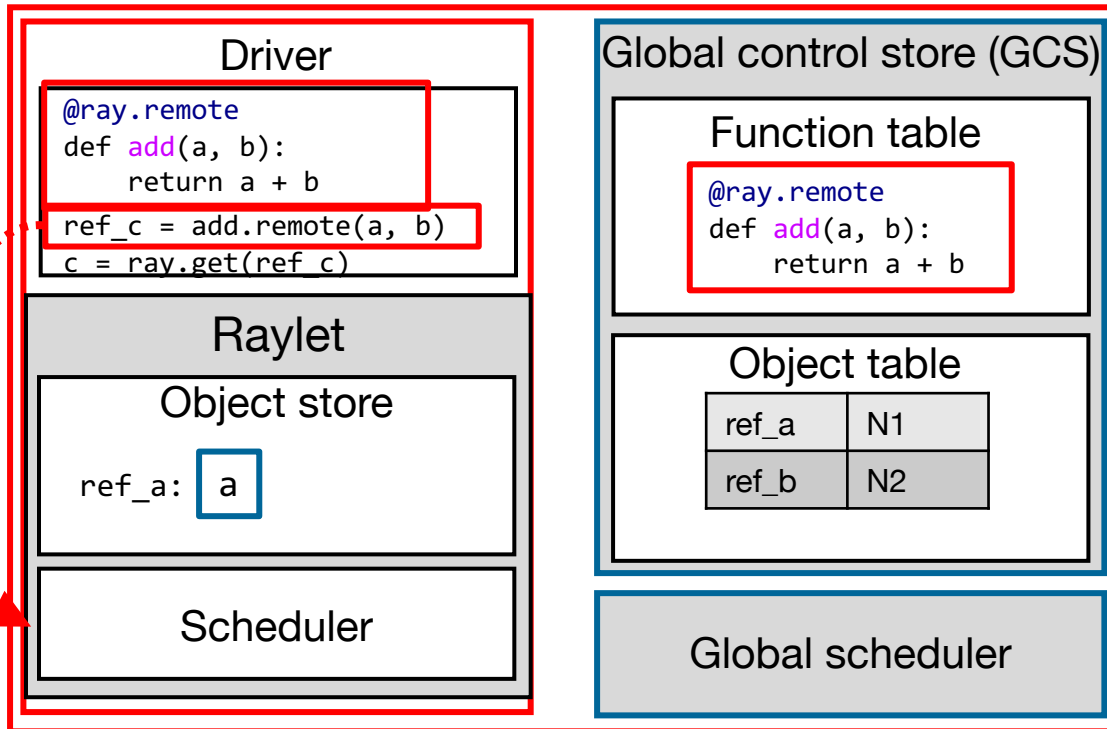
Cluster of machines

Step 0: Ray automatically registers each initialized remote function and distributed it to every worker in the cluster.


Executing a task remotely

```
@ray.remote  
def add(a, b):  
    return a + b  
ref_c = add.remote(a, b)  
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



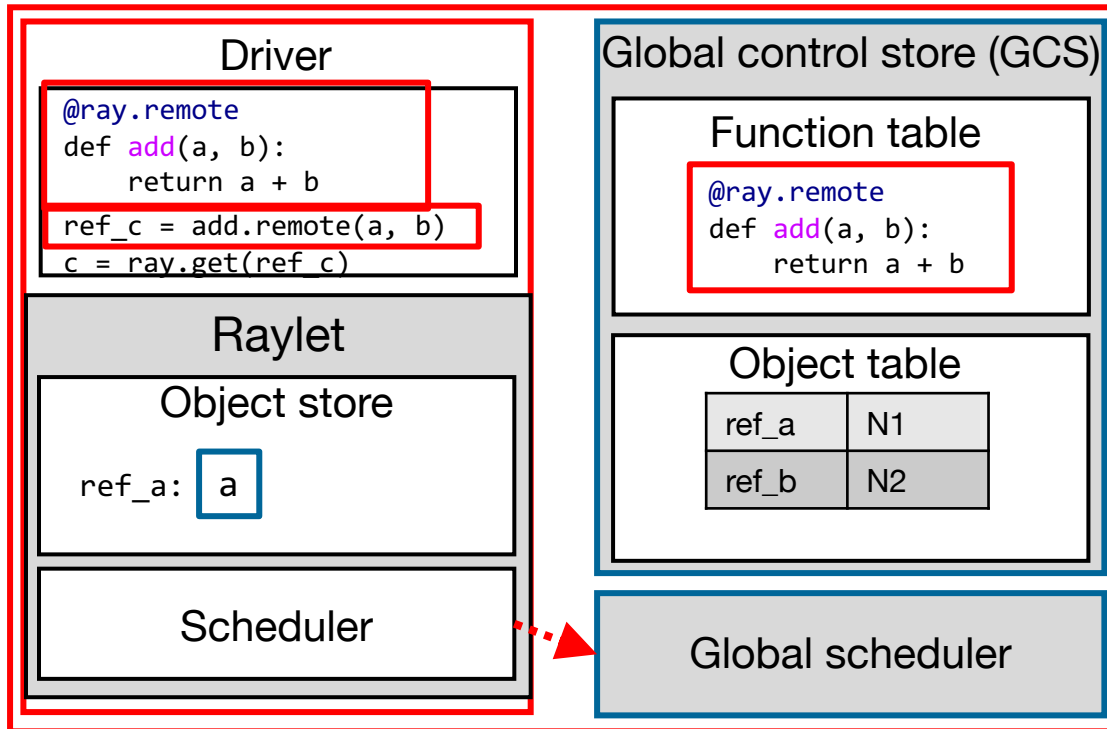
Cluster of machines

Step 1: Driver contacts N1's local scheduler to find out the ownership of object b (which node holds b).

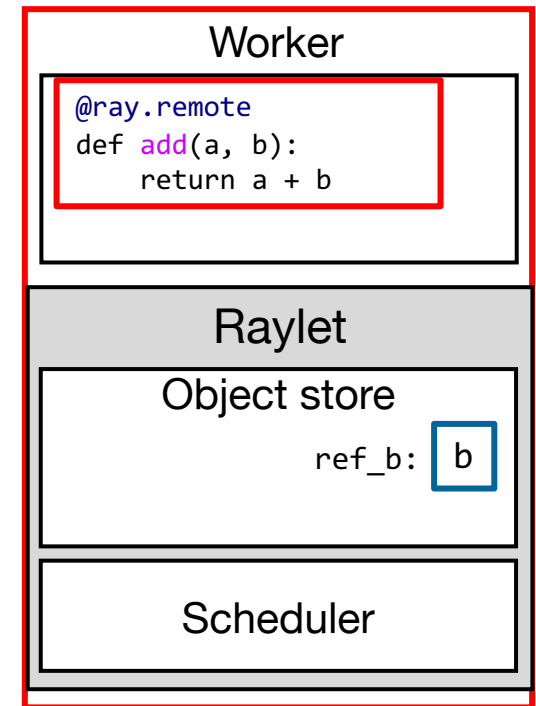
Executing a task remotely

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



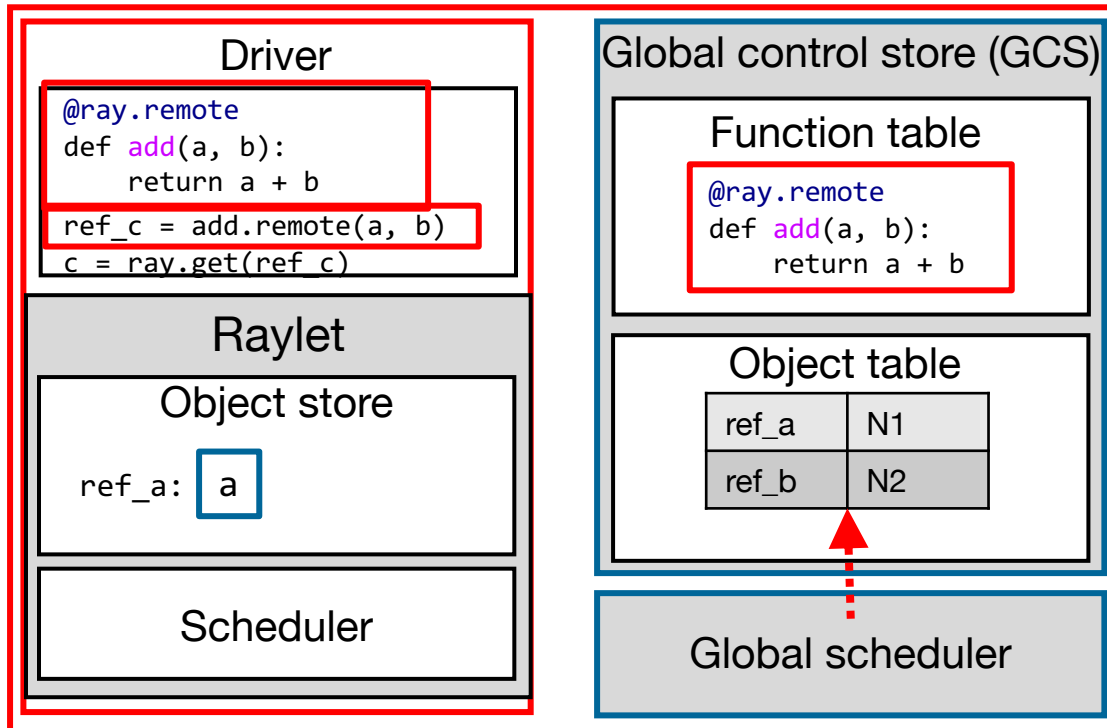
Cluster of machines

Step 2: N1's local scheduler (located on N1) contacts global scheduler.

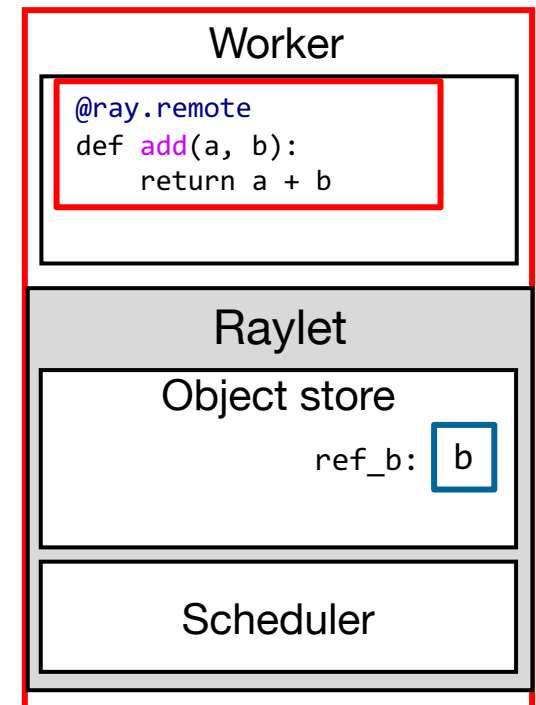
Executing a task remotely

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



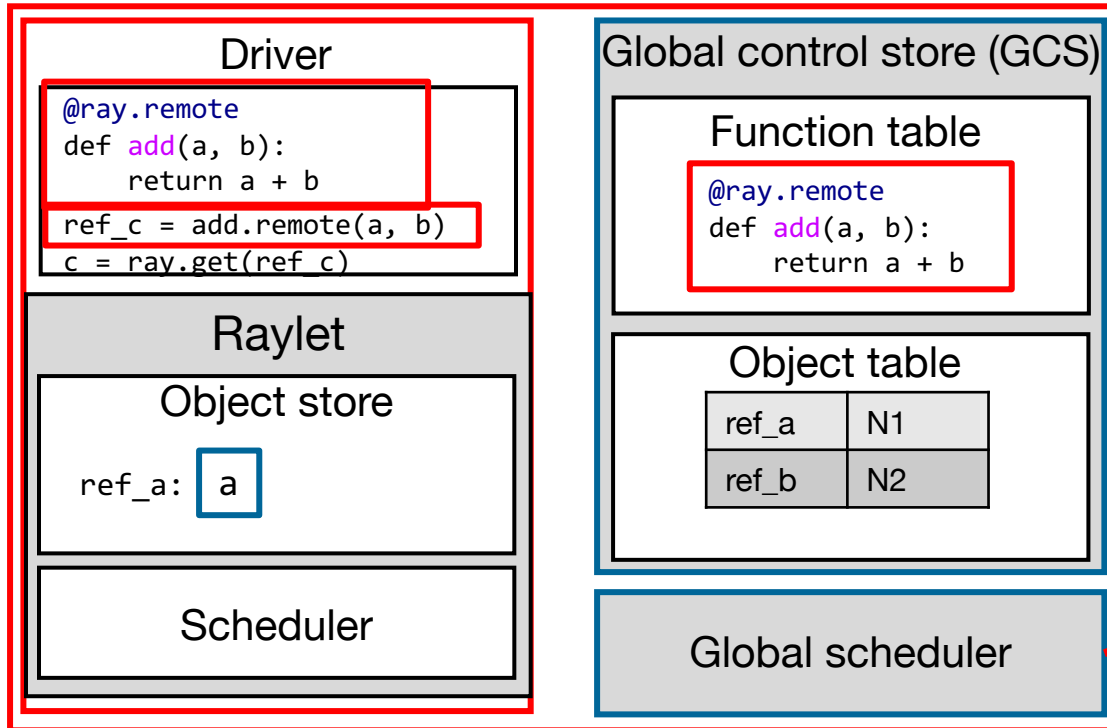
Cluster of machines

Step 3: Global scheduler performs an object table lookup and finds N2 holds b.

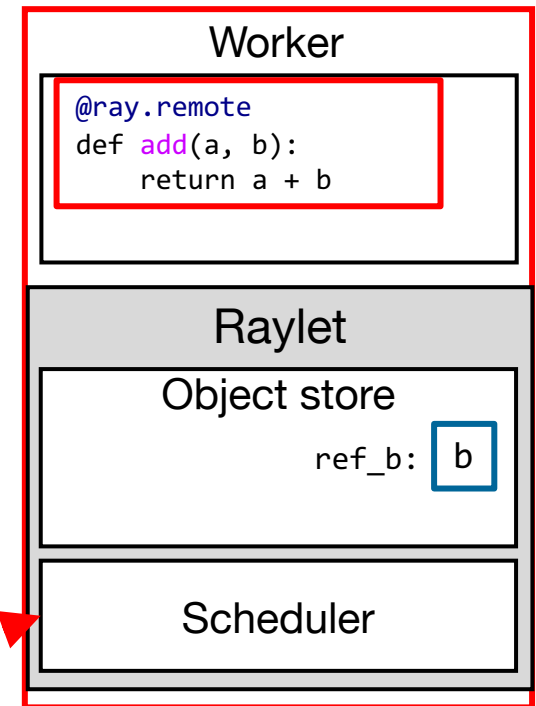
Executing a task remotely

```
@ray.remote  
def add(a, b):  
    return a + b  
ref_c = add.remote(a, b)  
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



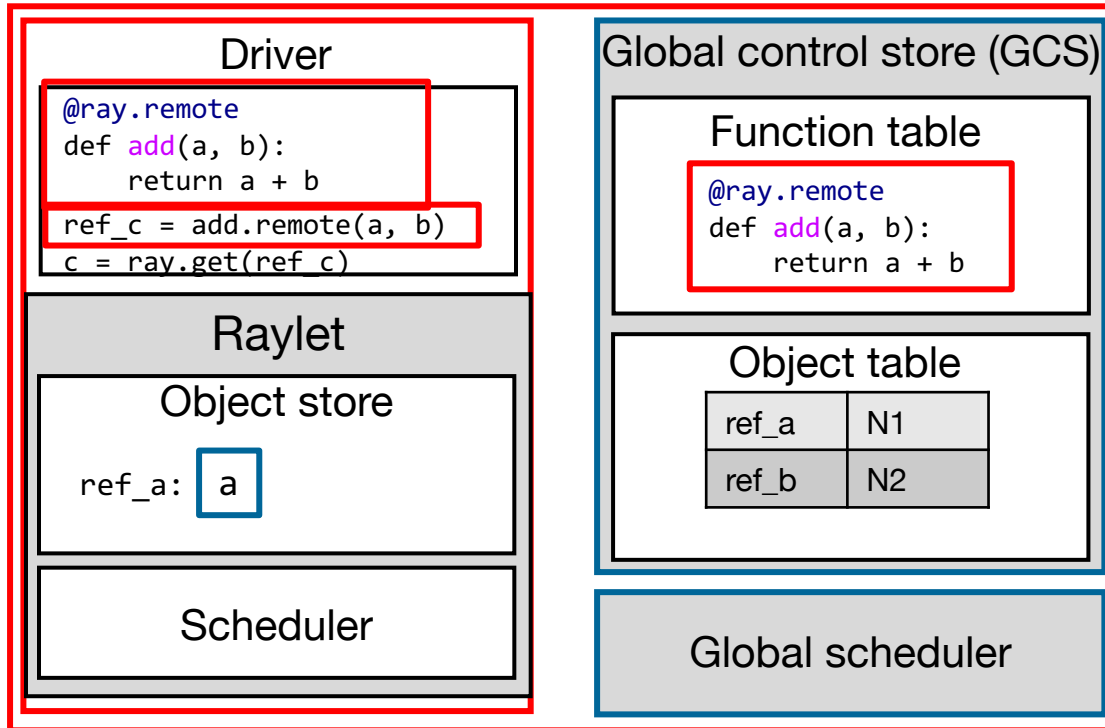
Cluster of machines

Step 4: Global scheduler decides to schedule the task on N2.

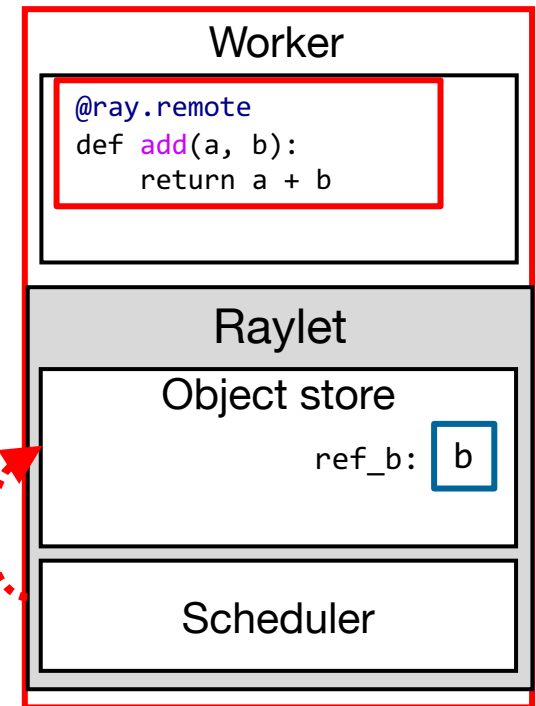
Executing a task remotely

```
@ray.remote  
def add(a, b):  
    return a + b  
ref_c = add.remote(a, b)  
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



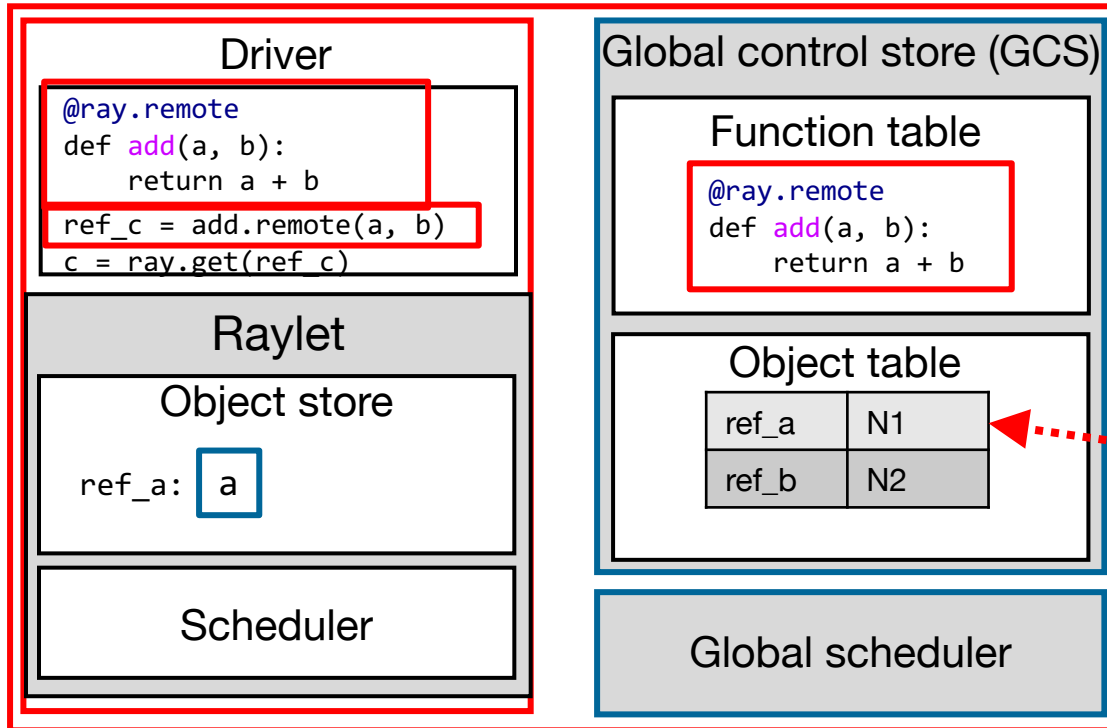
Cluster of machines

Step 5: N2's local scheduler checks whether the local object store contains `add(a, b)`'s arguments.

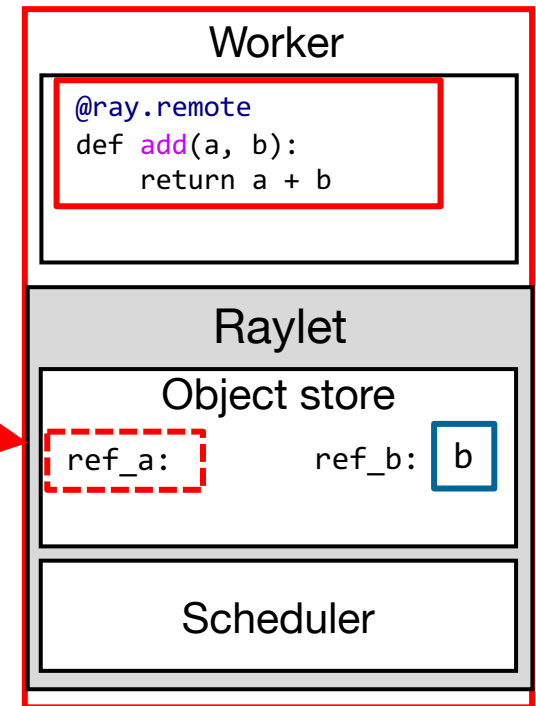
Executing a task remotely

```
@ray.remote  
def add(a, b):  
    return a + b  
ref_c = add.remote(a, b)  
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



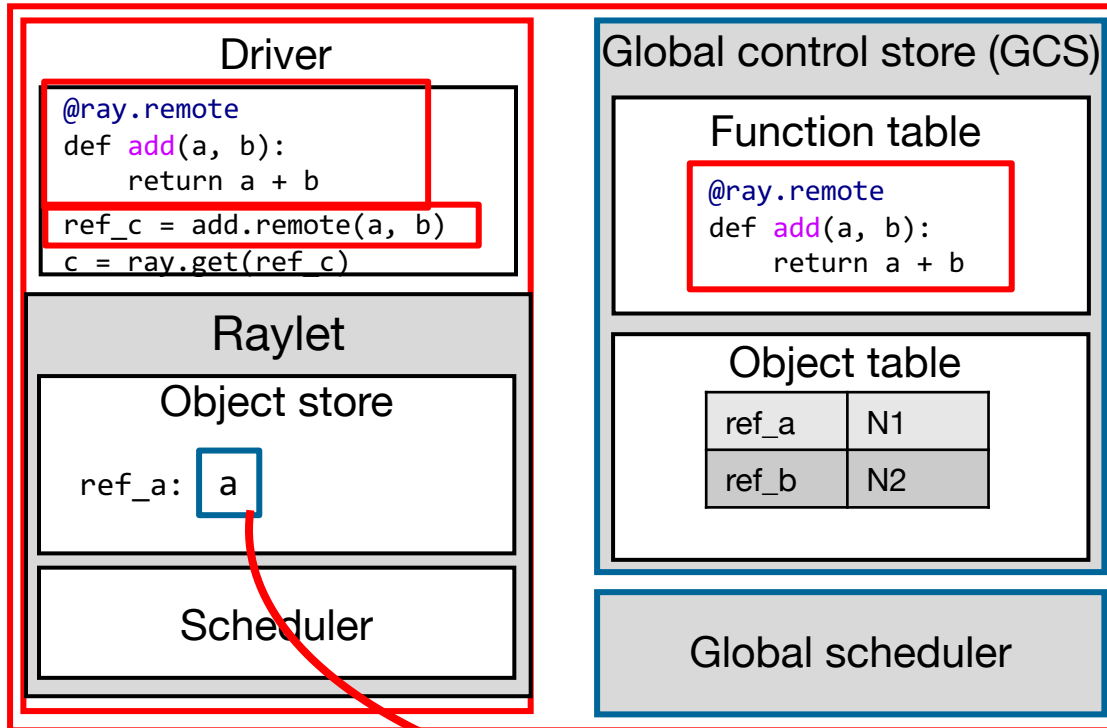
Cluster of machines

Step 6: N2 looks up a's location in the GCS.

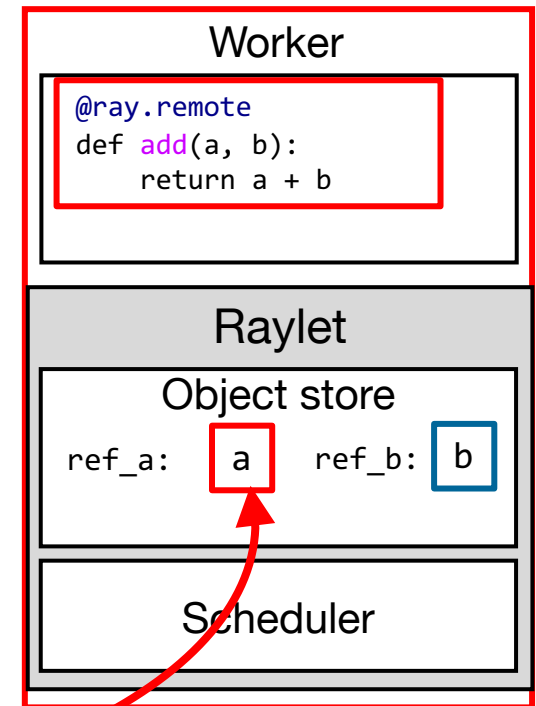
Executing a task remotely

```
@ray.remote  
def add(a, b):  
    return a + b  
→ ref_c = add.remote(a, b)  
   c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



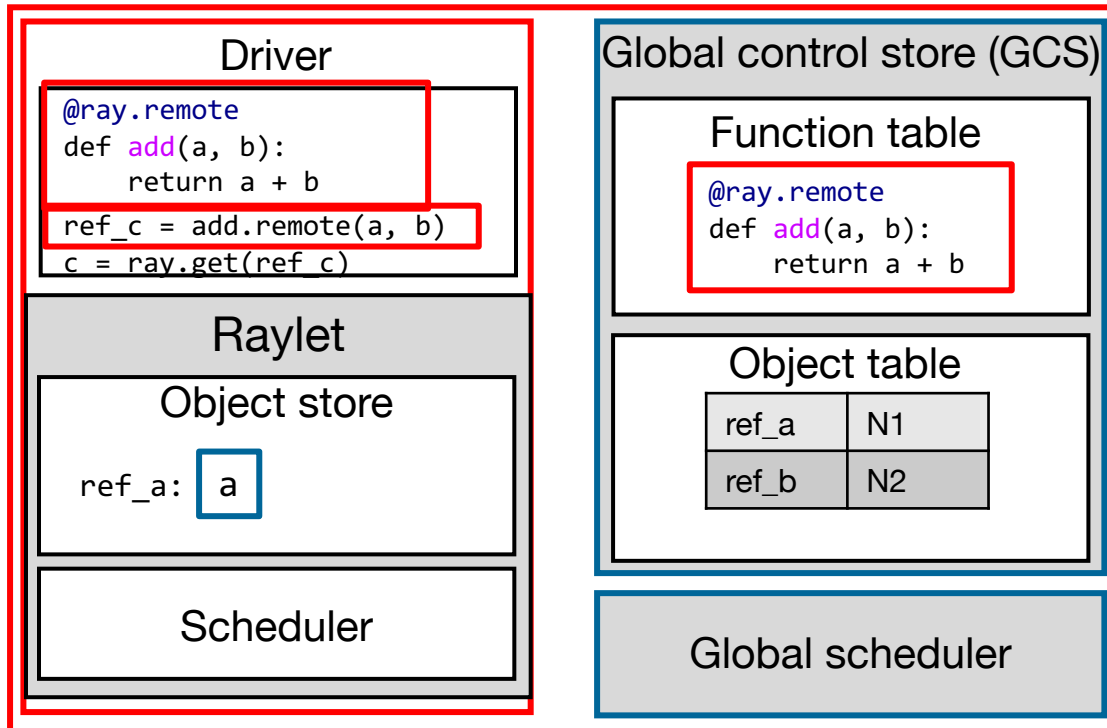
Cluster of machines

Step 7: Learning that N1 holds a, N2 fetches object a from N1's object store and replicates it locally.

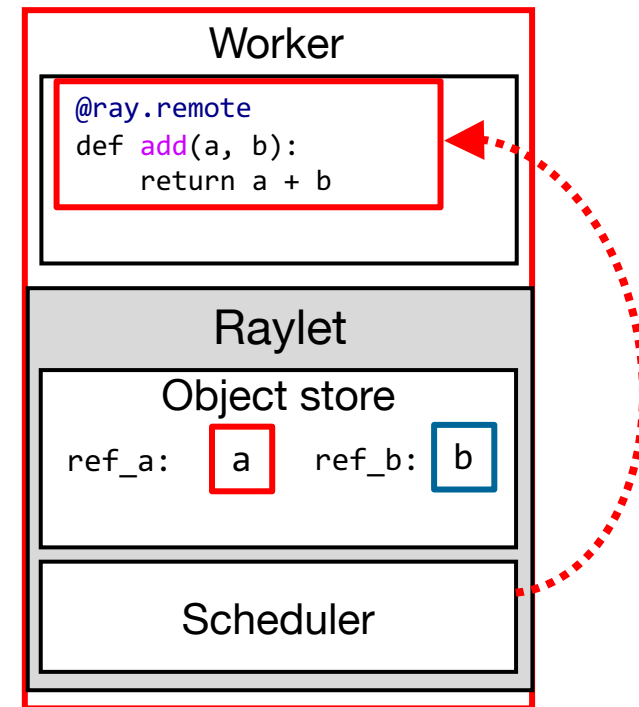
Executing a task remotely

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



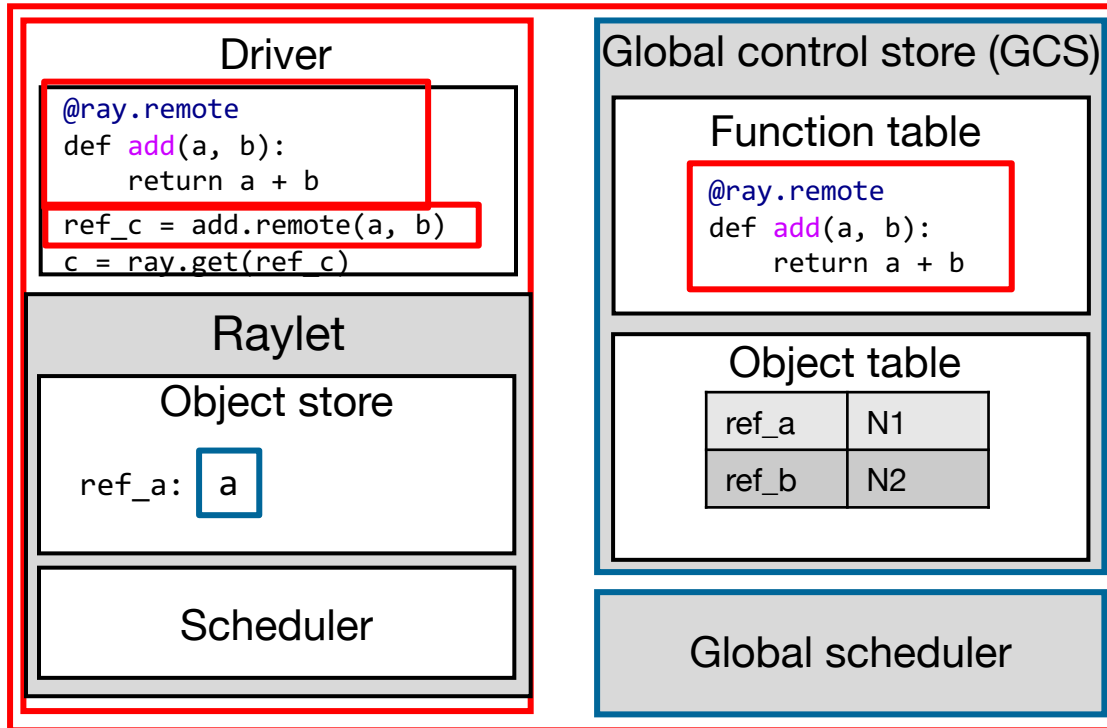
Cluster of machines

Step 8: N2's local scheduler invokes the task function `add()` at N2's local worker.

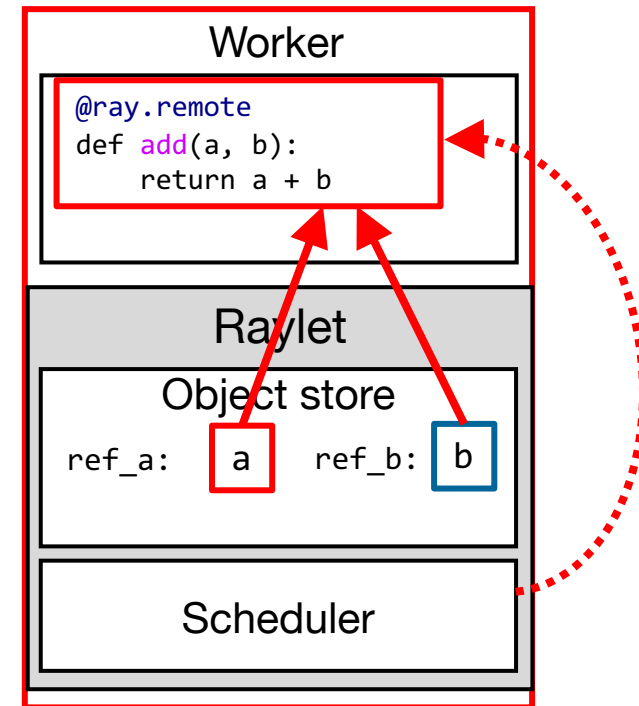
Executing a task remotely

```
@ray.remote  
def add(a, b):  
    return a + b  
→ ref_c = add.remote(a, b)  
   c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



Cluster of machines

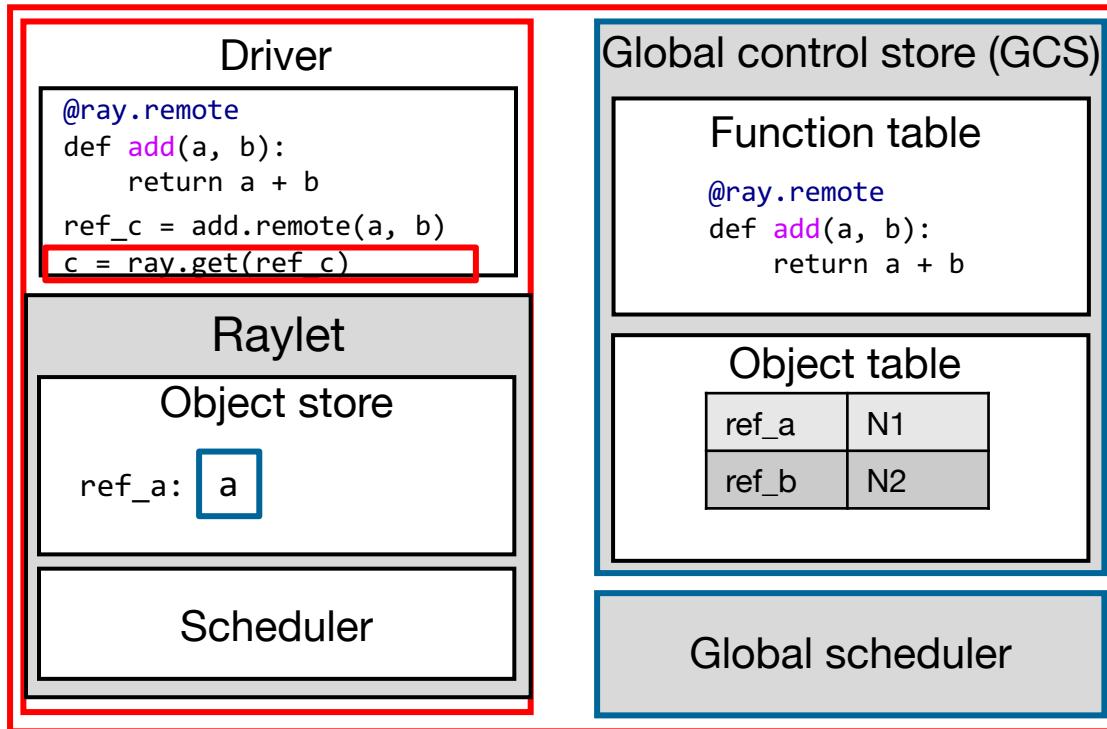
Step 9: N2's worker process executes the function code by accessing locally stored object a and b.

Getting the result of a remote task w/ `ray.get()`

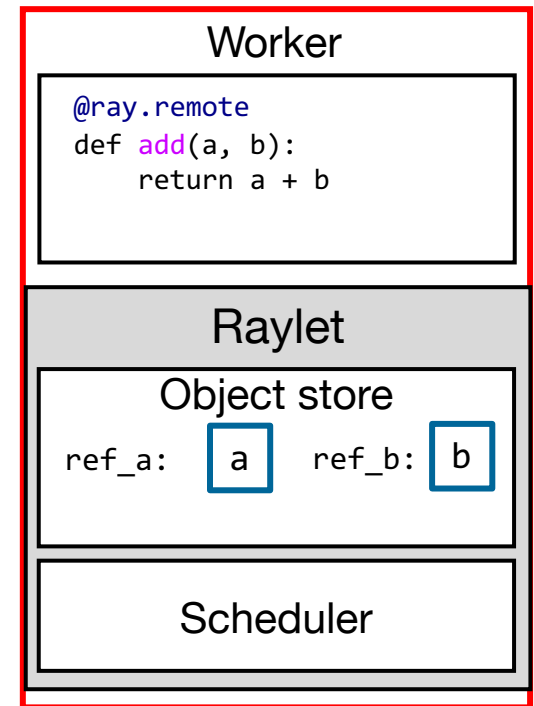
Getting the result of a remote task w/ `ray.get()`

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



Cluster of machines

Now, executing `ray.get(c)`...

Getting the result of a remote task w/ `ray.get()`

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)

Worker node (N2)

Driver

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Raylet

Object store

ref_a: a

Scheduler

Global control store (GCS)

Function table

```
@ray.remote
def add(a, b):
    return a + b
```

Object table

ref_a	N1
ref_b	N2

Global scheduler

Worker

```
@ray.remote
def add(a, b):
    return a + b
```

Raylet

Object store

ref_a: a ref_b: b

Scheduler

Cluster of machines

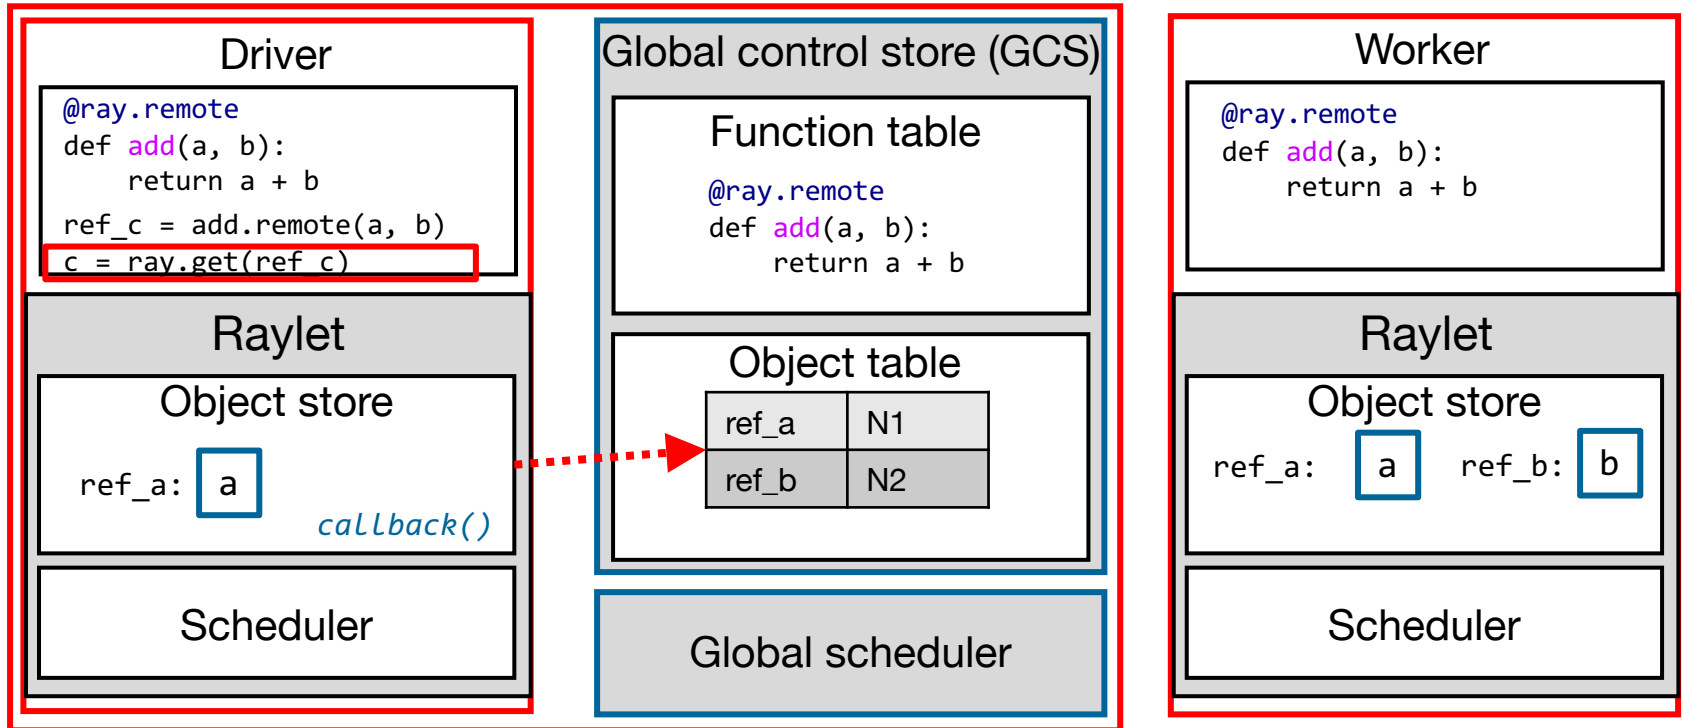
Step 1: Driver checks local object store for object c using the future ref of c.

Getting the result of a remote task w/ `ray.get()`

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)

Worker node (N2)



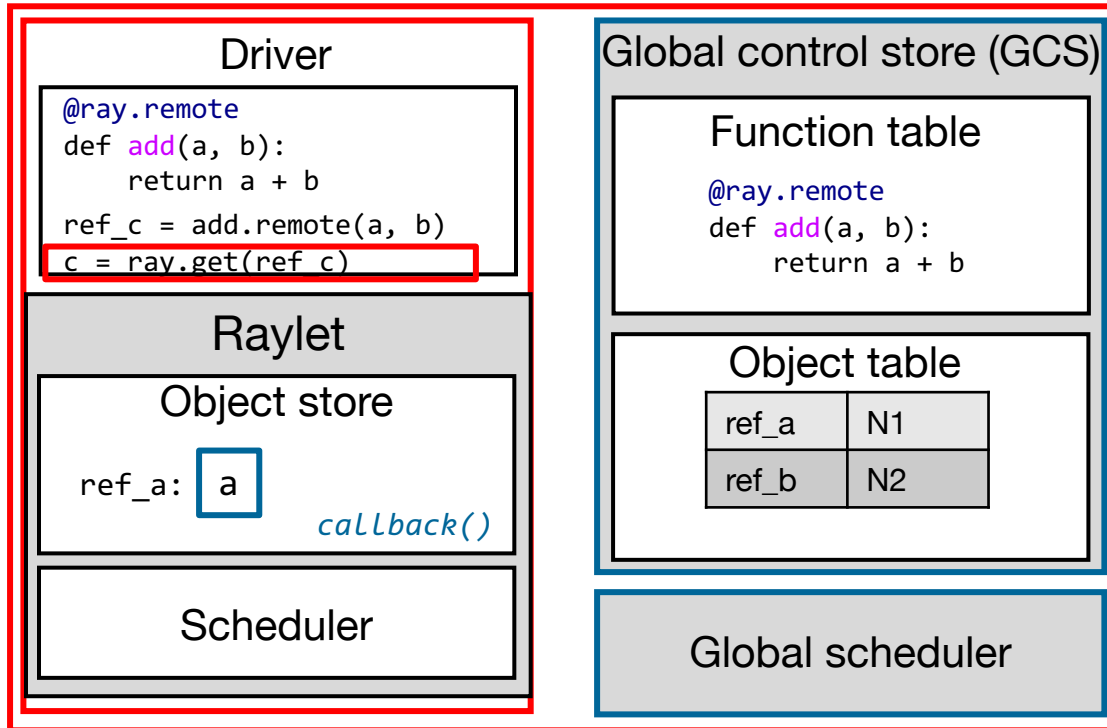
Cluster of machines

Step 2: N1's local object store looks up `c`'s location in GCS. GCS does not have an entry for `c` yet. Therefore, N1 registers a callback with GCS' object table to be triggered when `c`'s entry is created.

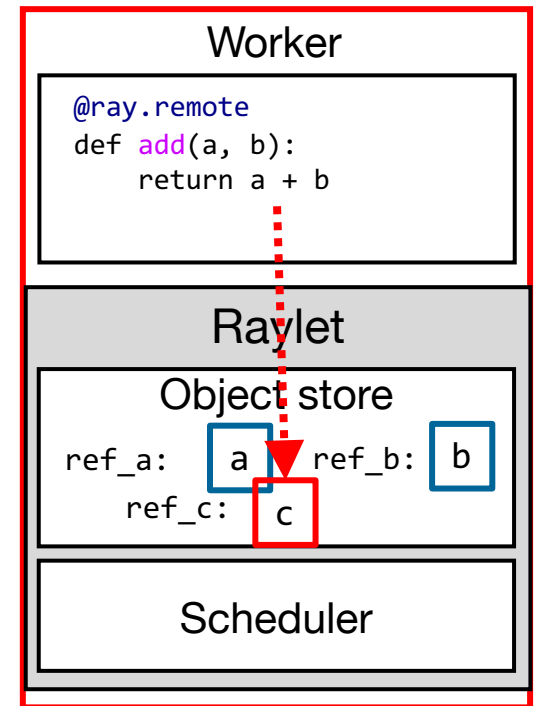
Getting the result of a remote task w/ `ray.get()`

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



Cluster of machines

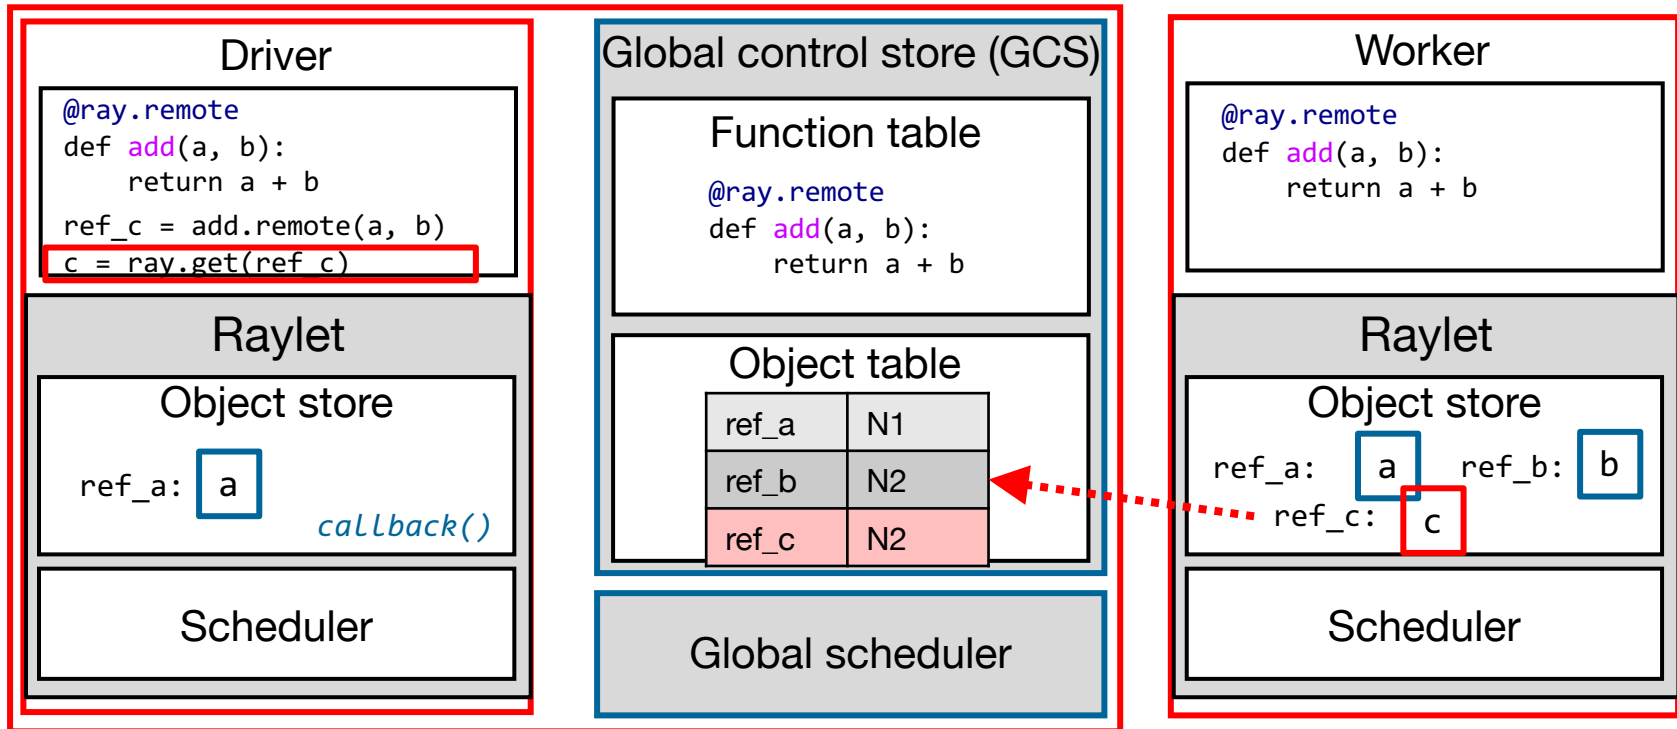
Step 3: N2's worker completes the execution of `add()` and stores the result `c` to the local object store.

Getting the result of a remote task w/ `ray.get()`

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)

Worker node (N2)



Cluster of machines

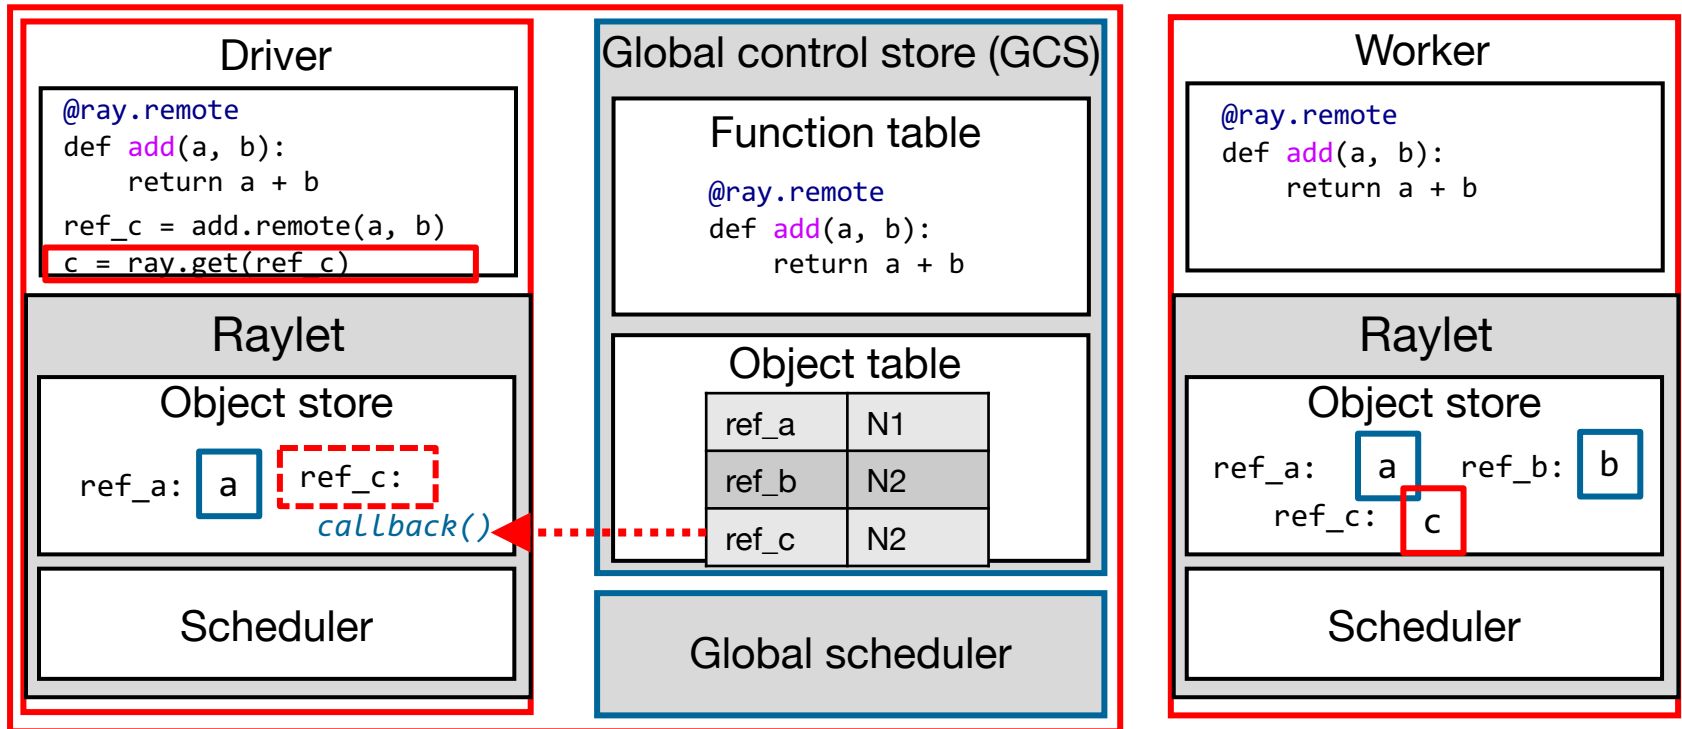
Step 4: N2's local object store in turn adds `c`'s entry to GCS.

Getting the result of a remote task w/ `ray.get()`

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)

Worker node (N2)



Cluster of machines


Step 5: GCS triggers the previously registered callback to N1's object store with c's entry.

Getting the result of a remote task w/ `ray.get()`

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)

Worker node (N2)

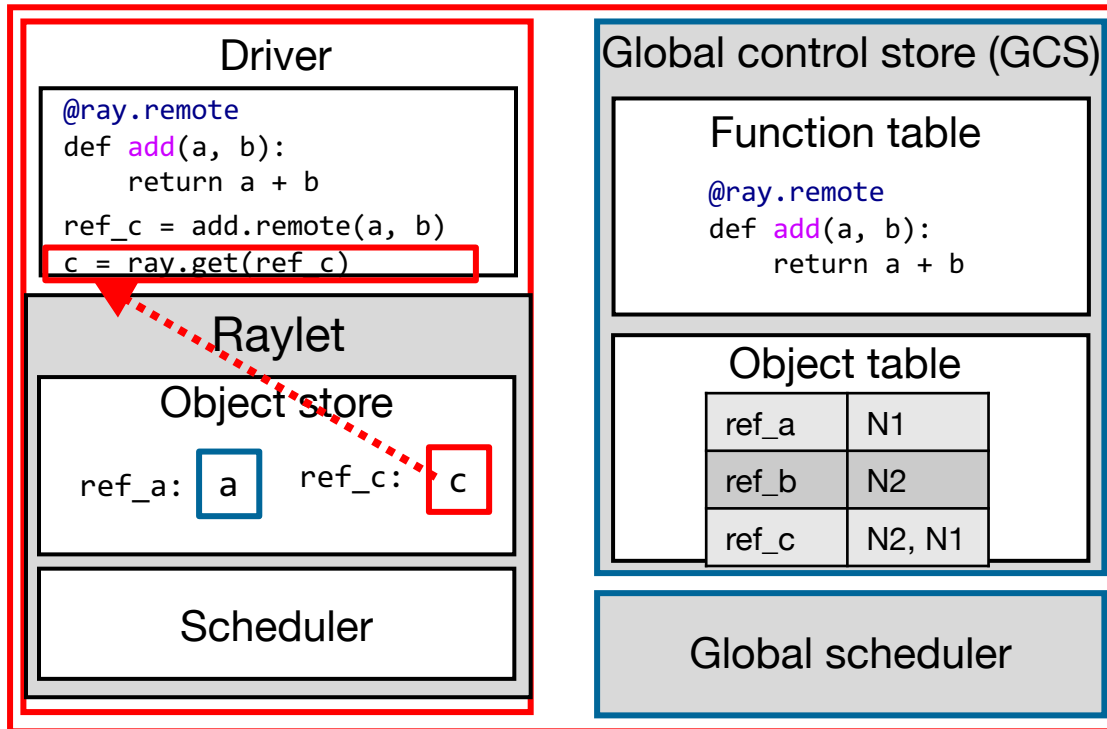


Step 6: N1 fetches `c` from N2's object store and replicates it in N1's local object store.

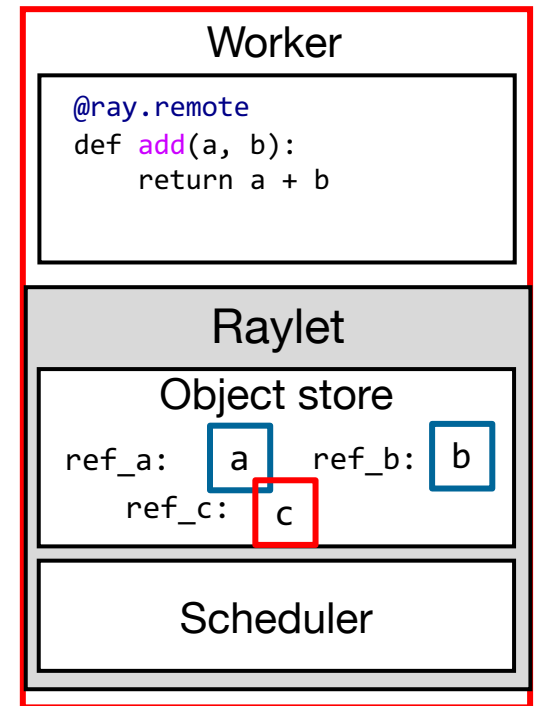
Getting the result of a remote task w/ `ray.get()`

```
@ray.remote
def add(a, b):
    return a + b
ref_c = add.remote(a, b)
c = ray.get(ref_c)
```

Head node (N1)



Worker node (N2)



Cluster of machines

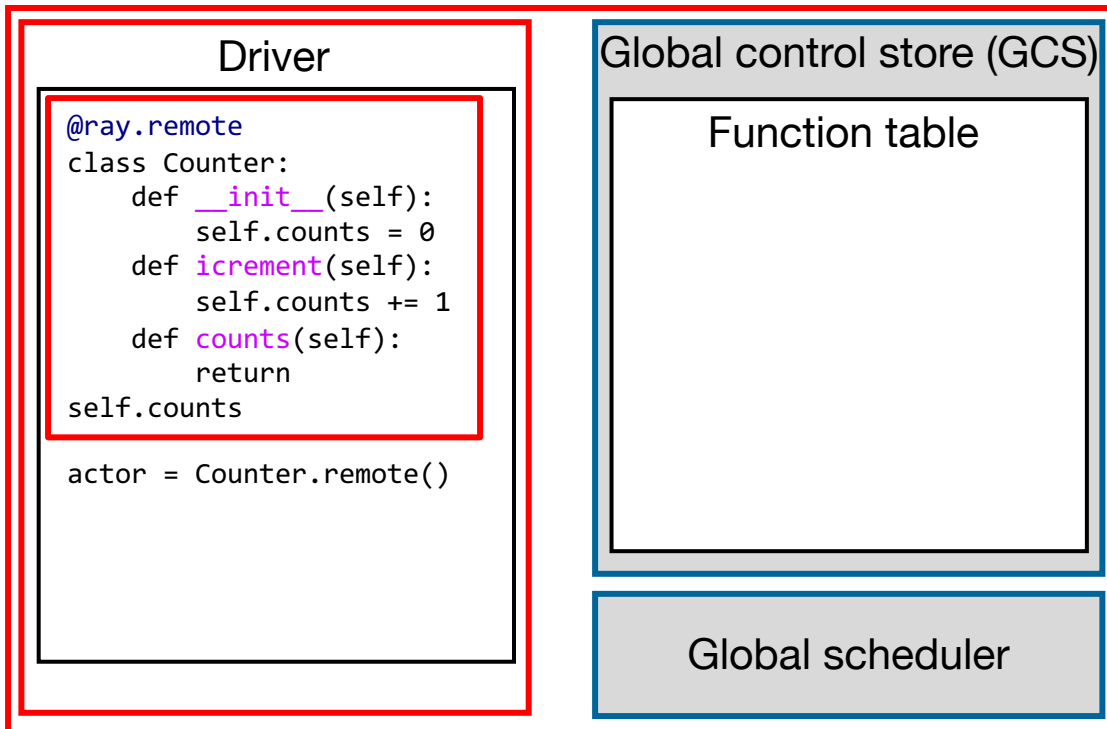
Step 7: N1's object store returns `c` to `ray.get()`.

Actor management


Actor creation

```
@ray.remote
class Counter:
    def __init__(self):
        self.counts = 0
    def increment(self):
        self.counts += 1
    def counts(self):
        return self.counts
```

Head node (N1)



Worker node (N2)



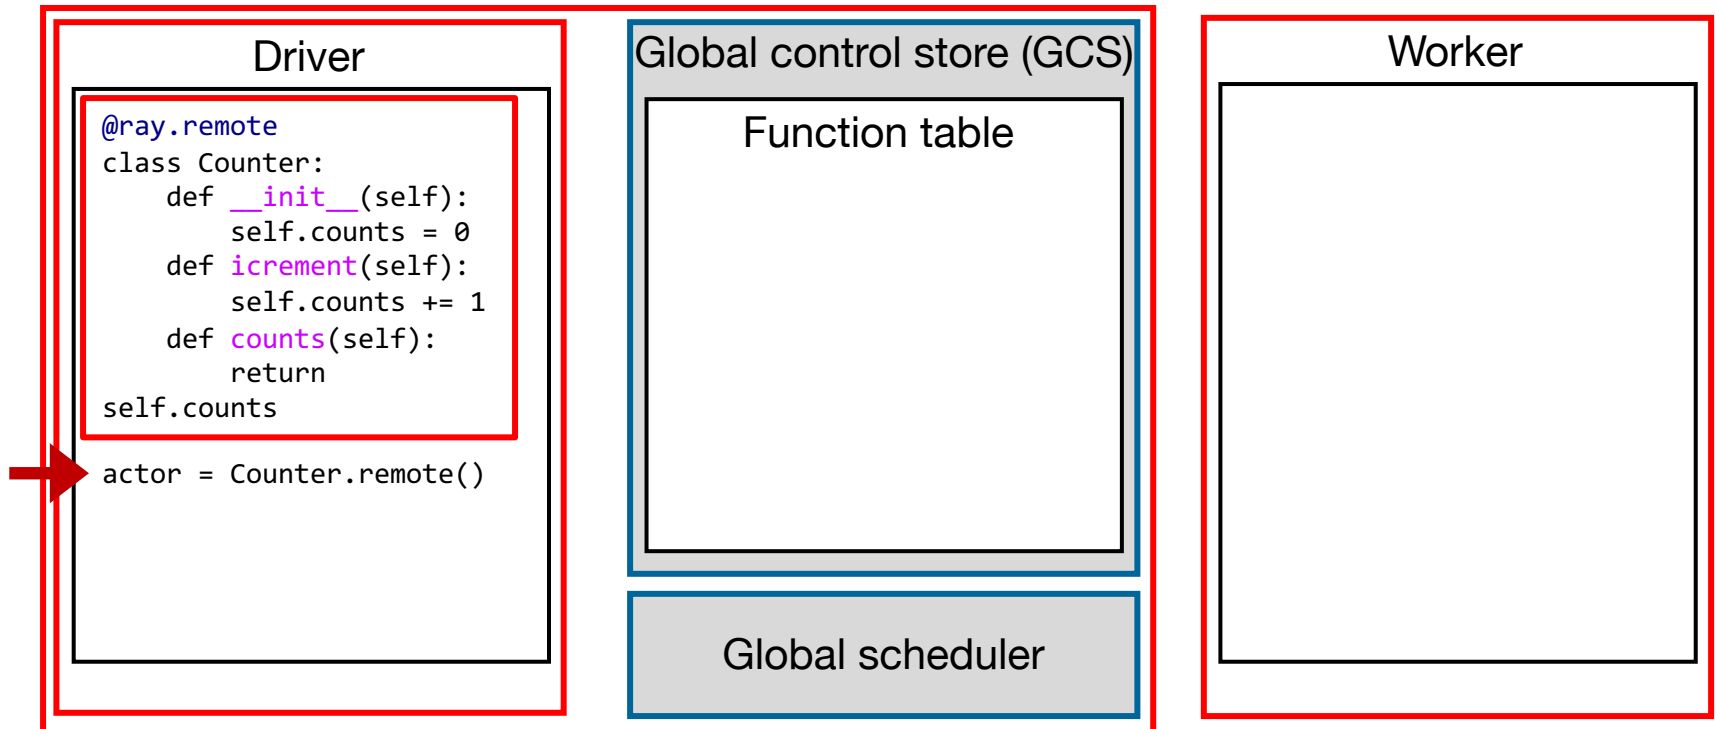
Cluster of machines

Actor creation

```
@ray.remote
class Counter:
    def __init__(self):
        self.counts = 0
    def increment(self):
        self.counts += 1
    def counts(self):
        return self.counts
```

Head node (N1)

Worker node (N2)

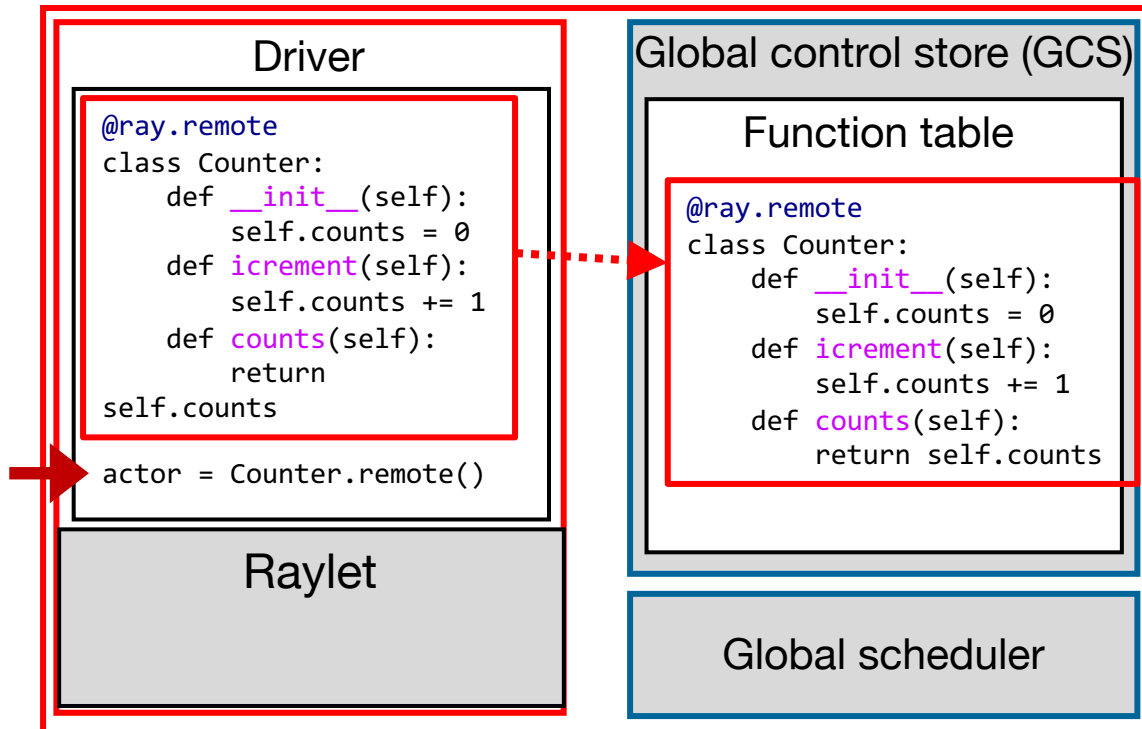


Cluster of machines


Actor creation

```
@ray.remote
class Counter:
    def __init__(self):
        self.counts = 0
    def increment(self):
        self.counts += 1
    def counts(self):
        return self.counts
```

Head node (N1)



Worker node (N2)



Cluster of machines

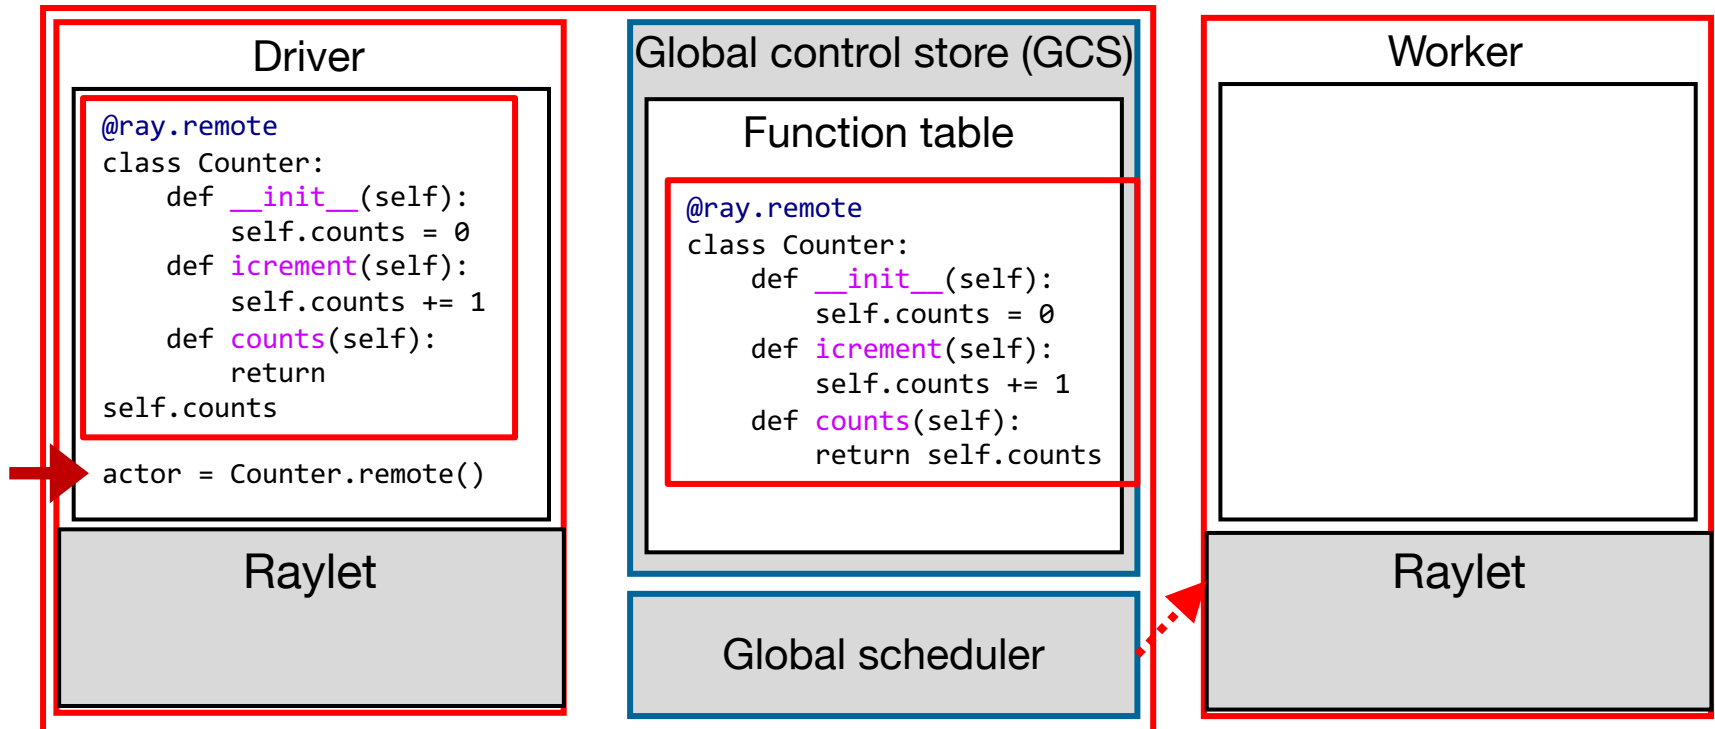
Step 1: Driver registers the actor with GCS.

Actor creation

```
@ray.remote
class Counter:
    def __init__(self):
        self.counts = 0
    def increment(self):
        self.counts += 1
    def counts(self):
        return self.counts
```

Head node (N1)

Worker node (N2)



Cluster of machines

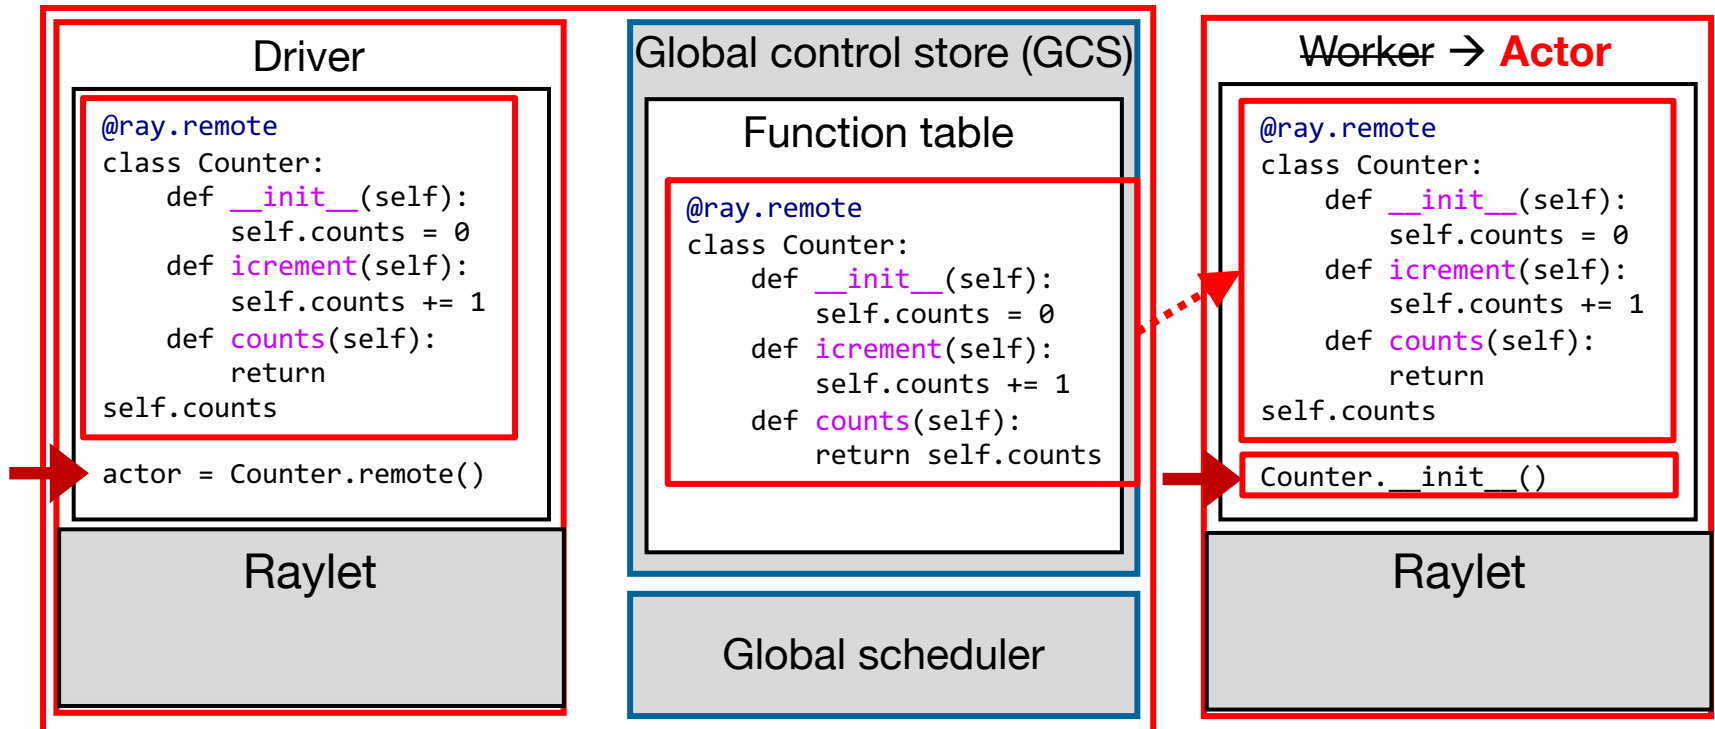
Step 2: Global scheduler selects a worker's raylet (N2), enqueue the actor creation request, and waits for the raylet to grant a resource lease.

Actor creation

```
@ray.remote
class Counter:
    def __init__(self):
        self.counts = 0
    def increment(self):
        self.counts += 1
    def counts(self):
        return self.counts
```

Head node (N1)

Worker node (N2)



Cluster of machines

Step 3: Once resource is granted on N2, GCS schedules the actor creation task on N2.
N2 now is effectively an Actor.

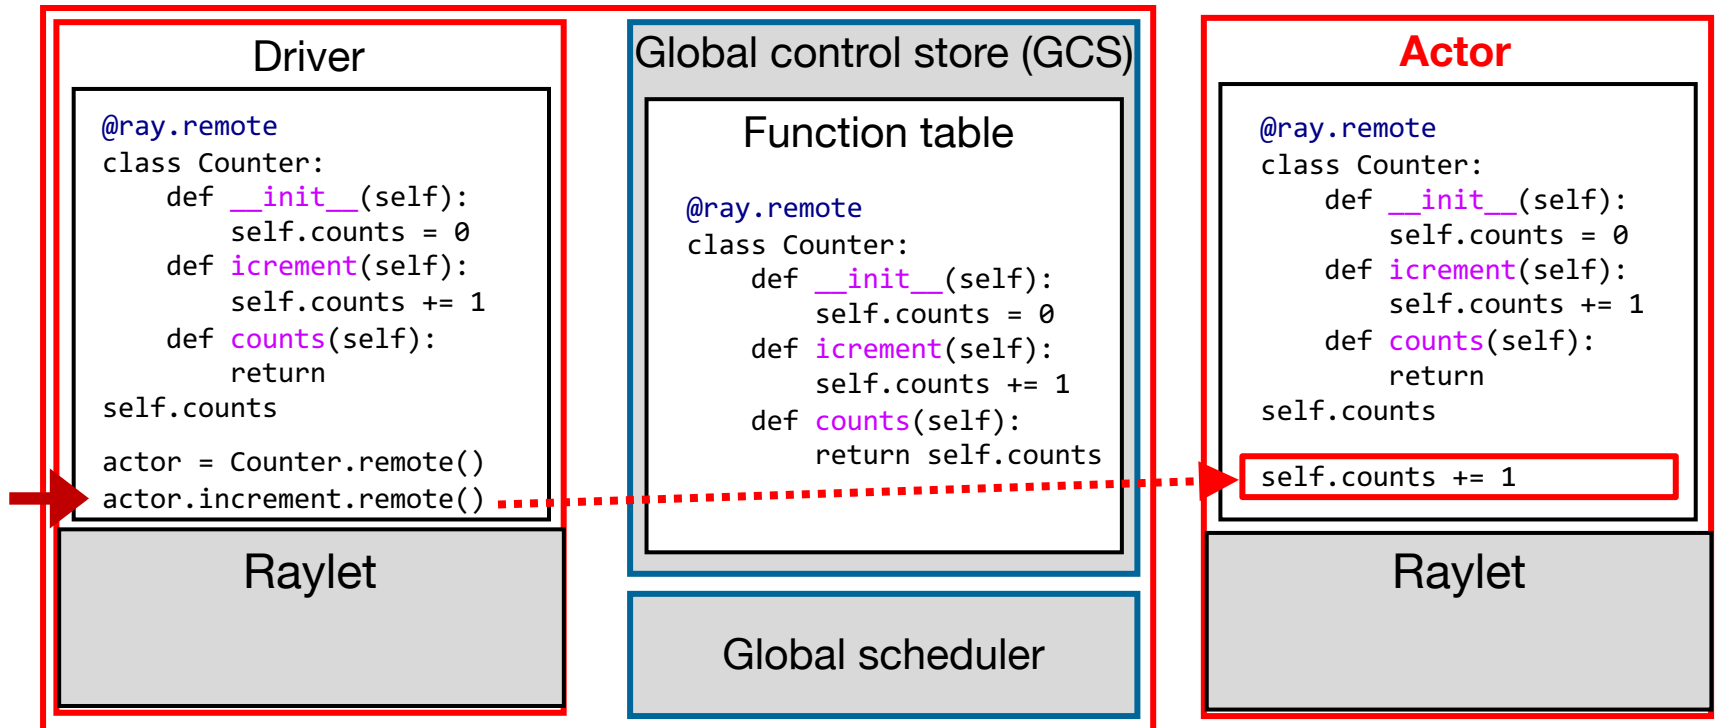
Actor task execution

Actor task execution

`actor.increment.remote()`

Head node (N1)

Worker node (N2)



Cluster of machines

Actor tasks are sent via remote function calls to the Actor process (N2).

Demo ...