

Cloud Computing

DS 5110/CS 5501: Big Data Systems

Spring 2024

Lecture 8a

Yue Cheng



Some material taken/derived from:

- Wisconsin CS 320 by Tyler Caraza-Harter.

@ 2024 released for use under a [CC BY-SA](https://creativecommons.org/licenses/by-sa/4.0/) license.

Learning objectives

- Understand basic cloud pricing model
- Know IaaS and PaaS and their differences
 - The cloud offering that we've been using through this semester is IaaS (EC2).
 - PaaS cloud offerings are similar to the open-source data systems that we have been learning this semester.
- Get to know some basic ideas behind containerization and container orchestration

Background

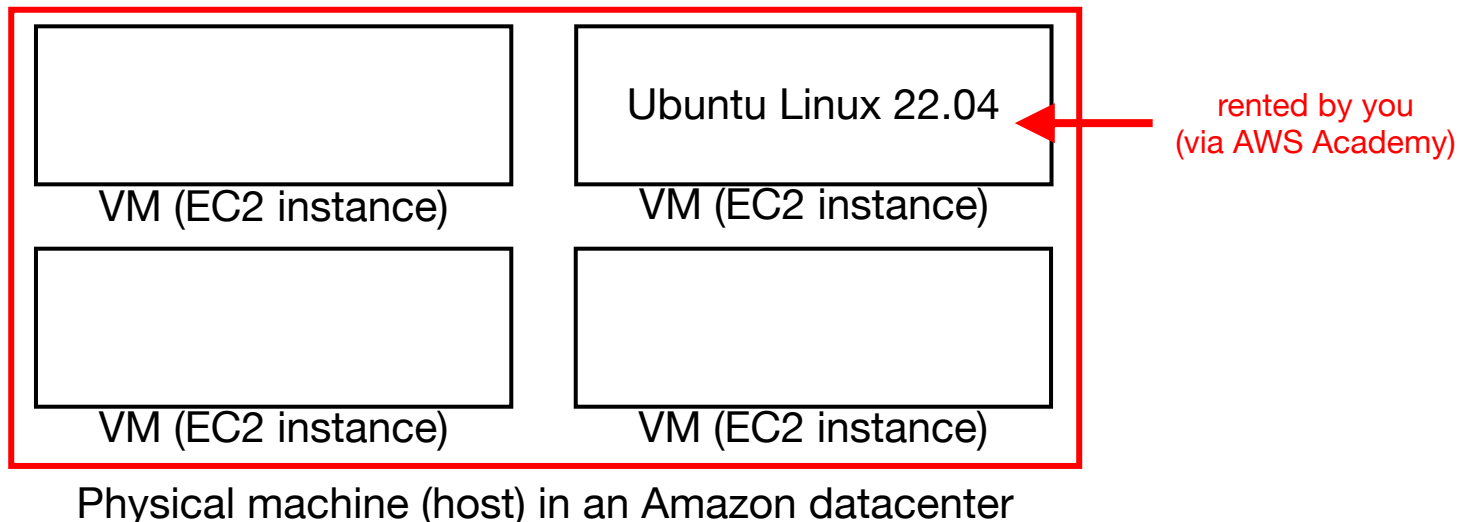
The beginning

“Sometimes you need a lot of processing power; and sometimes you need just a little. Sometimes you need a lot, but you only need it for a limited amount of time.”

-- Jeff Barr (https://aws.amazon.com/blogs/aws/amazon_ec2_beta/)

Amazon Web Services (AWS)

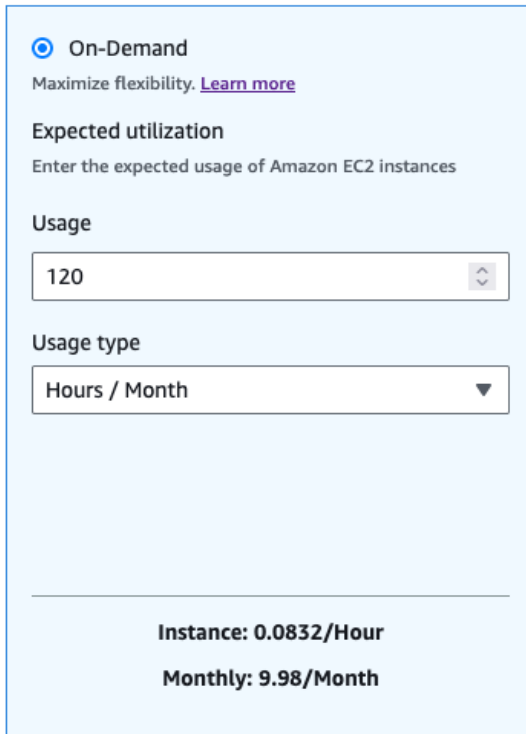
- Elastic Computing Cloud (EC2), rented VMs, launched in 2006
- “Infrastructure as a Service” (IaaS): rent infrastructure (compute, storage, network) instead of owning the hardware yourself



VM hours

Pricing summary

t3.large | Family: t3 | 2vCPU | 8 GiB Memory



The screenshot shows the AWS Pricing Calculator interface for an Amazon EC2 On-Demand instance. The instance type is t3.large. The usage is set to 120 hours per month. The pricing summary at the bottom shows an instance cost of 0.0832 per hour and a monthly cost of 9.98.

☒ On-Demand
Maximize flexibility. [Learn more](#)

Expected utilization
Enter the expected usage of Amazon EC2 instances

Usage
120

Usage type
Hours / Month

Instance: 0.0832/Hour
Monthly: 9.98/Month

Amazon EC2 On-Demand instances cost (Monthly): 9.98

Amazon Elastic Block Store (EBS) total cost (Monthly): 1.28

AWS pricing calculator: <https://calculator.aws/#/>

Pricing comparison

- **one VM for a month**: about \$10
- about 120 hours a month (4*30)
- **120 VMs for an hour**: about \$10
- same computation + storage resources
- very different wait time

Be careful!!

- programmers previously optimized when things were **too slow**
- now we need to optimize when it is **too expensive**
- cost is not always obvious at the moment you're running a job (need to do "back of the envelope" estimates before you deploy the resources)


Other cloud services

- AWS now has > 200 services beyond EC2 (and growing)
- **IaaS** (Infrastructure as a Service)
 - EC2, other services that feel closer to raw hardware
 - Virtual disks, virtual network, some storage systems, etc.
 - Cheap + flexible – you can deploy & run anything on it (Spark, Ray, etc.)
- **PaaS** (Platform as a Service)
 - Cloud providers has deployed systems on the infrastructure; you pay to use the deployed system
 - Databases, application framework/platforms, ML training/deployment systems
 - Less flexible, easier to use
 - Often more expensive (though not necessarily more than doing it yourself due to efficiencies available to cloud provider but not you)
- Line between IaaS and PaaS distinction is a bit subjective.

Lock-in


- Customers (tenants) worry: what if the cloud provider increases the price? If it's hard to move to a competing cloud, you're “**locked in**”
- PaaS: services are often unique, and it would be hard to move to a different cloud providers
- IaaS: services like VMs are more uniform – it would be easier to switch to a different cloud to find the cheapest place to rent VMs
- **Data**: cloud providers often make it free to bring data into the cloud (ingress) but expensive to take it out (egress)

Case study: Dropbox



- A data sync startup founded back in 2008
- Became popular so quickly
 - Peak number of users: 500+ Million
 - Overall amount of data stored: 500 PB
- Initially stored all data on public clouds (AWS)
- Seriously considered to move data out of AWS
- Cloud vendor lock in
 - **Enormous** egress costs
- Now still parts of its data services sitting on AWS

Major cloud providers today



<https://www.srgresearch.com/articles/q3-cloud-spending-up-over-11-billion-from-2021-despite-major-headwinds-google-increases-its-market-share>

Cloud economics and billing models

Tenants: Pay-as-you-go?



- (**Claimed**) pay-as-you-go pricing
 - Usage-based?
 - Most (compute) services charged per minute
 - Storage and network services charged per byte
 - No minimum or upfront fee

Tenants: Pay-as-you-go?

- **(Claimed)** pay-as-you-go pricing
 - Usage-based?
 - Most (compute) services charged per minute
 - Storage and network services charged per byte
 - No minimum or upfront fee

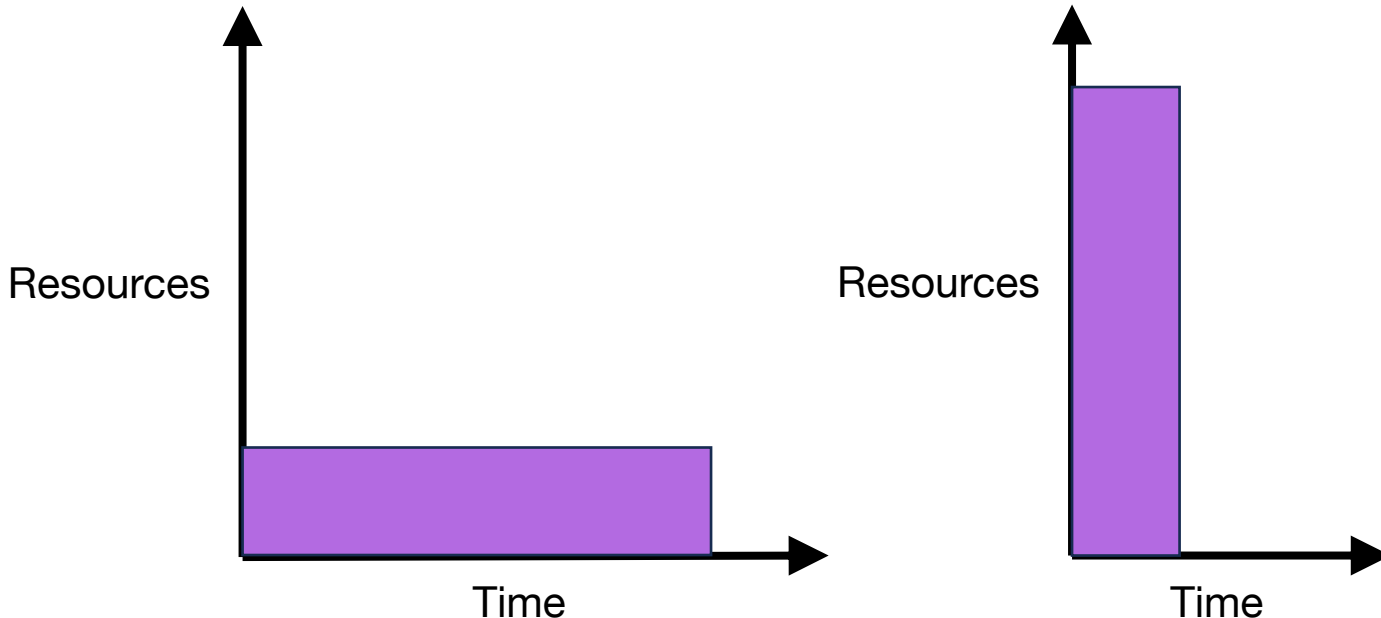
Q: Is the cloud pricing truly pay-as-you-go?

- **Problem:** How to perform strategic planning?



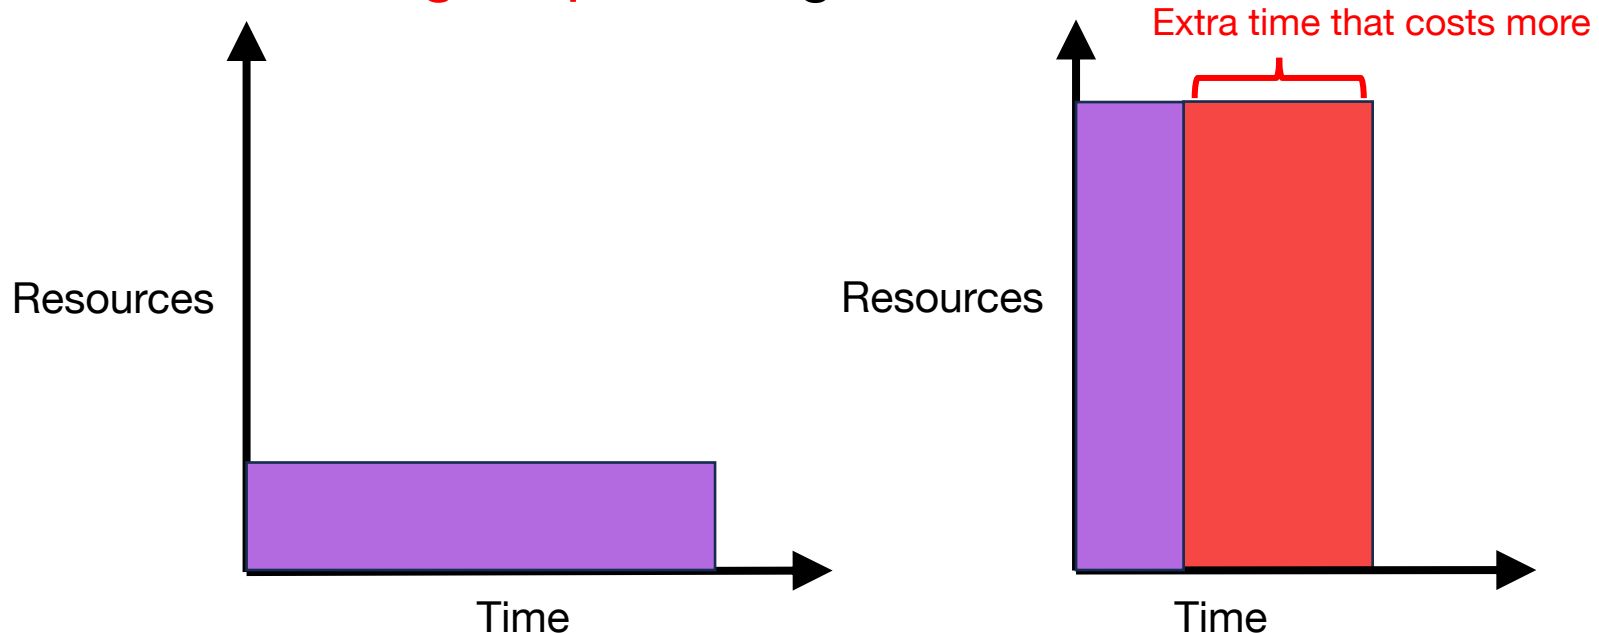
Tenants: Scalability gained?

- (**Ideally**) Linear scalability & perfect elasticity
 - Using 1000 servers for 1 hour costs the same as 1 server for 1000 hours
 - Same price to get a result faster




In practice, it really depends, case by case.
Likely the speedup of the computation is much lower than 1000X!

- **(In reality)** Scalability is sublinear and VM auto-scaling is slow.
 - Using 1000 servers for 1+N hour costs **N times** more than 1 server for 1000 hours
 - Often **higher price** to get a result faster



Providers: On-demand vs. spot instances



- How to create **incentives** for tenants?
 - Use less at peak time
 - Use more at low times
- Two VM deployment options
 - **On-demand instances**: Constant (high) price. Can generally get a VM. Won't be taken away from you arbitrarily. Used when capacity is needed at specific times.
 - **Spot instances**: Price varies throughout day. If you're not willing to pay enough, your computation waits for a cheaper price. VM might be interrupted ("preempted") once started. Excellent for once-a-day batch jobs.

Providers: Free tier, discounts at scale

AWS Lambda Pricing

Region: US East (N. Virginia) ↕

Architecture	Duration
x86 Price	
First 6 Billion GB-seconds / month	\$0.0000166667 for every GB-second
Next 9 Billion GB-seconds / month	\$0.000015 for every GB-second
Over 15 Billion GB-seconds / month	\$0.0000133334 for every GB-second

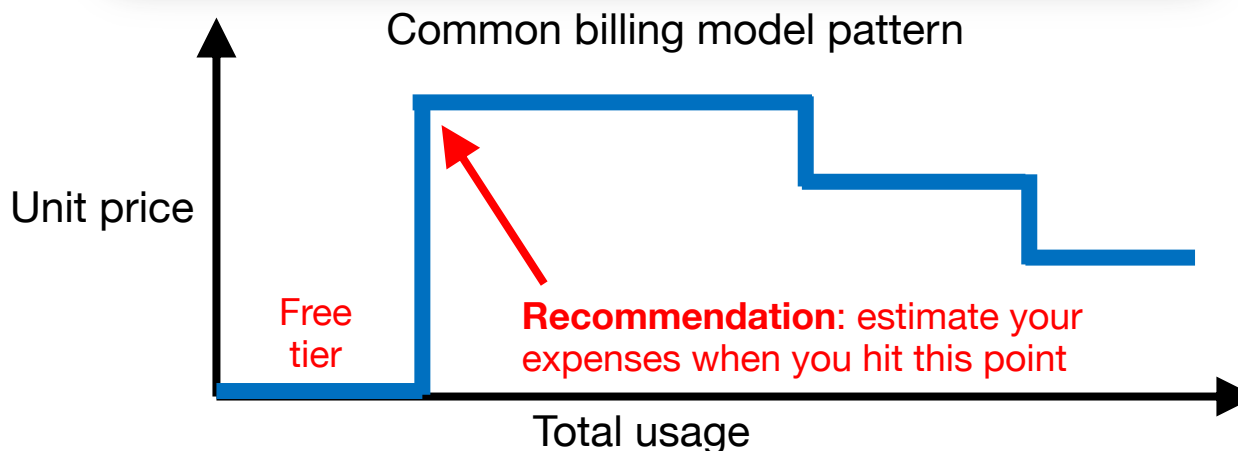
AWS Lambda example

“The AWS Lambda **free tier** includes one million free requests per month and 400,000 GB-seconds of compute time per month.”

(<https://aws.amazon.com/lambda/pricing/>)

“Duration is calculated from the time your code begins executing until it returns or otherwise terminates, **rounded up to the nearest 1 ms.**”

Recommendation: check if you have a large number of small ops getting rounded up



Virtualization and container orchestration

Virtualization: Providing isolation

- We don't want different applications running on the same hardware to interfere with each other – we want them to be **isolated**. Concerns:
 - Malicious programs, buggy programs, fairness
- Ways to interfere
 - **Directly**: Seeing/modifying data of another process
 - **Indirectly**: Inflicting bad performance on another process
- Some operating system isolation features with a long history:
 - Virtual memory: Can't see another process's data (namespace isolation)
 - Schedulers: Can't hog the whole CPU (performance isolation)

Problem: CPU and memory are not the only resources

Goal: Both namespace AND performance isolation for EVERY kind of resource

Linux features: cgroups and namespaces

- **cgroup** types (performance isolation)
 - cpu, memory, cpuacct, cpuset, freezer, net_cls, blkio, perf_event, net_prio, hugetlb, pids, rdma
- **namespace** types (namespace isolation)
 - network, mount, time, user, cgroup, IPC, PID, UTS
- Both cgroups and namespaces apply to sets of processes. Configuring all this by hand is **VERY** complicated.
- One reason Docker is popular: “docker run ...” starts a process using all these features, each with reasonable configurations.
- “**Containers**” definition: Set of processes using a combination of cgroup / namespace / other features.




Kubernetes (k8s)

8 letters

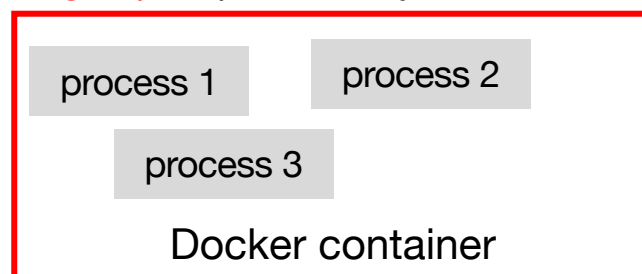
- cgroups and namespaces are very flexible: Docker's approach is just ONE way to use them to build containers



namespaces: mount, network, etc.
cgroups: cpu, memory, etc.




namespaces: mount, network, etc.
cgroups: cpu, memory, etc.



namespaces: network, etc.


namespaces: mount, etc.
cgroups: cpu, etc.



namespaces: network, etc.

namespaces: mount, etc.
cgroups: cpu, etc.

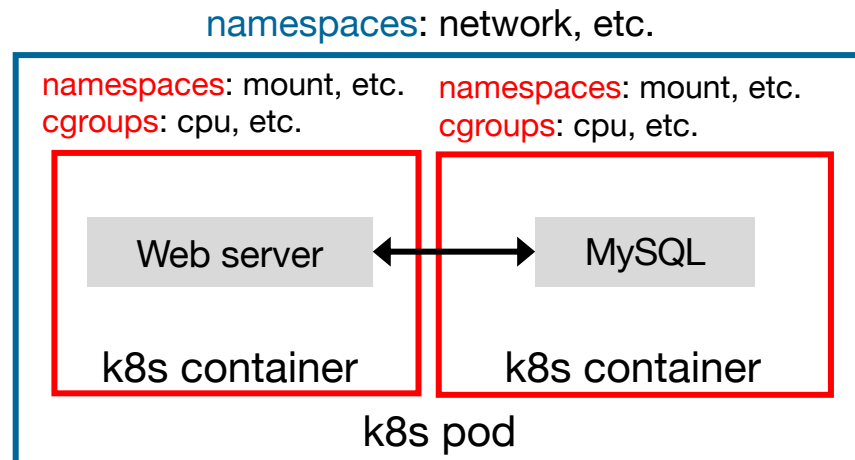
namespaces: mount, etc.
cgroups: cpu, etc.



Kubernetes (k8s)

8 letters

- **Motivation:** We often want to deploy multiple applications that “work together”
- **Shared** between containers in same pod
 - Same VM, IP, port visibility
- **Not shared**
 - CPU/memory resources (etc.)
 - Files (great! Each can have their own Linux distro, packages versions, etc.)



Container orchestration

- Kubernetes currently is the most popular container orchestrator.
 - A container orchestrator can launch many containers in a cluster (of VMs or physical machines).
- Other orchestrators:
 - **Docker compose**: only launches containers on one node (so not necessarily an “orchestrator” depending on definition)
 - **Docker swarm**: built from compose to support multiple nodes
 - **Nomad**: simpler alternative to Kubernetes

Docker demo