

Serverless Computing

DS 5110/CS 5501: Big Data Systems

Spring 2024

Lecture 8b

Yue Cheng



Some material taken/derived from:

- Berkeley CS 262a (Spring '18) by Ali Ghodsi and Ion Stoica;
- Tyler Harter's HotCloud '18 OpenLambda talk;

@ 2024 released for use under a [CC BY-SA](https://creativecommons.org/licenses/by-sa/4.0/) license.

Learning objectives

- Understand the motivation behind serverless computing
- Know the different generations of cloud computing
 - Virtual machines
 - Containers
 - Serverless functions
- Know FaaS' current limitations

Motivation

When to use the cloud?

- Data
 - Large amounts of data – can't store locally
 - Shared data across users
 - Long-term storage
- Compute
 - Need lots of CPUs for data processing
 - Varying computing demands (resources)
 - No admin (for managing your local hardware)

When to use the cloud?

- Data
 - Large amounts of data – can't store
 - Shared data across users
 - Long-term storage
- Compute
 - Need lots of CPUs for data processing
 - Varying computing demands (resources)
 - No admin



Why is there no “cloud button”?



EC2 [RDS](#)

Region: US East (N. Virginia) -

Cost: Hourly -

Reserved: 1-year - No Upfront -

Columns -

Compare Selected

Clear Filters

CSV

Filter: Min Memory (GiB): Min vCPUs: Min Storage (GiB):

| Name | API Name | Memory | vCPUs | Instance Storage | Network Performance | Linux On Demand cost | Linux Reserved cost | Windows On Demand cost | Windows Reserved cost |
|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|
| <input type="text" value="Search"/> | <input type="text" value="Search"/> | <input type="text" value="Search"/> | <input type="text" value="Search"/> | <input type="text" value="Search"/> | <input type="text" value="Search"/> | <input type="text" value="Search"/> | <input type="text" value="Search"/> | <input type="text" value="Search"/> | <input type="text" value="Search"/> |
| M5DN Extra Large | m5dn.xlarge | 16.0 GiB | 4 vCPUs | 150 GiB NVMe SSD | Up to 25 Gigabit | \$0.272000 hourly | \$0.173000 hourly | \$0.456000 hourly | \$0.357000 hourly |
| M5A Double Extra Large | m5a.2xlarge | 32.0 GiB | 8 vCPUs | EBS only | Up to 10 Gigabit | \$0.344000 hourly | \$0.219000 hourly | \$0.712000 hourly | \$0.587000 hourly |
| R5N 12xlarge | r5n.12xlarge | 384.0 GiB | 48 vCPUs | EBS only | 50 Gigabit | \$3.576000 hourly | \$2.253000 hourly | \$5.784000 hourly | \$4.461000 hourly |
| R5AD Extra Large | r5ad.xlarge | 32.0 GiB | 4 vCPUs | 150 GiB NVMe SSD | 10 Gigabit | \$0.262000 hourly | \$0.166000 hourly | \$0.446000 hourly | \$0.350000 hourly |
| R5N Extra Large | r5n.xlarge | 32.0 GiB | 4 vCPUs | EBS only | Up to 25 Gigabit | \$0.298000 hourly | \$0.188000 hourly | \$0.482000 hourly | \$0.372000 hourly |
| I3EN 12xlarge | i3en.12xlarge | 384.0 GiB | 48 vCPUs | 30000 GiB (4 * 7500 GiB NVMe SSD) | 50 Gigabit | \$5.424000 hourly | \$3.694000 hourly | \$7.632000 hourly | \$5.902000 hourly |
| I3EN Metal | i3en.metal | 768.0 GiB | 96 vCPUs | 60000 GiB (8 * 7500 GiB NVMe SSD) | 100 Gigabit | \$10.848000 hourly | \$7.388000 hourly | \$15.264000 hourly | \$11.804000 hourly |
| R5DN Extra Large | r5dn.xlarge | 32.0 GiB | 4 vCPUs | 150 GiB NVMe SSD | Up to 25 Gigabit | \$0.334000 hourly | \$0.211000 hourly | \$0.518000 hourly | \$0.395000 hourly |
| I2 Extra Large | i2.xlarge | 30.5 GiB | 4 vCPUs | 800 GiB SSD | Moderate | \$0.853000 hourly | \$0.424000 hourly | \$0.973000 hourly | \$0.565000 hourly |
| M5N 16xlarge | m5n.16xlarge | 256.0 GiB | 64 vCPUs | EBS only | 75 Gigabit | \$3.808000 hourly | \$2.419000 hourly | \$6.752000 hourly | \$5.363000 hourly |
| T2 Micro | t2.micro | 1.0 GiB | 1 vCPUs <u>for a 2h 24m burst</u> | EBS only | Low to Moderate | \$0.011600 hourly | \$0.007200 hourly | \$0.016200 hourly | \$0.011800 hourly |
| D2 Eight Extra Large | d2.8xlarge | 244.0 GiB | 36 vCPUs | 48000 GiB (24 * 2000 GiB HDD) | 10 Gigabit | \$5.520000 hourly | \$3.216000 hourly | \$6.198000 hourly | \$3.300000 hourly |
| I3EN 3xlarge | i3en.3xlarge | 96.0 GiB | 12 vCPUs | 7500 GiB NVMe SSD | Up to 25 Gigabit | \$1.356000 hourly | \$0.924000 hourly | \$1.908000 hourly | \$1.476000 hourly |
| Z1D 3xlarge | z1d.3xlarge | 96.0 GiB | 12 vCPUs | 450 GiB NVMe SSD | Up to 10 Gigabit | \$1.116000 hourly | \$0.705000 hourly | \$1.668000 hourly | \$1.257000 hourly |
| X1E 16xlarge | x1e.16xlarge | 1952.0 GiB | 64 vCPUs | 1920 GiB SSD | 10 Gigabit | \$13.344000 hourly | \$8.223000 hourly | \$16.288000 hourly | \$11.167000 hourly |
| R5N 24xlarge | r5n.24xlarge | 768.0 GiB | 96 vCPUs | EBS only | 100 Gigabit | \$7.152000 hourly | \$4.506000 hourly | \$11.568000 hourly | \$8.922000 hourly |
| I2 Eight Extra Large | i2.8xlarge | 244.0 GiB | 32 vCPUs | 6400 GiB (8 * 800 GiB SSD) | 10 Gigabit | \$6.820000 hourly | \$3.392000 hourly | \$7.782000 hourly | \$4.521000 hourly |
| R5A Eight Extra Large | r5a.8xlarge | 256.0 GiB | 32 vCPUs | EBS only | Up to 10 Gigabit | \$1.808000 hourly | \$1.141000 hourly | \$3.280000 hourly | \$2.613000 hourly |
| A1 Metal | a1.metal | 32.0 GiB | 16 vCPUs | EBS only | Up to 10 Gigabit | \$0.408000 hourly | \$0.257000 hourly | unavailable | unavailable |
| I2 Double Extra Large | i2.2xlarge | 61.0 GiB | 8 vCPUs | 1600 GiB (2 * 800 GiB SSD) | High | \$1.705000 hourly | \$0.848000 hourly | \$1.946000 hourly | \$1.131000 hourly |
| I3EN Double Extra Large | i3en.2xlarge | 64.0 GiB | 8 vCPUs | 5000 GiB (2 * 2500 GiB NVMe SSD) | Up to 25 Gigabit | \$0.904000 hourly | \$0.616000 hourly | \$1.272000 hourly | \$0.984000 hourly |
| M5A Extra Large | m5a.xlarge | 16.0 GiB | 4 vCPUs | EBS only | Up to 10 Gigabit | \$0.172000 hourly | \$0.109000 hourly | \$0.356000 hourly | \$0.293000 hourly |
| P3 Double Extra Large | p3.2xlarge | 61.0 GiB | 8 vCPUs | EBS only | Up to 10 Gigabit | \$3.060000 hourly | \$2.088000 hourly | \$3.428000 hourly | \$2.456000 hourly |
| T2 Double Extra Large | t2.2xlarge | 32.0 GiB | 8 vCPUs <u>for a 4h 4.8m burst</u> | EBS only | Moderate | \$0.371200 hourly | \$0.230000 hourly | \$0.433200 hourly | \$0.292000 hourly |
| H1 Eight Extra Large | h1.8xlarge | 128.0 GiB | 32 vCPUs | 8000 GiB (4 * 2000 GiB HDD) | 10 Gigabit | \$1.872000 hourly | \$1.272000 hourly | \$3.344000 hourly | \$2.744000 hourly |
| R5D 24xlarge | r5d.24xlarge | 768.0 GiB | 96 vCPUs | 3600 GiB (4 * 900 GiB NVMe SSD) | 25 Gigabit | \$6.912000 hourly | \$4.362000 hourly | \$11.328000 hourly | \$8.778000 hourly |
| I3EN 6xlarge | i3en.6xlarge | 192.0 GiB | 24 vCPUs | 15000 GiB (2 * 7500 GiB NVMe SSD) | 25 Gigabit | \$2.712000 hourly | \$1.847000 hourly | \$3.816000 hourly | \$2.951000 hourly |
| R4 High-Memory Eight Extra Large | r4.8xlarge | 244.0 GiB | 32 vCPUs | EBS only | 10 Gigabit | \$2.128000 hourly | \$1.344000 hourly | \$3.600000 hourly | \$2.816000 hourly |
| T2 Large | t2.large | 8.0 GiB | 2 vCPUs <u>for a 7h 12m burst</u> | EBS only | Low to Moderate | \$0.092800 hourly | \$0.057500 hourly | \$0.120800 hourly | \$0.085500 hourly |
| X1 Extra High-Memory 16xlarge | x1.16xlarge | 976.0 GiB | 64 vCPUs | 1920 GiB SSD | High | \$6.669000 hourly | \$4.110000 hourly | \$9.613000 hourly | \$7.054000 hourly |
| M5A 16xlarge | m5a.16xlarge | 256.0 GiB | 64 vCPUs | EBS only | 12 Gigabit | \$2.752000 hourly | \$1.751000 hourly | \$5.696000 hourly | \$4.695000 hourly |
| R5 Metal | r5.metal | 768.0 GiB | 96 vCPUs | EBS only | 25 Gigabit | \$6.048000 hourly | \$3.810000 hourly | \$10.464000 hourly | \$8.226000 hourly |
| R5A Large | r5a.large | 16.0 GiB | 2 vCPUs | EBS only | 10 Gigabit | \$0.113000 hourly | \$0.071000 hourly | \$0.205000 hourly | \$0.163000 hourly |
| C3 High-CPU Large | c3.large | 3.75 GiB | 2 vCPUs | 32 GiB (2 * 16 GiB SSD) | Moderate | \$0.105000 hourly | \$0.073000 hourly | \$0.188000 hourly | \$0.165000 hourly |
| R5A 24xlarge | r5a.24xlarge | 768.0 GiB | 96 vCPUs | EBS only | 20 Gigabit | \$5.424000 hourly | \$3.423000 hourly | \$9.840000 hourly | \$7.839000 hourly |
| G3 16xlarge | g3.16xlarge | 488.0 GiB | 64 vCPUs | EBS only | 20 Gigabit | \$4.560000 hourly | \$3.112200 hourly | \$7.504000 hourly | \$6.056200 hourly |
| A1 Double Extra Large | a1.2xlarge | 16.0 GiB | 8 vCPUs | EBS only | Up to 10 Gigabit | \$0.204000 hourly | \$0.128500 hourly | unavailable | unavailable |
| C4 High-CPU Extra Large | c4.xlarge | 7.5 GiB | 4 vCPUs | EBS only | High | \$0.199000 hourly | \$0.126000 hourly | \$0.383000 hourly | \$0.310000 hourly |
| X1E Quadruple Extra Large | x1e.4xlarge | 488.0 GiB | 16 vCPUs | 480 GiB SSD | Up to 10 Gigabit | \$3.336000 hourly | \$2.056000 hourly | \$4.072000 hourly | \$2.792000 hourly |
| M5AD Extra Large | m5ad.xlarge | 16.0 GiB | 4 vCPUs | 150 GiB NVMe SSD | Up to 10 Gigabit | \$0.206000 hourly | \$0.132000 hourly | \$0.390000 hourly | \$0.316000 hourly |

EC2 Instances (724)

Based on your inputs, this is the lowest-cost EC2 instance: **t4g.nano**

#thecloudistoodamnhard

1. What type of instances?
2. How many to spin up?
3. What base image?
4. On-demand or spot?
5. What storage service to use?
6. And then wait to start...
7. Not the end of the horror story:
 1. When to scale out?
 2. When to scale in?
 3. When to switch to different instance types?
8. Go back to Step 1...

EC2Instances.info Easy Amazon EC2 Instance Comparison

EC2 RDS

Region: US East (N. Virginia) - Cost: Hourly - Reserved: 1-year - No Upfront - Columns - Compare Selected Clear Filters CSV

Filter: Min Memory (GiB): 0 Min vCPUs: 0 Min Storage (GiB): 0

| Name | API Name | Memory | vCPUs | Instance Storage |
|----------------------------------|---------------|------------|----------------------------|----------------------------|
| Search | Search | Search | Search | Search |
| MSDN Extra Large | m5dn.xlarge | 16.0 GiB | 4 vCPUs | 150 GiB NV |
| MSA Double Extra Large | m5a.2xlarge | 32.0 GiB | 8 vCPUs | E |
| R5N 12xlarge | r5n.12xlarge | 384.0 GiB | 48 vCPUs | E |
| R5AD Extra Large | r5ad.xlarge | 32.0 GiB | 4 vCPUs | 150 GiB NV |
| R5N Extra Large | r5n.xlarge | 32.0 GiB | 4 vCPUs | E |
| I3EN 12xlarge | i3en.12xlarge | 384.0 GiB | 48 vCPUs | 30000 GiB (4 * 7500 GiB NV |
| I3EN Metal | i3en.metal | 768.0 GiB | 96 vCPUs | 60000 GiB (8 * 7500 GiB NV |
| R5DN Extra Large | r5dn.xlarge | 32.0 GiB | 4 vCPUs | 150 GiB NV |
| I2 Extra Large | i2.xlarge | 30.5 GiB | 4 vCPUs | 800 |
| MSN 16xlarge | m5n.16xlarge | 256.0 GiB | 64 vCPUs | E |
| T2 Micro | t2.micro | 1.0 GiB | 1 vCPUs for a 2h 24m burst | E |
| D2 Eight Extra Large | d2.8xlarge | 244.0 GiB | 36 vCPUs | 48000 GiB (24 * 2000 G |
| I3EN 3xlarge | i3en.3xlarge | 96.0 GiB | 12 vCPUs | 7500 GiB NV |
| Z1D 3xlarge | z1d.3xlarge | 96.0 GiB | 12 vCPUs | 450 GiB NV |
| X1E 16xlarge | x1e.16xlarge | 1952.0 GiB | 64 vCPUs | 1920 |
| R5N 24xlarge | r5n.24xlarge | 768.0 GiB | 96 vCPUs | E |
| I2 Eight Extra Large | i2.8xlarge | 244.0 GiB | 32 vCPUs | 6400 GiB (8 * 800 G |
| R5A Eight Extra Large | r5a.8xlarge | 256.0 GiB | 32 vCPUs | E |
| A1 Metal | a1.metal | 32.0 GiB | 16 vCPUs | E |
| I2 Double Extra Large | i2.2xlarge | 61.0 GiB | 8 vCPUs | 1600 GiB (2 * 800 G |
| I3EN Double Extra Large | i3en.2xlarge | 64.0 GiB | 8 vCPUs | 5000 GiB (2 * 2500 GiB NV |
| MSA Extra Large | m5a.xlarge | 16.0 GiB | 4 vCPUs | E |
| P3 Double Extra Large | p3.2xlarge | 61.0 GiB | 8 vCPUs | E |
| T2 Double Extra Large | t2.2xlarge | 32.0 GiB | 8 vCPUs for a 4h 48m burst | E |
| H1 Eight Extra Large | h1.8xlarge | 128.0 GiB | 32 vCPUs | 8000 GiB (4 * 2000 G |
| R5D 24xlarge | r5d.24xlarge | 768.0 GiB | 96 vCPUs | 3600 GiB (4 * 900 GiB NV |
| I3EN 6xlarge | i3en.6xlarge | 192.0 GiB | 24 vCPUs | 15000 GiB (2 * 7500 GiB NV |
| R4 High-Memory Eight Extra Large | r4.8xlarge | 244.0 GiB | 32 vCPUs | E |
| T2 Large | t2.large | 8.0 GiB | 2 vCPUs for a 7h 12m burst | E |
| X1 Extra High-Memory 16xlarge | x1.16xlarge | 976.0 GiB | 64 vCPUs | 1920 |
| MSA 16xlarge | m5a.16xlarge | 256.0 GiB | 64 vCPUs | E |
| R5 Metal | r5.metal | 768.0 GiB | 96 vCPUs | E |
| R5A Large | r5a.large | 16.0 GiB | 2 vCPUs | E |
| C3 High-CPU Large | c3.large | 3.75 GiB | 2 vCPUs | 32 GiB (2 * 16 |
| R5A 24xlarge | r5a.24xlarge | 768.0 GiB | 96 vCPUs | E |
| G3 16xlarge | g3.16xlarge | 488.0 GiB | 64 vCPUs | E |
| A1 Double Extra Large | a1.2xlarge | 16.0 GiB | 8 vCPUs | E |
| C4 High-CPU Extra Large | c4.xlarge | 7.5 GiB | 4 vCPUs | E |
| X1E Quadruple Extra Large | x1e.4xlarge | 488.0 GiB | 16 vCPUs | 480 |
| MSAD Extra Large | m5ad.xlarge | 16.0 GiB | 4 vCPUs | 150 GiB NV |

~~Decision paralysis??~~

Go for Serverless Computing!




Amazon
Lambda



Google Cloud Functions



Microsoft Azure Functions



What is serverless computing?

What is serverless computing?

Serverless computing (Function-as-a-Service, or FaaS) is a **programming abstraction** that enables users to upload programs, run them at **(virtually)** any scale, and **pay only for the resources used**

A car analogy

Car analogy



Own a car
(Bare metal servers)



Rent a car
(VPS)



City car-sharing
(Serverless)

Cars are parked **95%** of the time (loige.link/car-parked-95)

How much do you use the car?

<https://www.slideshare.net/loige/building-a-serverless-company-with-nodejs-react-and-the-serverless-framework-jsday-2017-verona>

Concept of serverless is not new

- Google App Engine
 - Fully managed platform as a service (PaaS) for developing and hosting web applications



- Google BigQuery
 - Fully managed data warehouse
 - “Arbitrarily” large data and queries
 - Pay per byte being processed
 - No concept of server or cluster




- AWS S3
 - Fully managed object storage service
 - Pay per byte being stored and written
 - No server maintenance or resource scaling




Cloud evolution history: A virtualization story


Classic cloud app stack




1st generation: virtual machine (VM)




1st generation: virtual machine (VM)




2nd generation: containers



3rd generation: serverless functions



3rd generation: serverless functions



Tradeoff discussion

Serverless functions
(AWS Lambdas)

Containers

VMs

Isolation?

Flexibility?

Overhead?

Above the surface: Core capability

1. (Provider) Manage a set of user-defined functions

Above the surface: Core capability

1. (Provider) Manage a set of user-defined functions
2. Take an event sent over HTTP or received from an event source

Above the surface: Core capability

1. (Provider) Manage a set of user-defined functions
2. Take an event sent over HTTP or received from an event source
3. Determine function(s) to which to dispatch the event

Above the surface: Core capability

1. (Provider) Manage a set of user-defined functions
2. Take an event sent over HTTP or received from an event source
3. Determine function(s) to which to dispatch the event
4. Find an existing instance of function or create a new one

Above the surface: Core capability

1. (Provider) Manage a set of user-defined functions
2. Take an event sent over HTTP or received from an event source
3. Determine function(s) to which to dispatch the event
4. Find an existing instance of function or create a new one
5. Send the event to the function instance

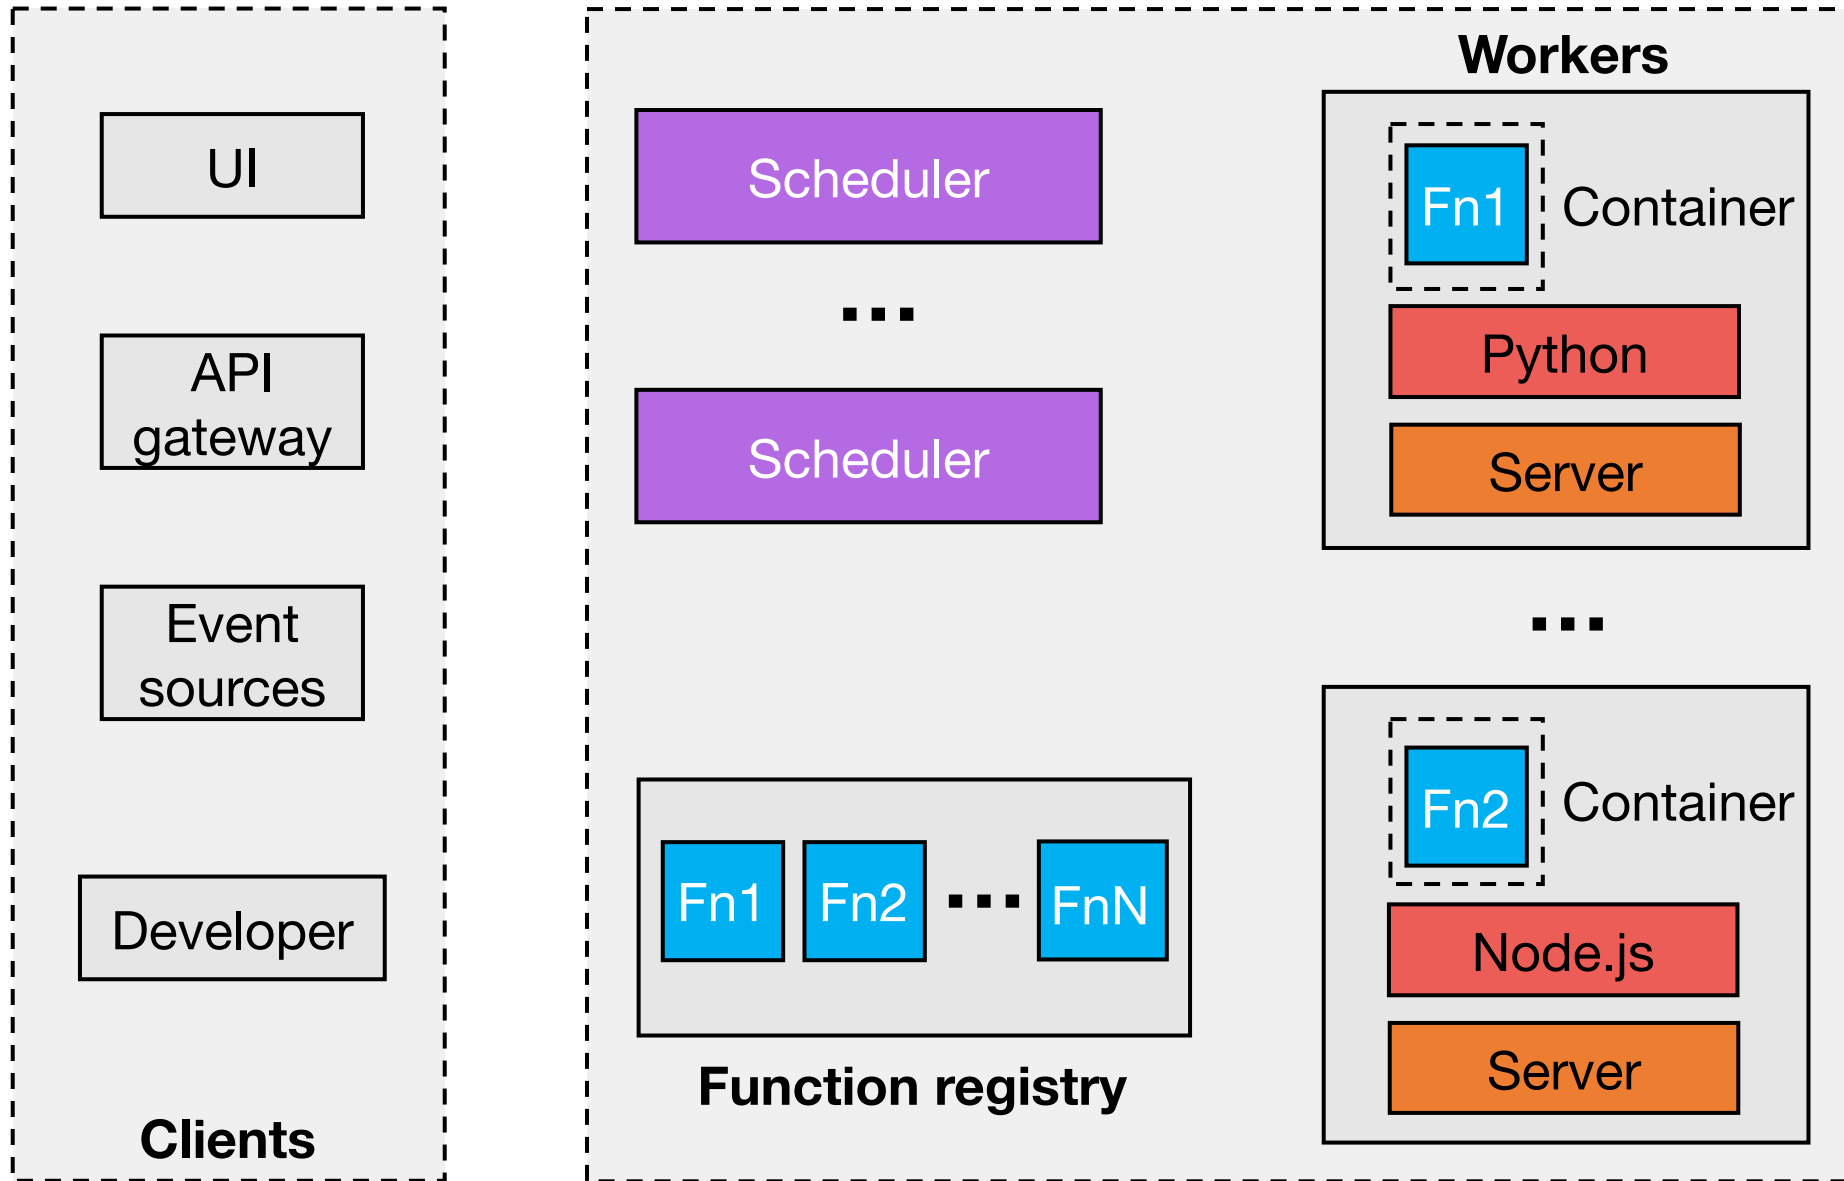
Above the surface: Core capability

1. (Provider) Manage a set of user-defined functions
2. Take an event sent over HTTP or received from an event source
3. Determine function(s) to which to dispatch the event
4. Find an existing instance of function or create a new one
5. Send the event to the function instance
6. Wait for a response

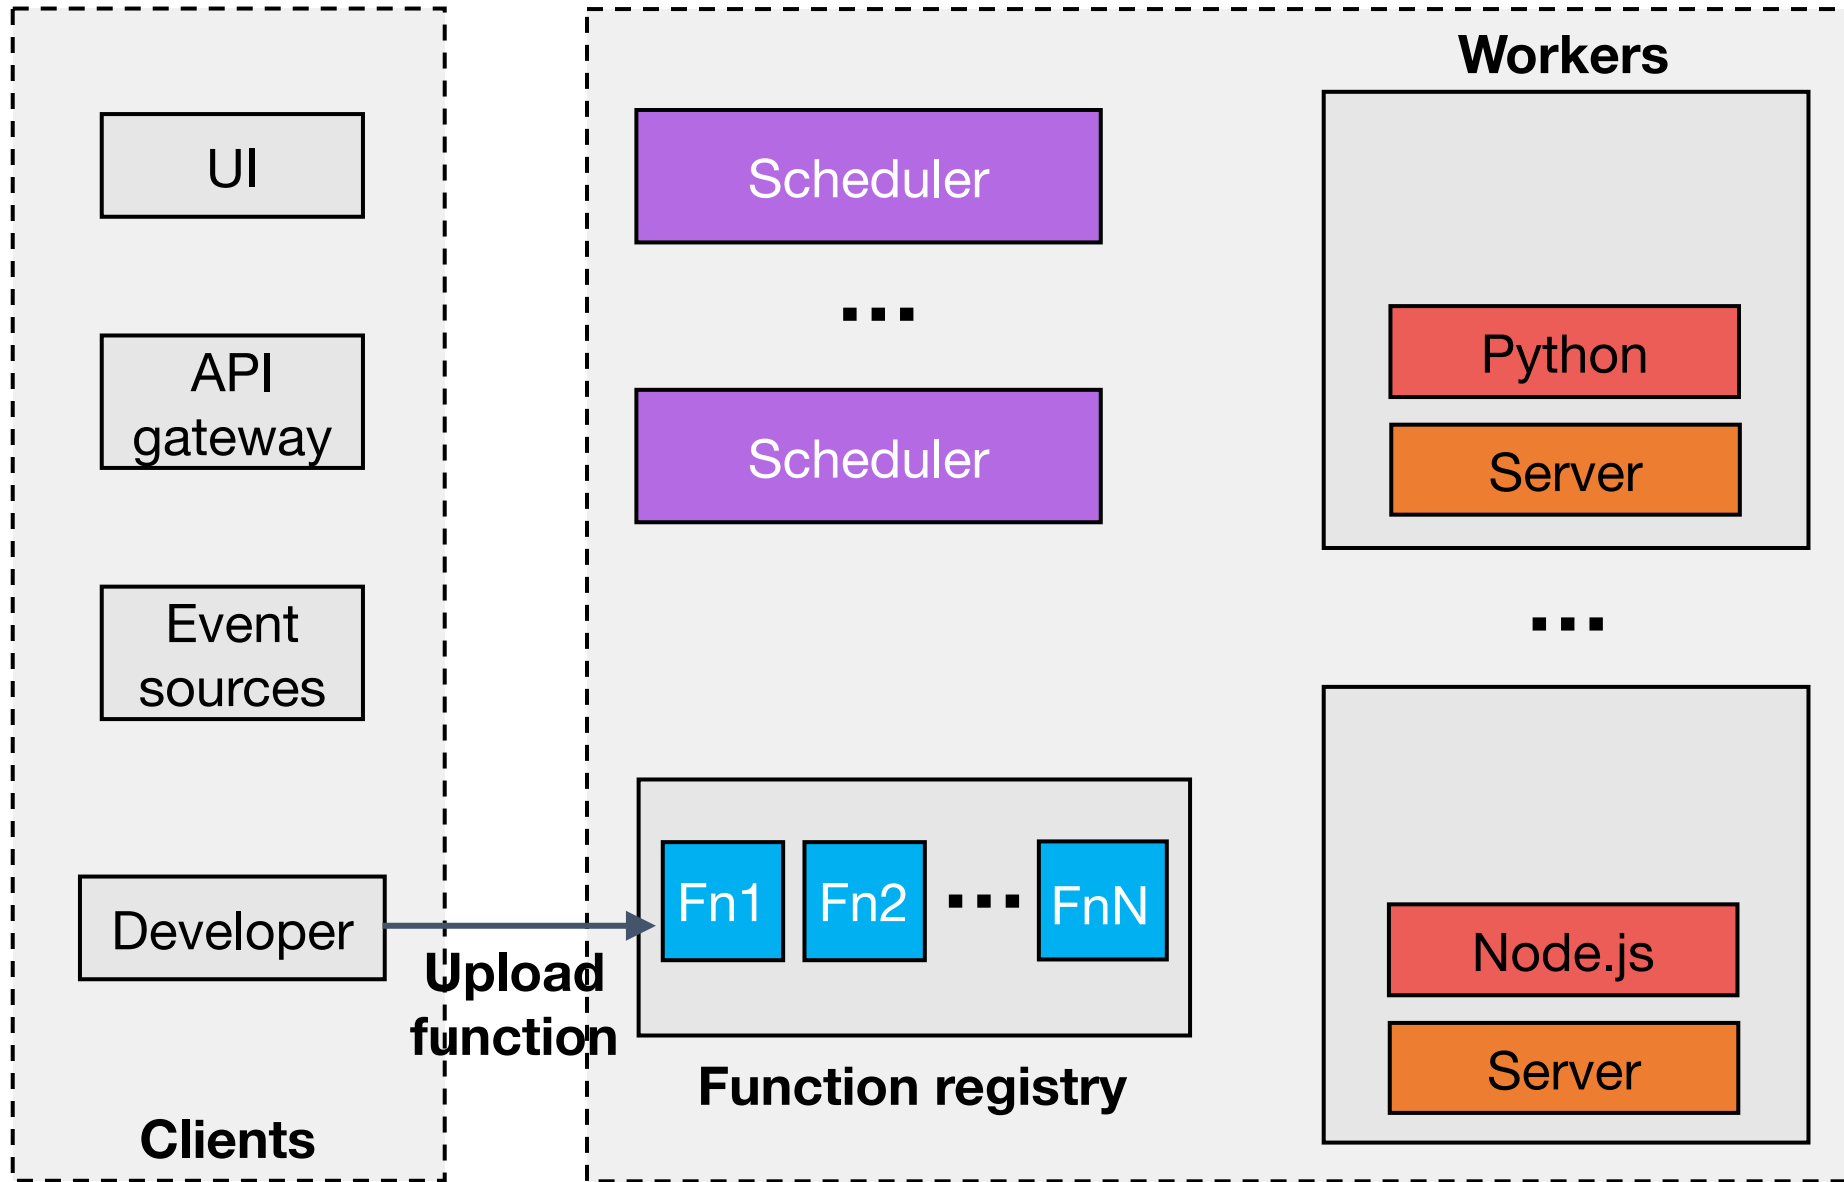
Above the surface: Core capability

1. (Provider) Manage a set of user-defined functions
2. Take an event sent over HTTP or received from an event source
3. Determine function(s) to which to dispatch the event
4. Find an existing instance of function or create a new one
5. Send the event to the function instance
6. Wait for a response
7. Gather execution logs
8. Make the response available to the user
9. Stop the function when the execution terminates

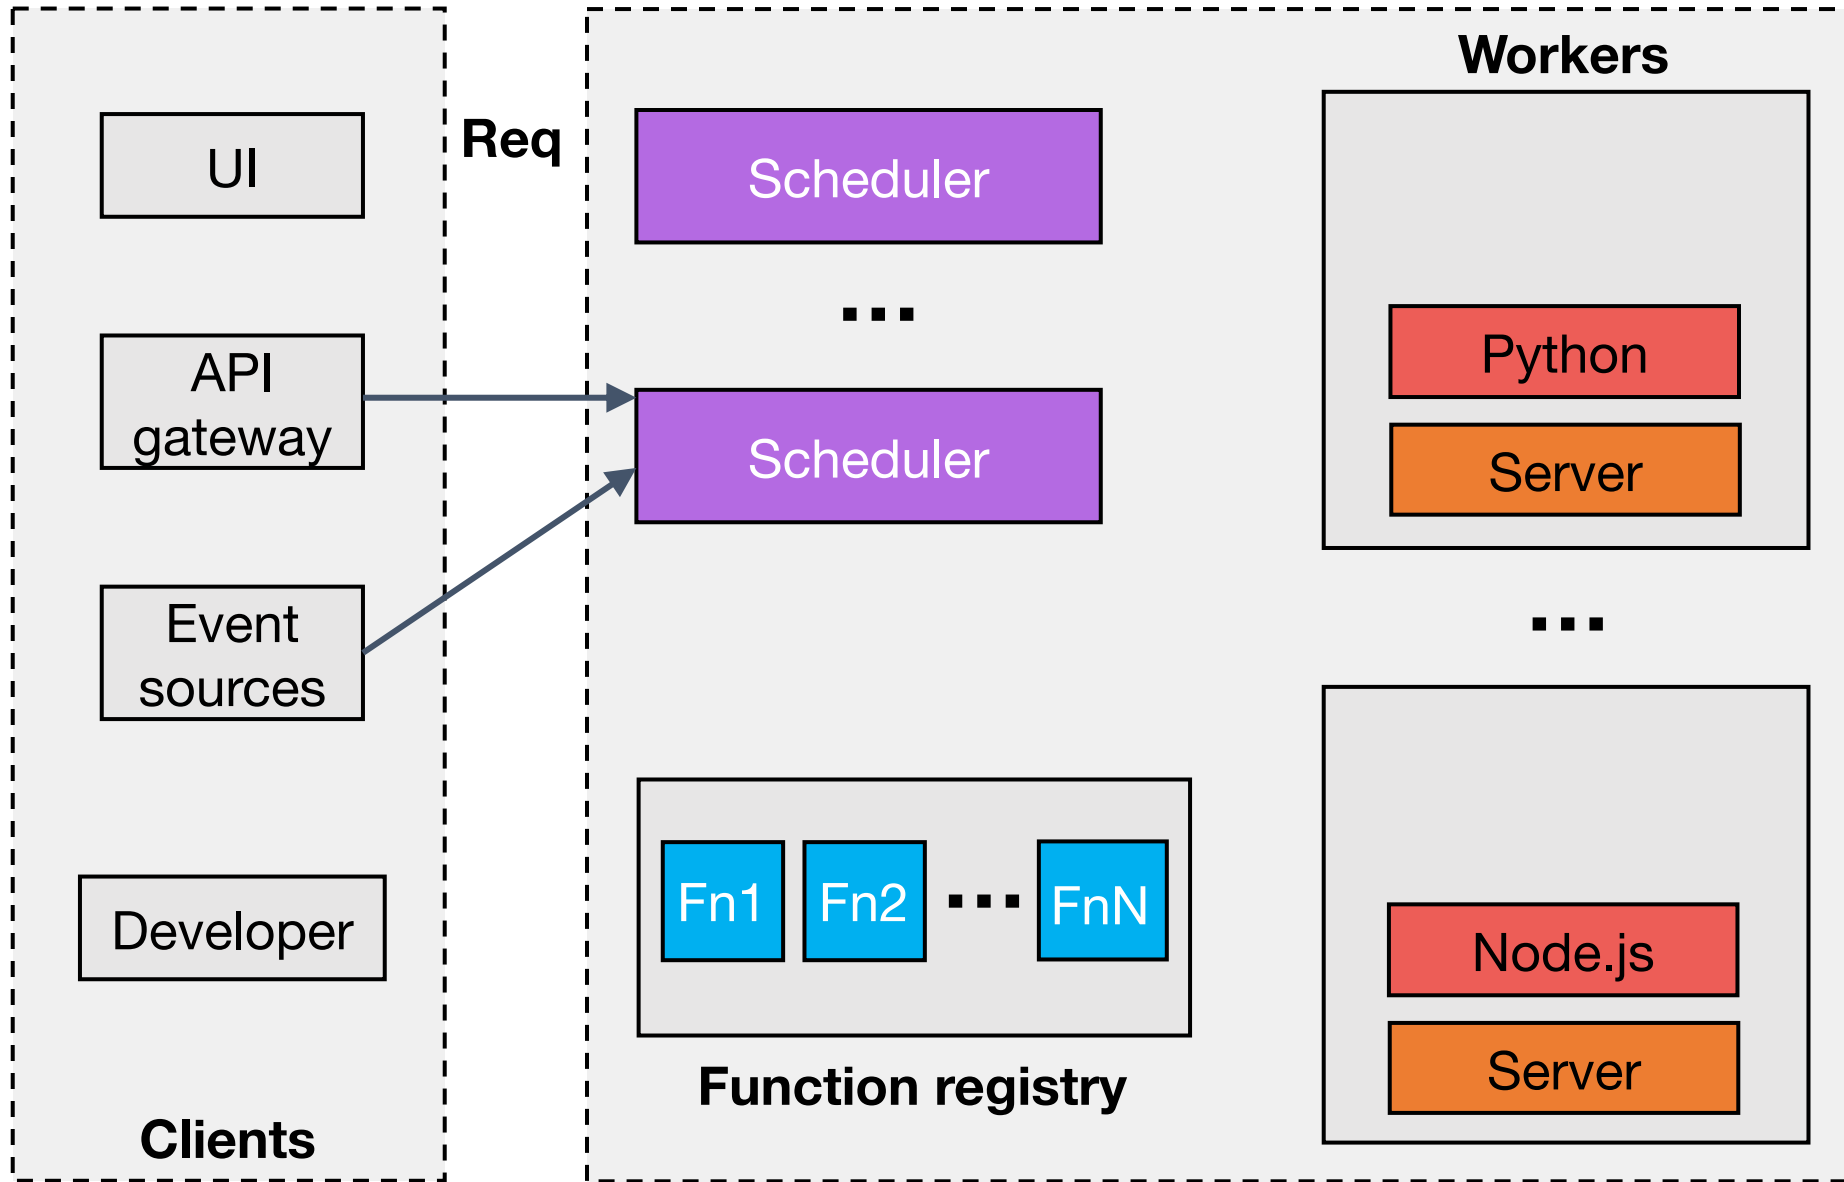
Under the hood: FaaS architecture



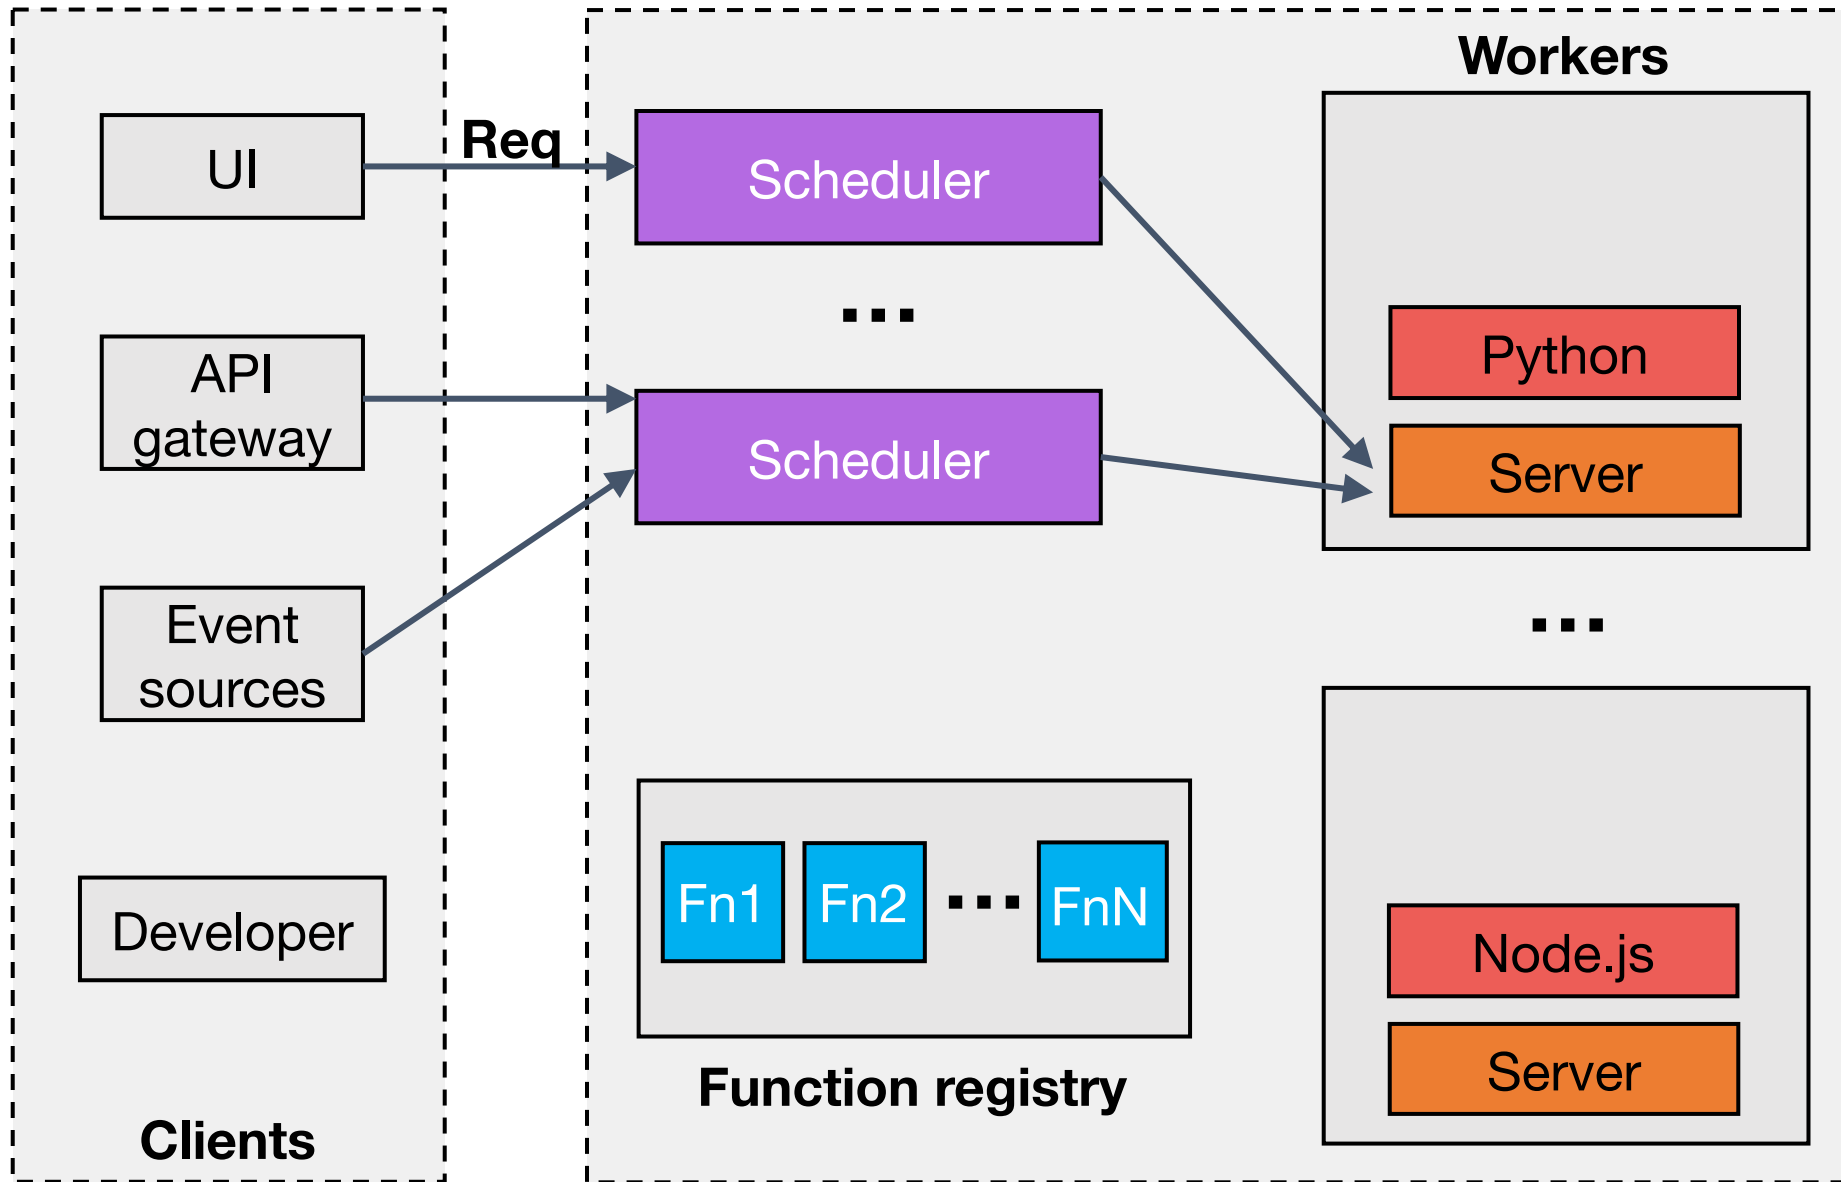
Under the hood: FaaS architecture



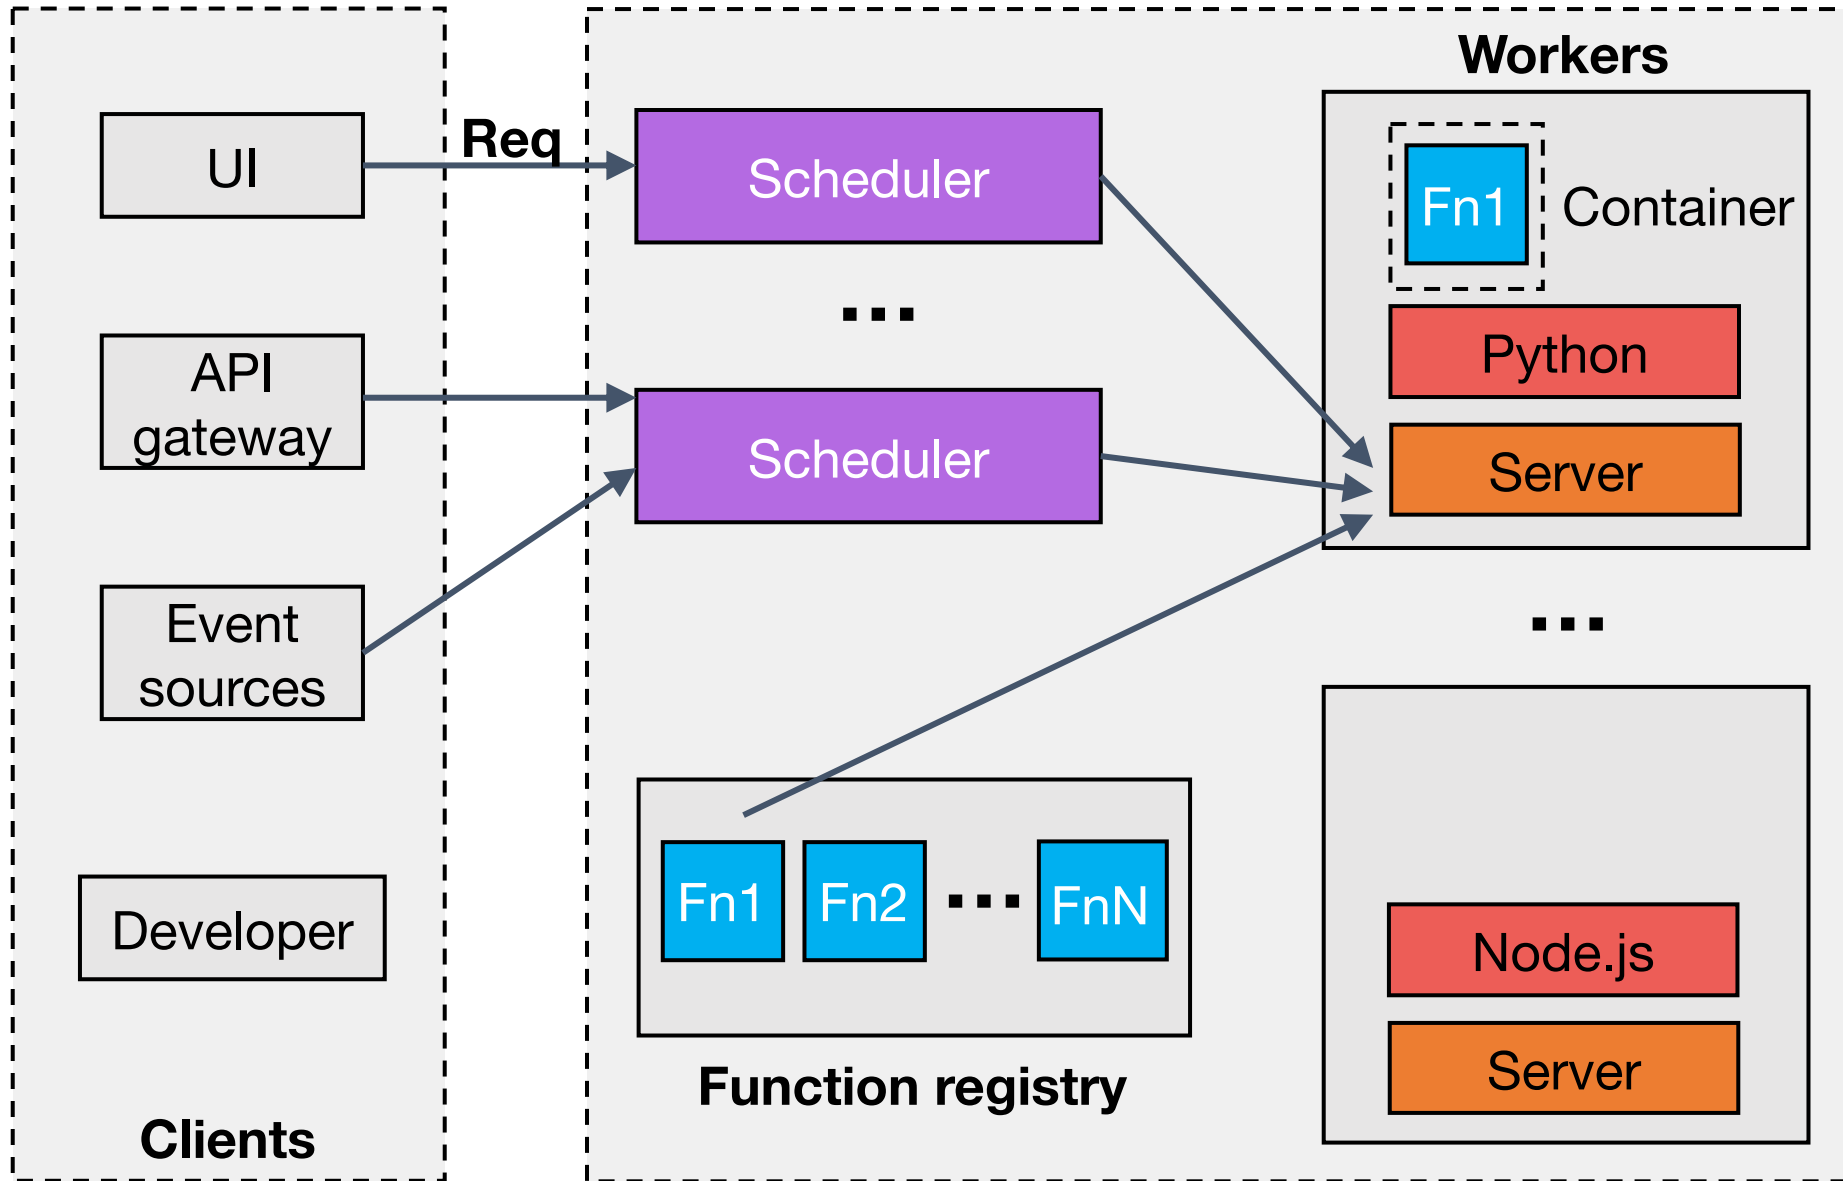
Under the hood: FaaS architecture



Under the hood: FaaS architecture

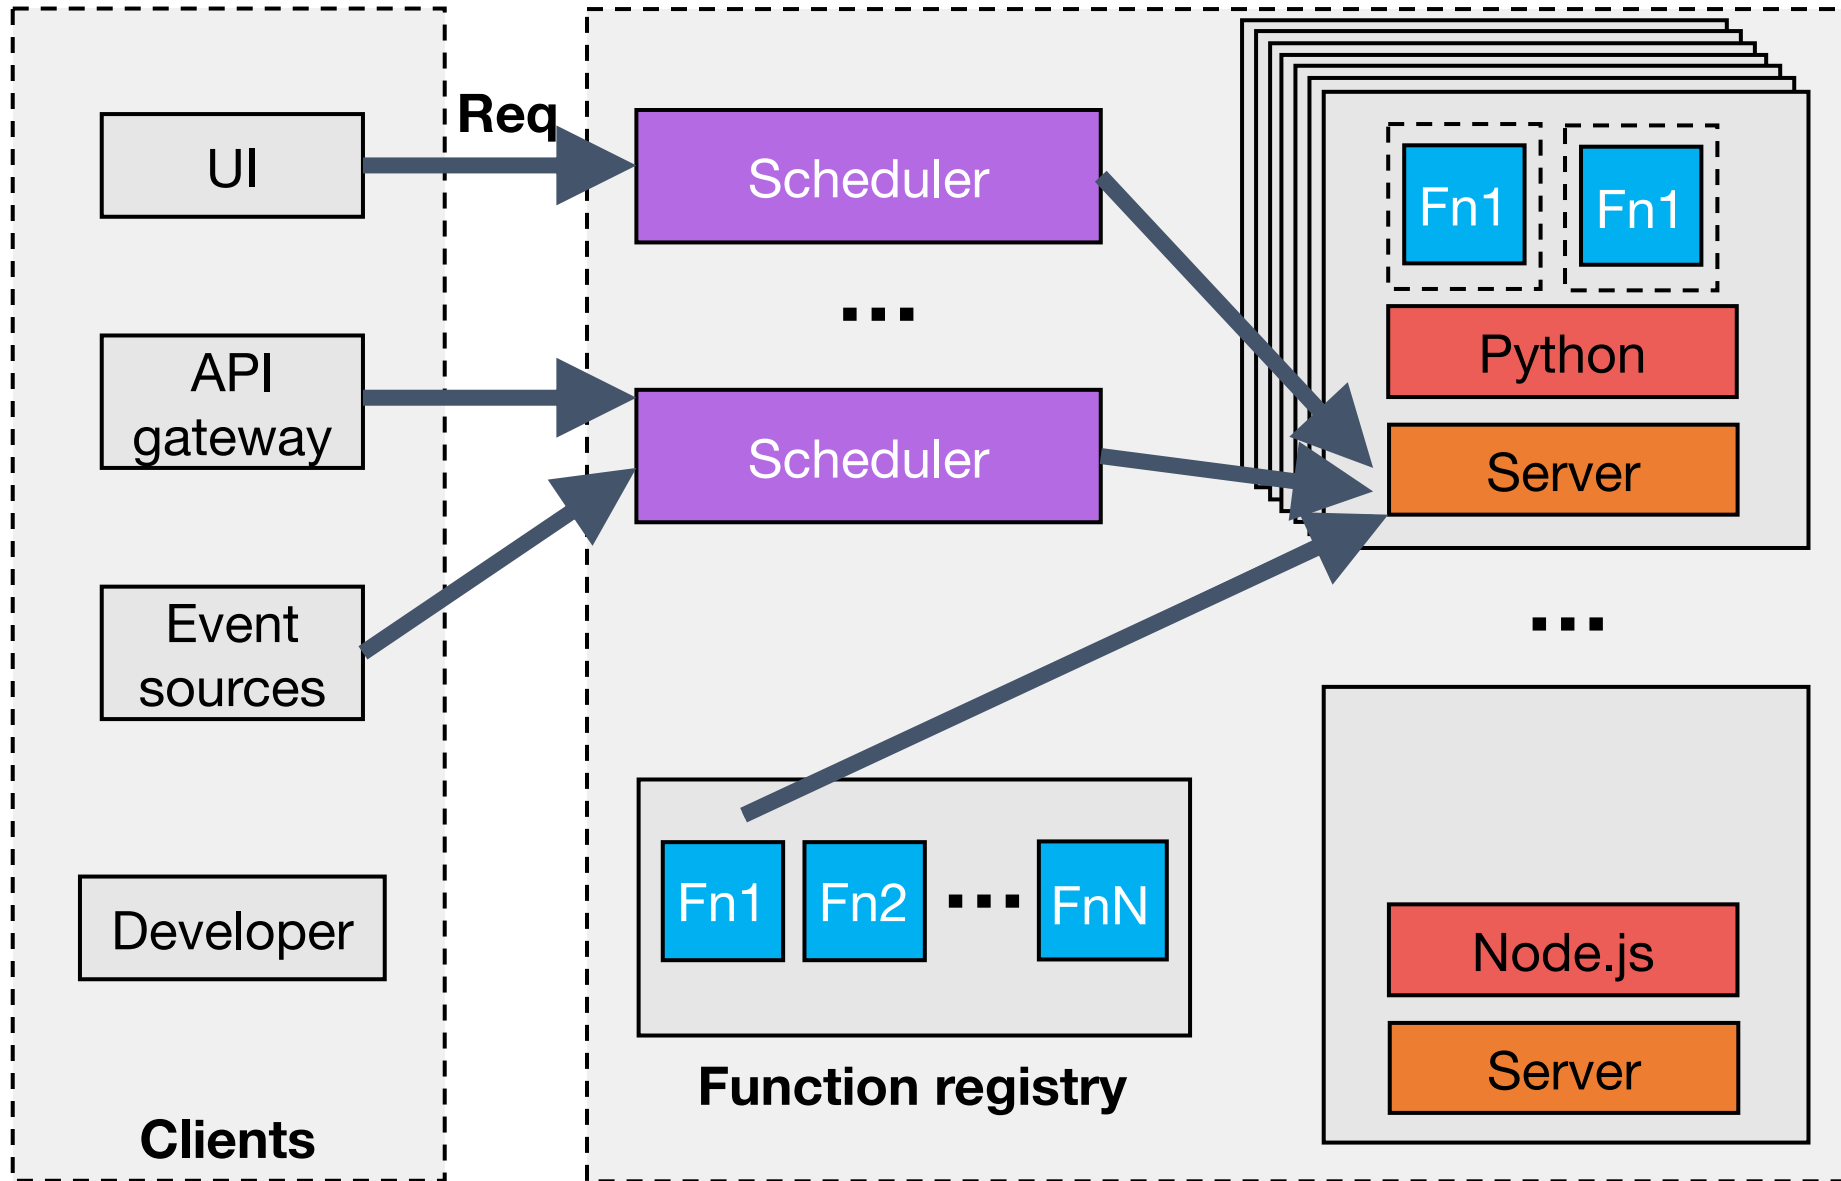


Under the hood: FaaS architecture



Under the hood: FaaS architecture

Servers are
auto-scaled



AWS Lambda

- Lambda capacity config keeps **evolving**:

1st gen { ~~300 seconds~~ 900 seconds (15 minutes)
~~single-core~~ **two-core** → up to 6 cores
~~1.5 GB~~ → 10 GB memory
~~512 MB~~ → up to 10GB of /tmp file system } **Current offering**

Python, Java, Node.js, Go, ...

Pricing:

- Fine-grained billing: **1-millisecond** billed duration
- \$0.20 per 1M requests (invocations charge \$)
- \$0.0000166667 for every GB-second (compute time charges \$\$)
- **6,000 1 GB Lambda functions for one second: 10¢**

Desirable properties

- (Near) zero administration overhead
 - No need to handle server provisioning, failure, etc.
- Elastic auto-scaling
 - Spin up / tear down functions quickly based on load
- Pay-per-use
 - Only pay for the resources used (CPU-mem bundle)

Limitations

- **Banned** inbound network
- **No guaranteed** data availability
- Lambdas are **resource-constrained**
- Lambdas have **limited execution time**
- **High cold startup cost** and **invocation cost**

AWS Lambda demo