

# Serverless Parallel Data Analytics

*DS 5110/CS 5501: Big Data Systems*

*Spring 2024*

Lecture 8c

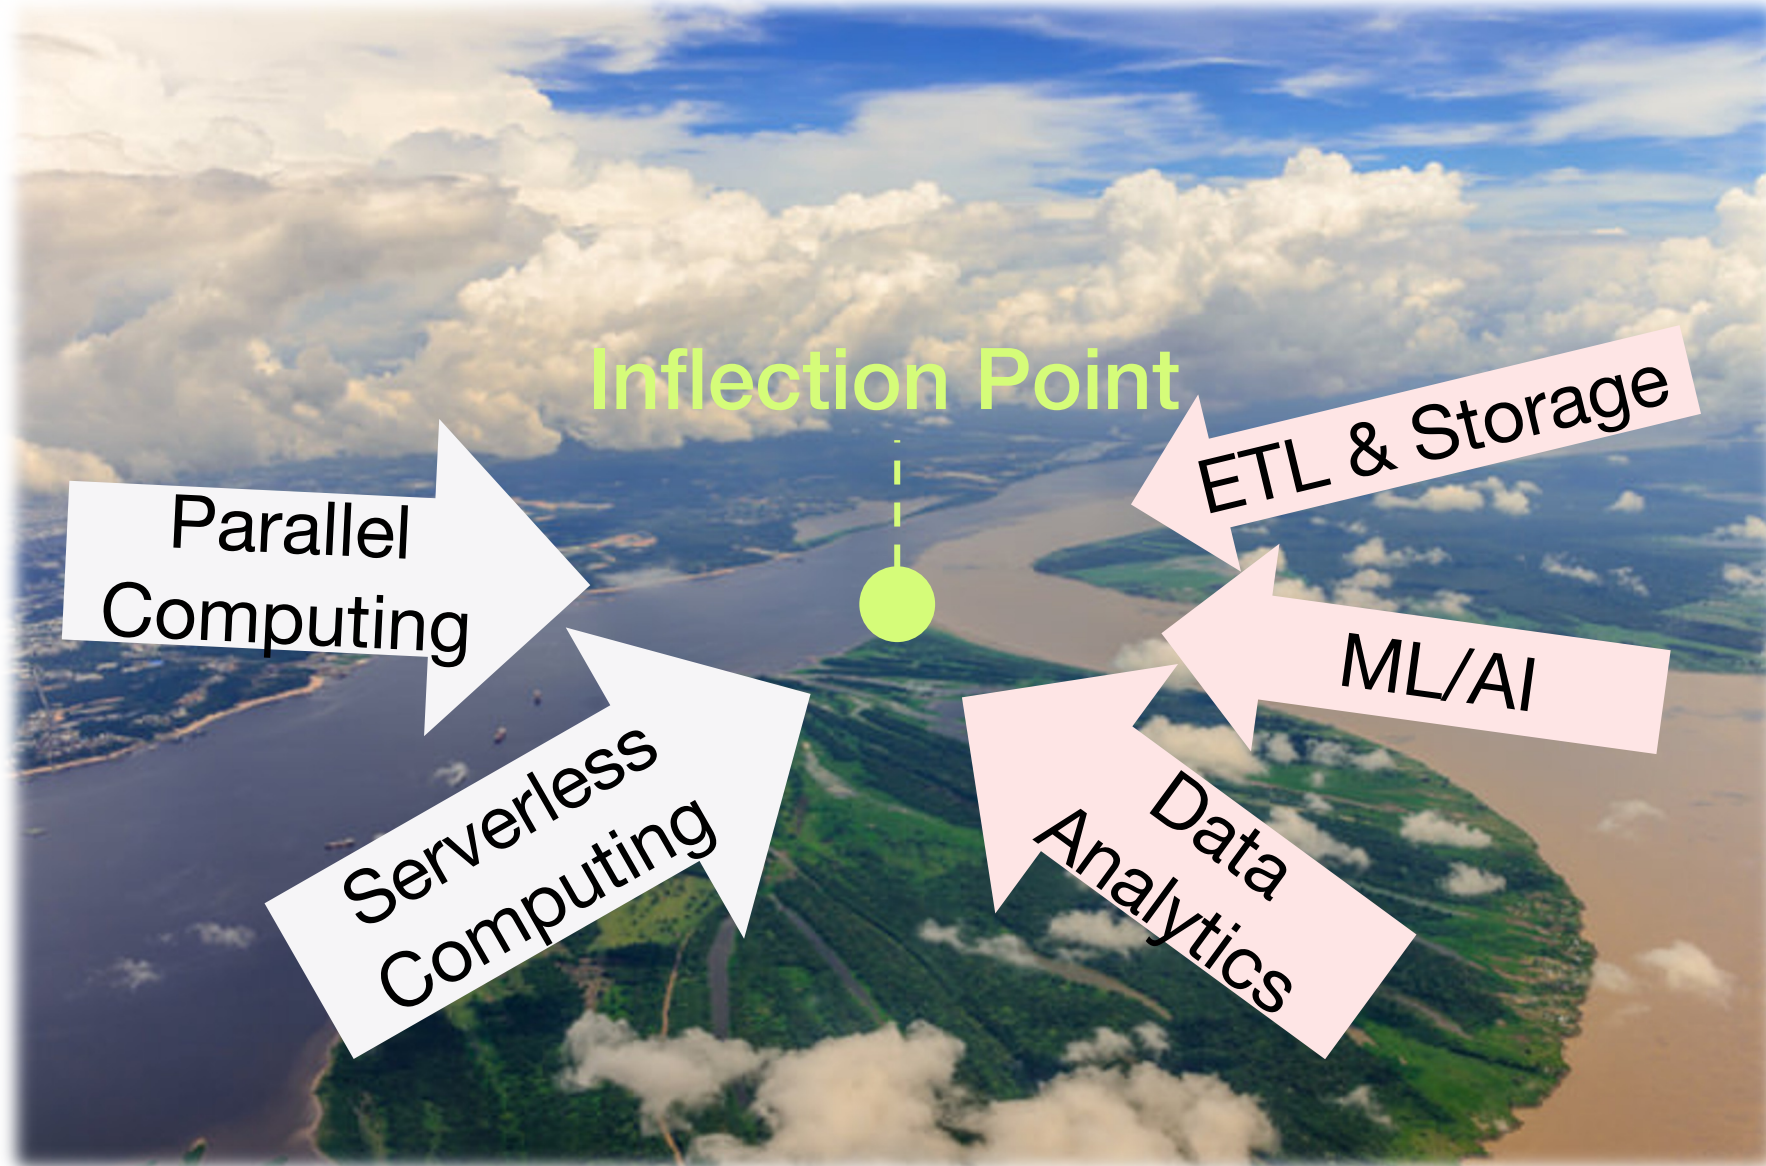
Yue Cheng



# Learning objectives

- Understand the challenges of supporting “stateful” computations on FaaS
- Know how PyWren works and its limitations
- Know how Wukong addresses some of PyWren’s limitations

# Confluence: When stateful apps meet serverless computing



# Today's data analytics landscape

Libraries efficient for  $O(1\text{MB})$



matplotlib

# Today's data analytics landscape

## Libraries efficient for O(1MB)



## Frameworks for O(100s GB)



# Today's data analytics landscape

Libraries efficient for  $O(1\text{MB})$



matplotlib

Frameworks for  $O(100\text{s GB})$




# Today's data analytics landscape

## Libraries efficient for $O(1\text{MB})$

- Easy to program (writing centralized code)
- Low barrier for environment setup (just installing libs)
- Well understood
- No scalability / elasticity
- Not able to efficiently handle large data

## Frameworks for $O(100\text{s GB})$

- Scale to 100s GB data
- Difficult to program and debug
  - Requires distributed systems knowledge
- No elasticity
- High barrier for environment setup
  - Requires low-level administration skills



# Today's data analytics landscape

## Libraries efficient for $O(1\text{MB})$

- **Easy-to-use**
- **Not scalable**
- **Not elastic**

## Frameworks for $O(100\text{s GB})$

- **Scalable**
- **Not easy-to-use**
- **Not elastic**



# Can we achieve all these desirable properties with **Serverless?**

Libraries efficient for  $O(1\text{MB})$

Frameworks for  $O(100\text{s GB})$

**Easy-to-use**

**Elastic**

**Pay-per-use**


**Scalable**

# Recap: What is serverless computing?

Many people define it many ways

A **programming abstraction** that enables users to upload programs, run them at **virtually** any scale, and pay **only for the resources used**

- **Function-as-a-Service (FaaS):** Cloud functions as a basic deployment unit




# Function-as-a-Service (FaaS)

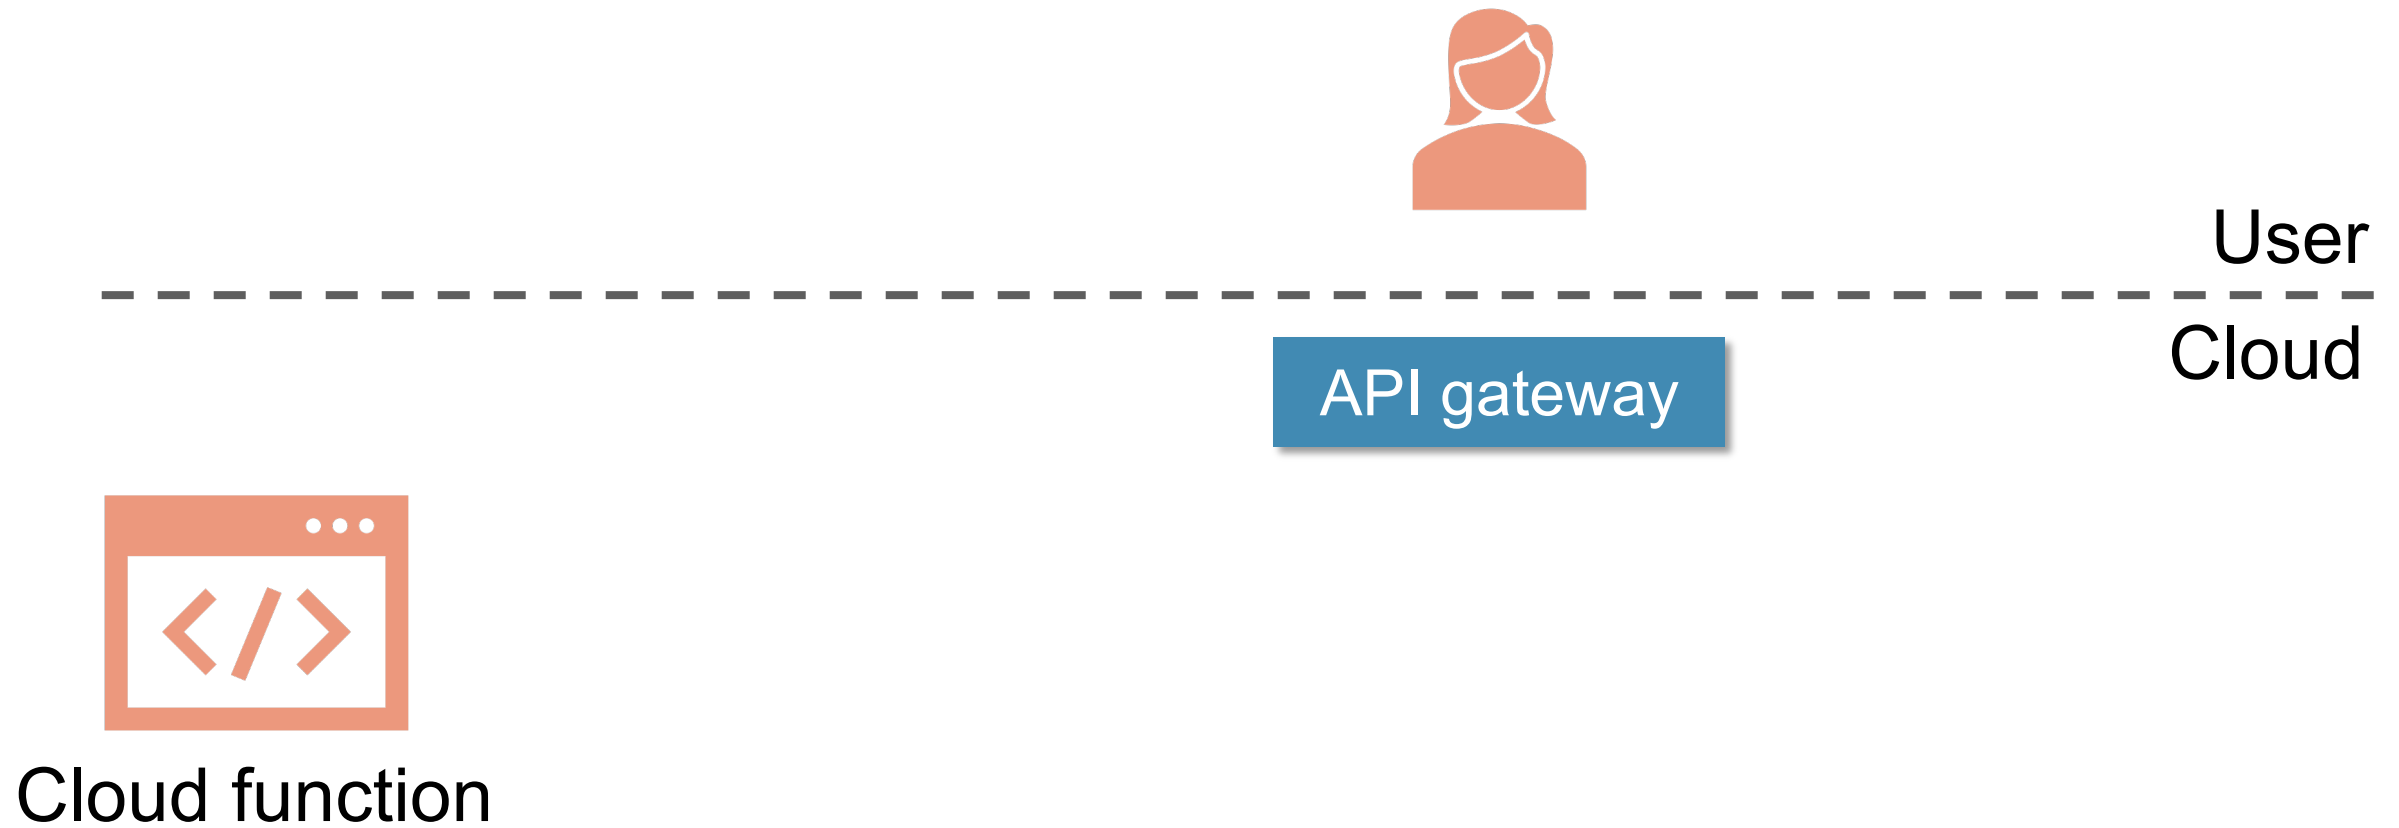


User  
Cloud

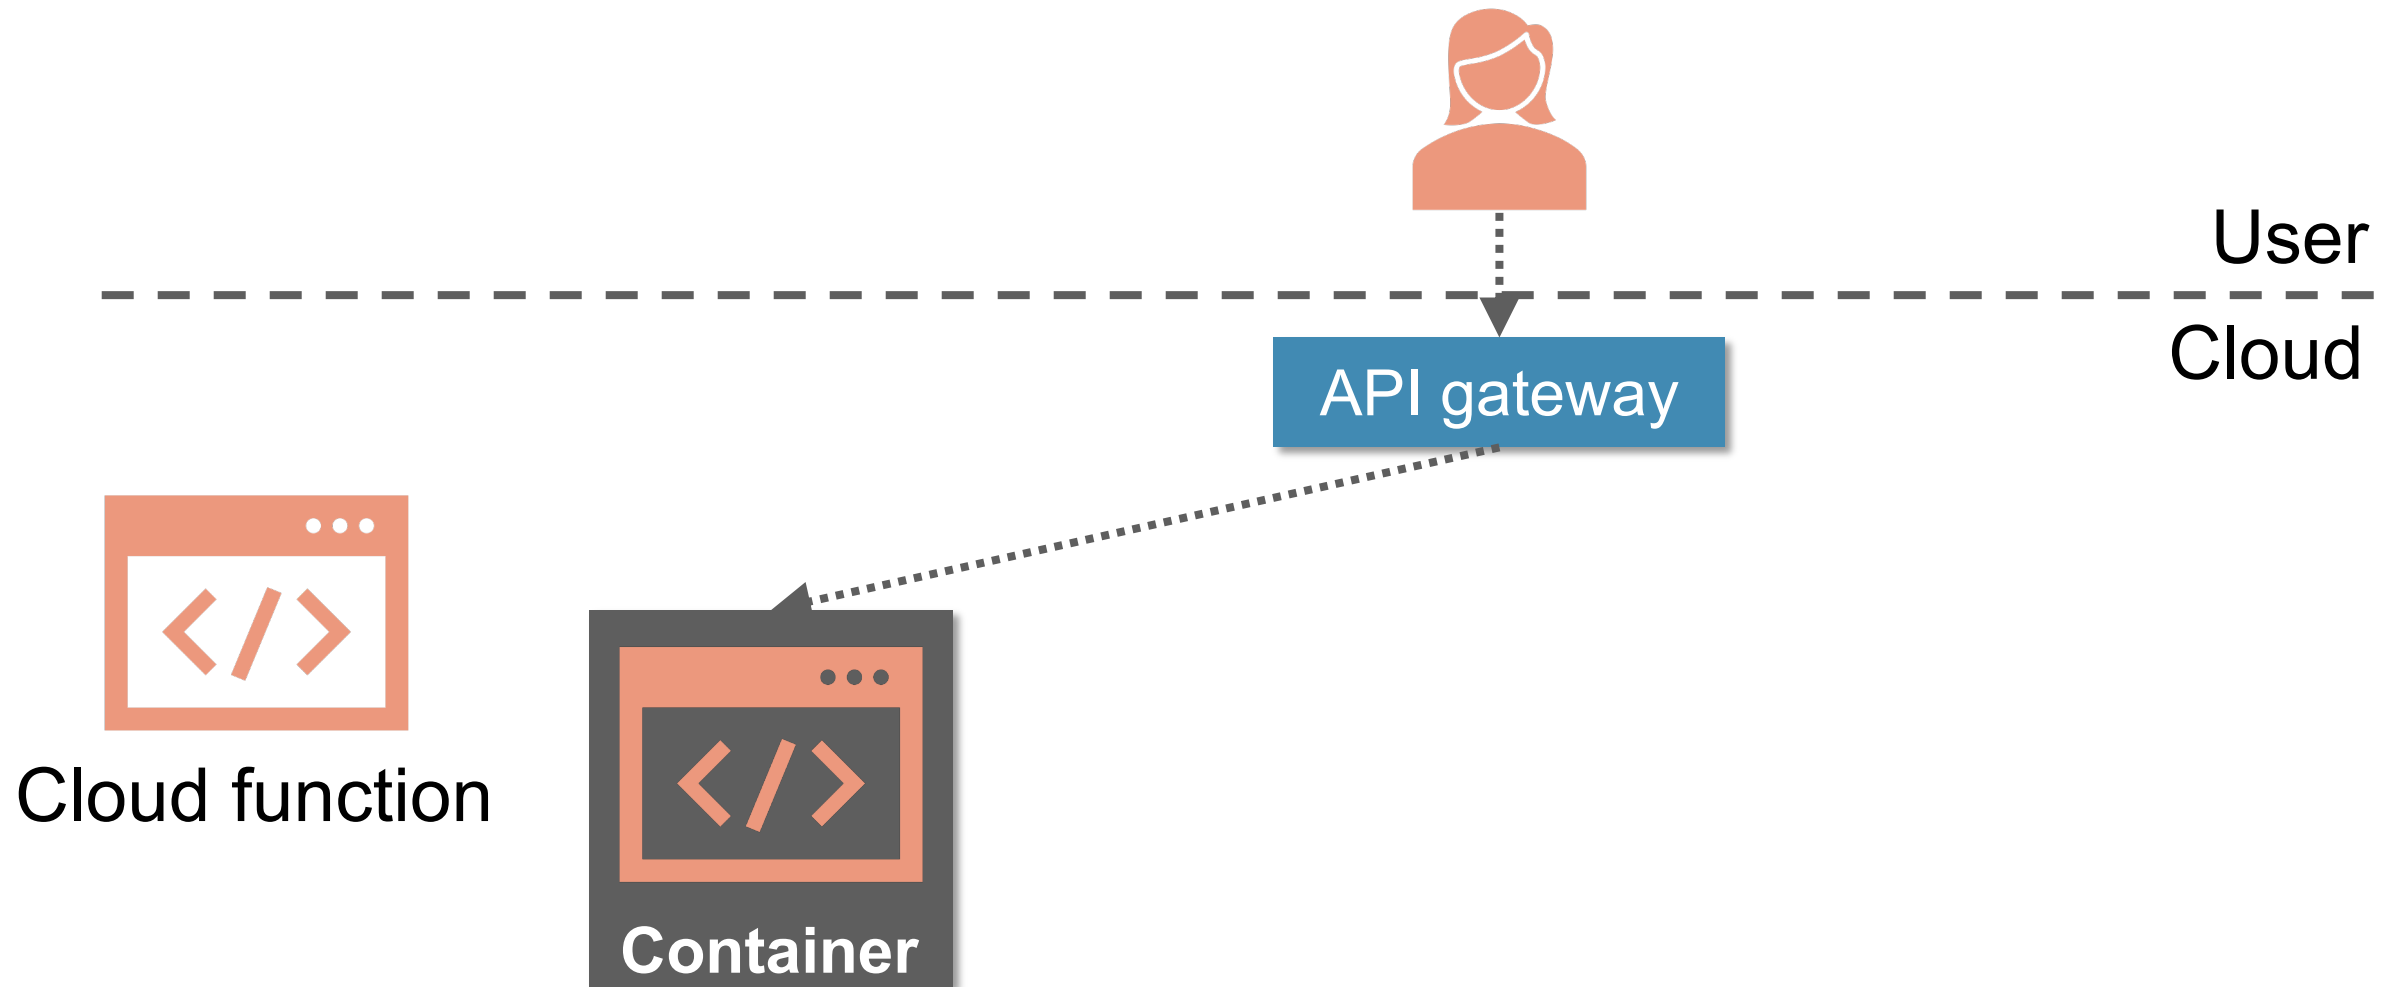
# Function-as-a-Service (FaaS)



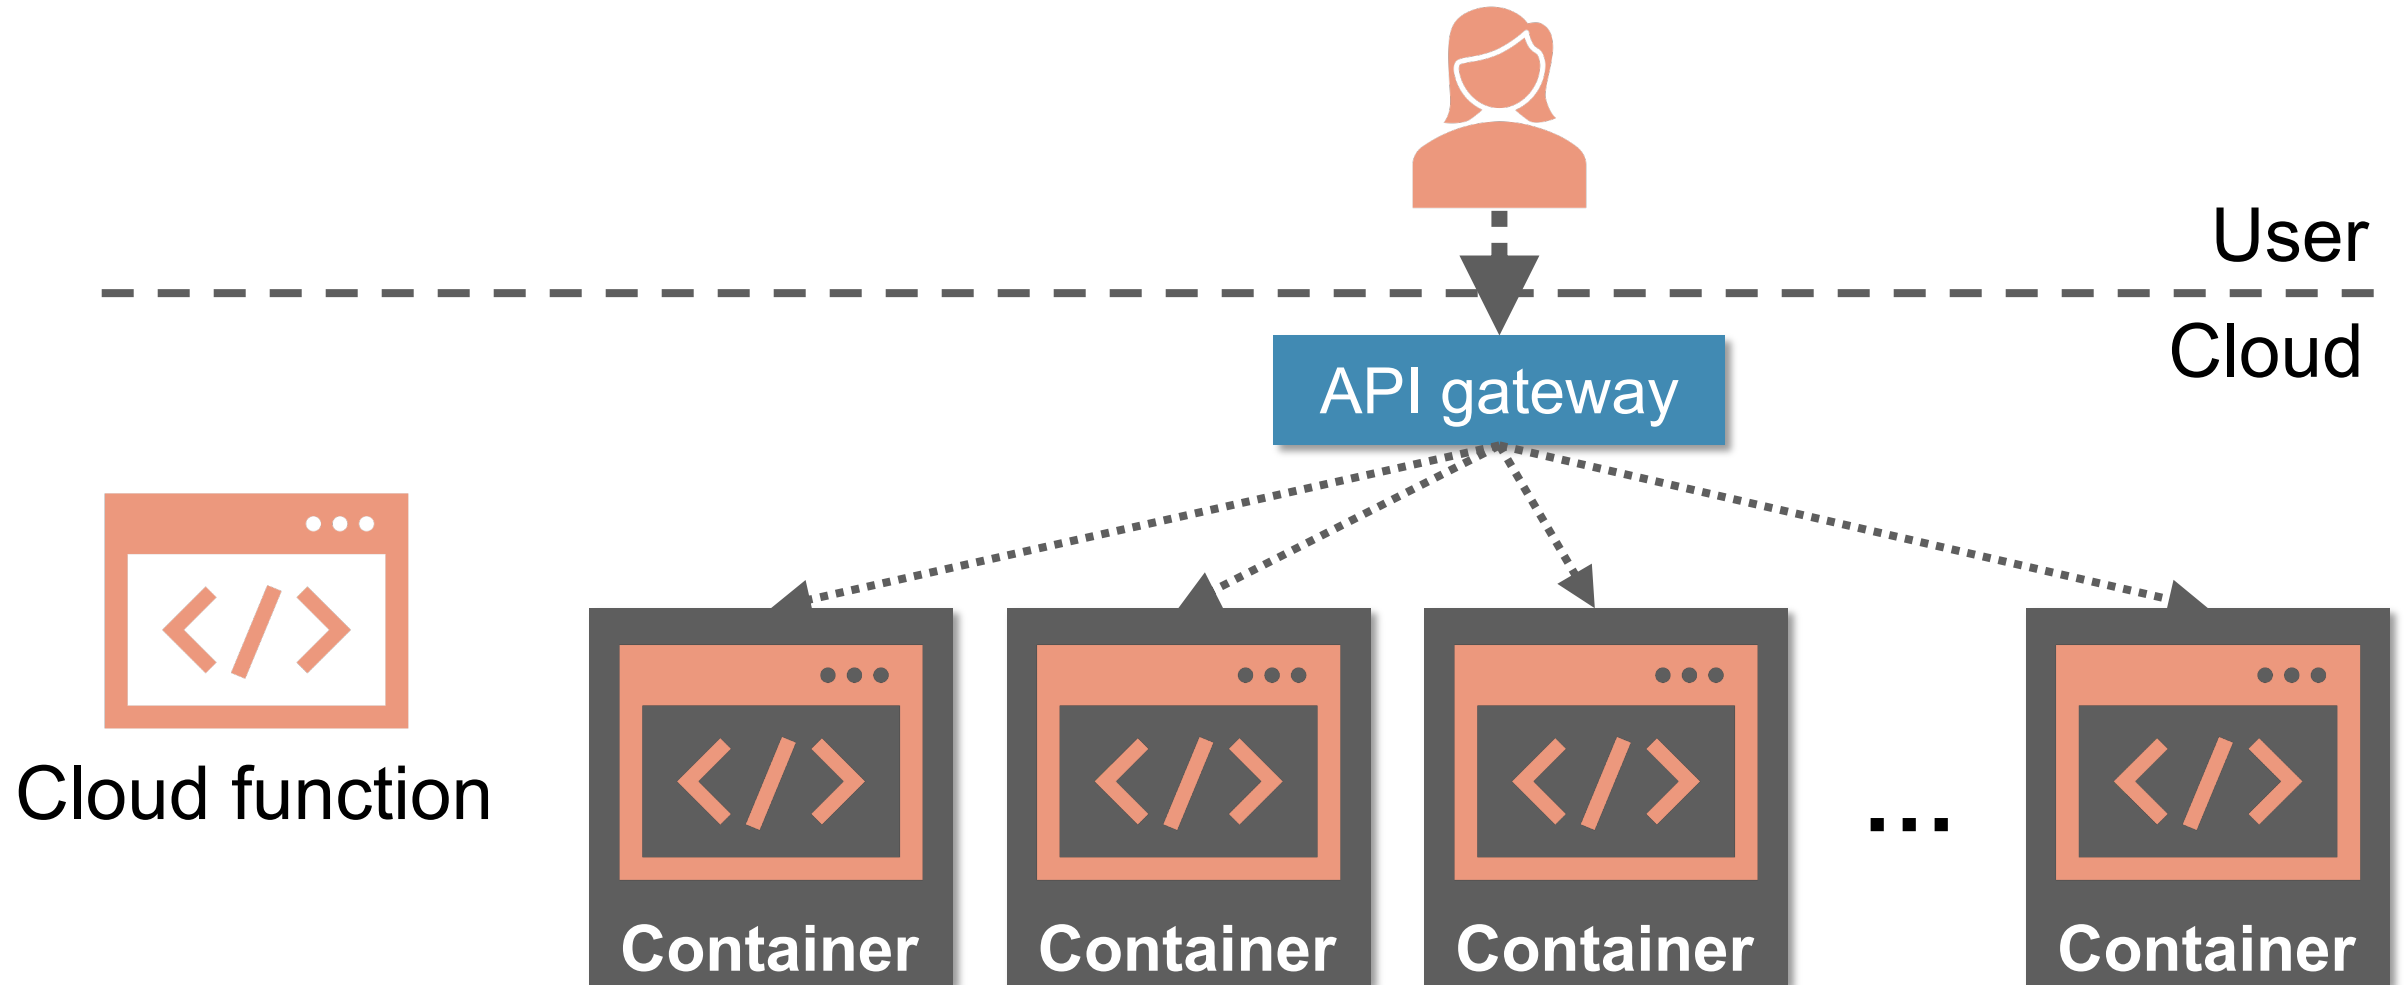
# Function-as-a-Service (FaaS)



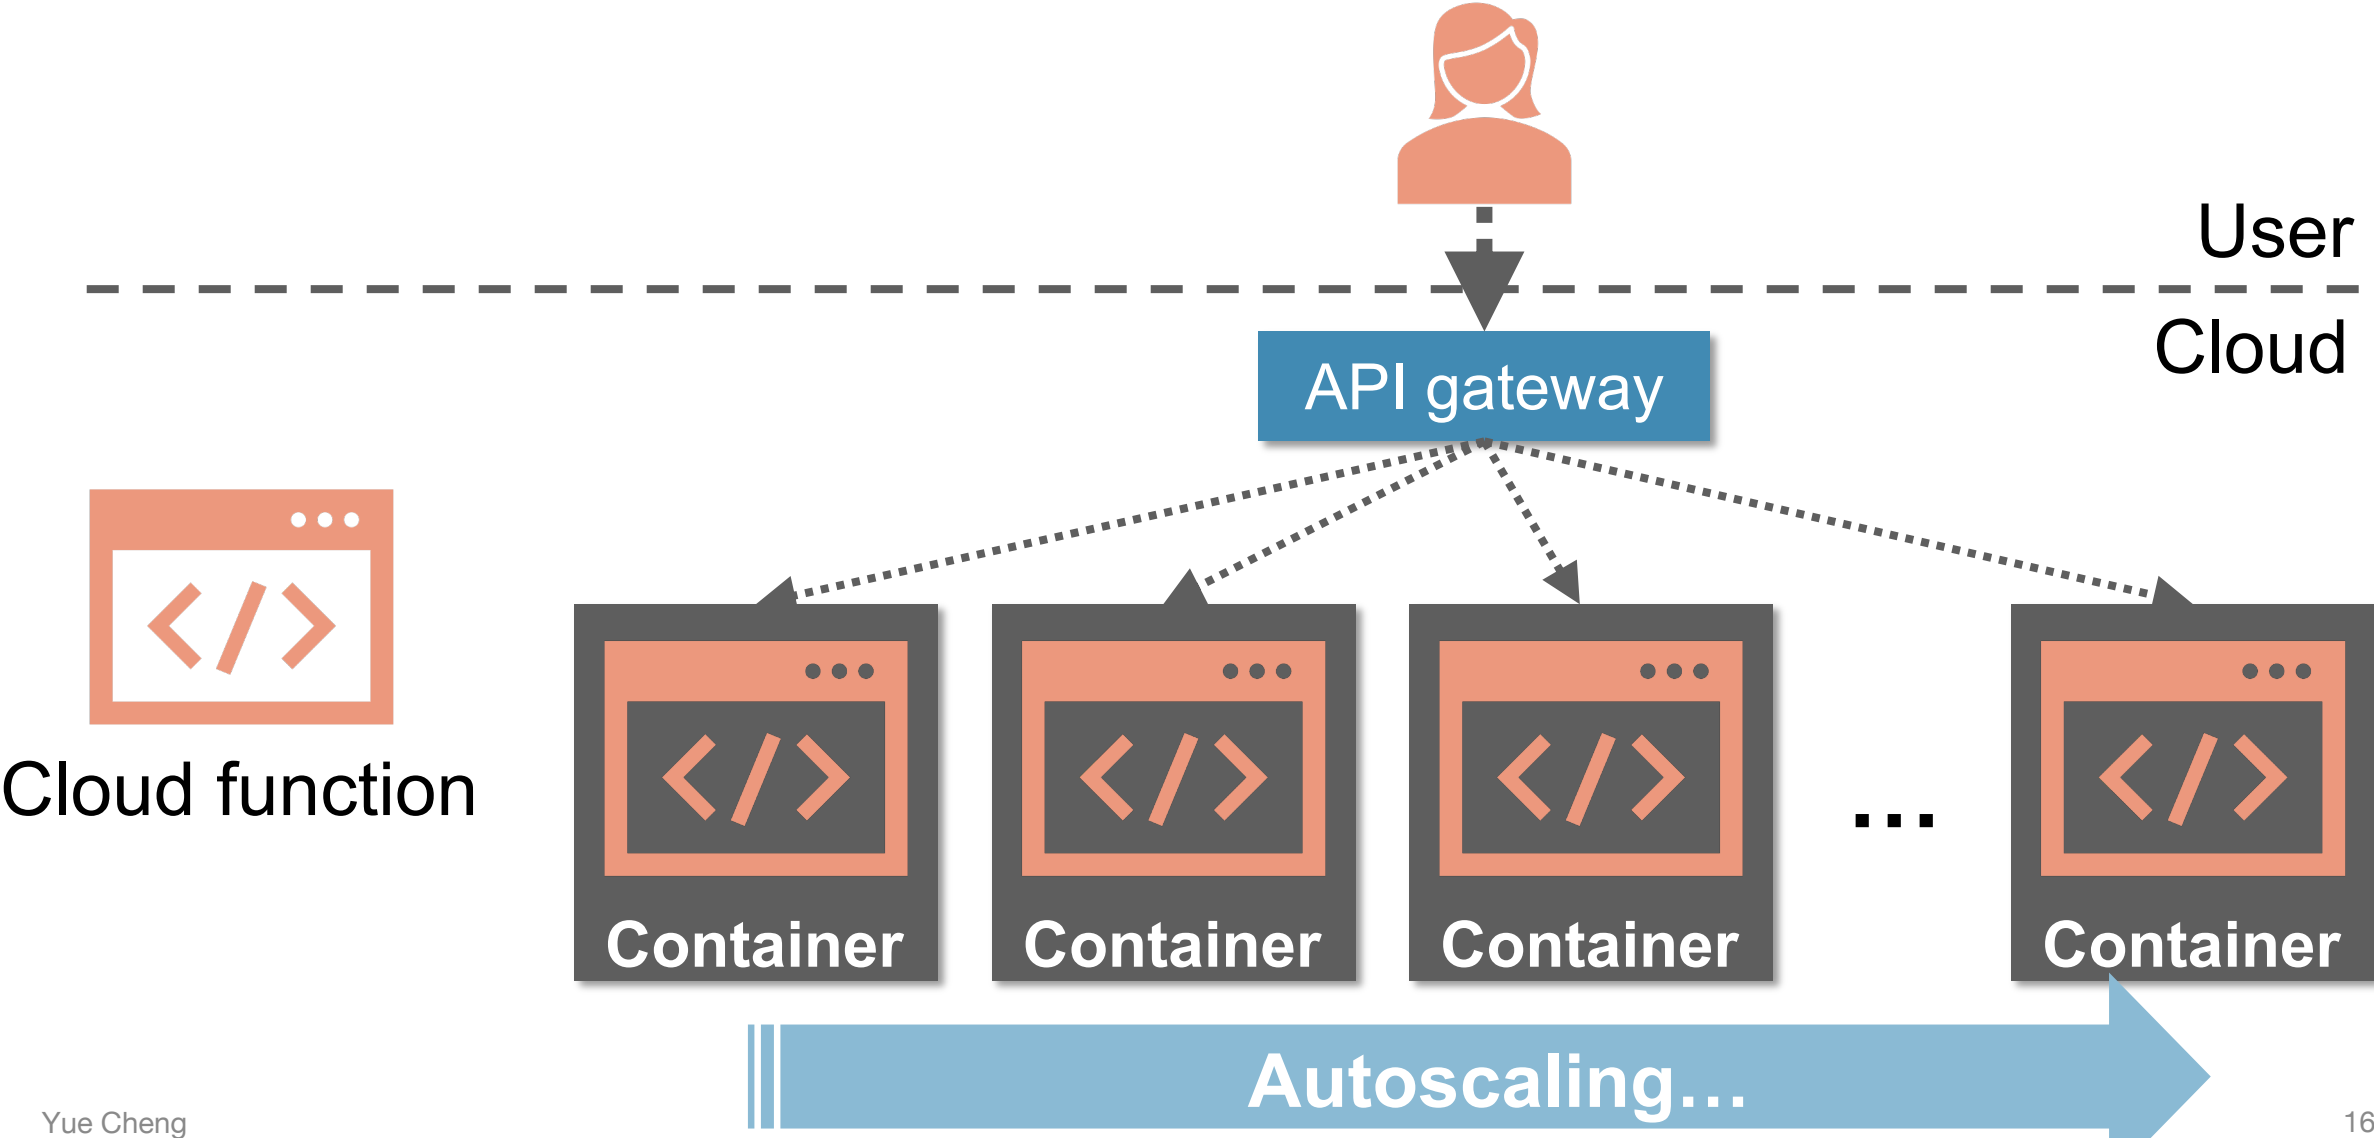
# Function-as-a-Service (FaaS)



# Function-as-a-Service (FaaS)




# Function-as-a-Service (FaaS)






# Python analytics: What we have today



# Python analytics: What we have today



User writes interactive analytics and runs it on a notebook server

- No autoscaling for large computations
- Too slow? OOM? Need to scale out manually!
- Too expensive? Idled resources charge \$\$

# Python analytics: What we have today




User writes interactive analytics and runs it on a notebook server

- No autoscaling for large computations
- Too slow? OOM? Need to scale out manually!
- Too expensive? Idled resources charge \$\$


**High barriers to enter for those who lack CS/systems background**

# Python analytics: What we would like to have



- User writes interactive analytics and runs it **on FaaS**
- Elastically and automatically scales to the right size
  - Pay-per-use with minimal \$\$ cost
  - Expertise of writing parallel programs **NOT required**
  - Manual cluster maintenance **NOT required**

# PyWren: Stateful computing over stateless serverless functions



# PyWren: How it works

```
def fn(input):  
    return input + 1
```


➔ `futures = runner.map(fn, dataset)`

```
print([f.result() for f in futures])
```

Scheduler

Your laptop

Cloud



# PyWren: How it works

```
def fn(input):  
    return input + 1
```

➔ `futures = runner.map(fn, dataset)`

```
print([f.result() for f in futures])
```

Scheduler

Your laptop

Cloud



# PyWren: How it works

```
def fn(input):  
    return input + 1
```

→ `futures = runner.map(fn, dataset)`

```
print([f.result() for f in futures])
```

Your laptop

Scheduler

Invoke

Fn

Invoke

Fn

Invoke

Fn

Invoke

Fn

...

Cloud

S3



# PyWren: How it works


```
def fn(input):  
    return input + 1
```

→ `futures = runner.map(fn, dataset)`

```
print([f.result() for f in futures])
```

Your laptop


Scheduler



Cloud

# PyWren: How it works

*Compute*



Cloud

```
def fn(input):  
    return input + 1  
  
futures = runner.map(fn, dataset)  
  
print([f.result() for f in futures])
```

Scheduler

Your laptop


# PyWren: How it works

```
def fn(input):  
    return input + 1  
  
futures = runner.map(fn, dataset)
```

➔ `print([f.result() for f in futures])`

Scheduler

Your laptop



Cloud

# PyWren: How it works

Lambda functions are terminated

```
def fn(input):  
    return input + 1  
  
futures = runner.map(fn, dataset)
```

→ `print([f.result() for f in futures])`

Scheduler

Fn

Fn

Fn

...

Fn

S3

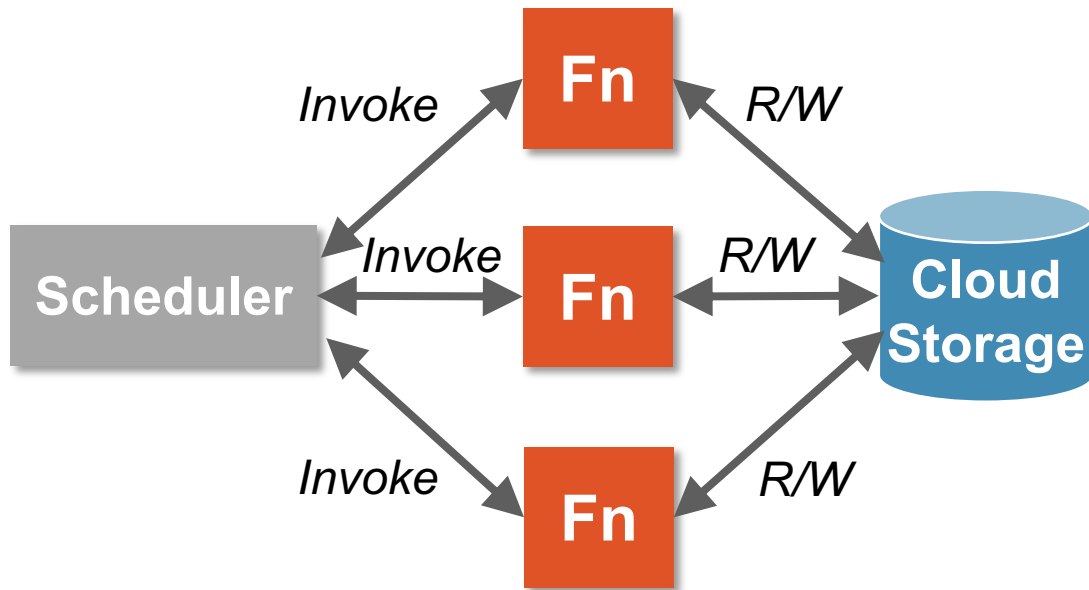
Your laptop

Cloud

# Quantifying the pain of FaaS

How FaaS adds huge amounts of **performance taxes**

# Python analytics on FaaS is slow!


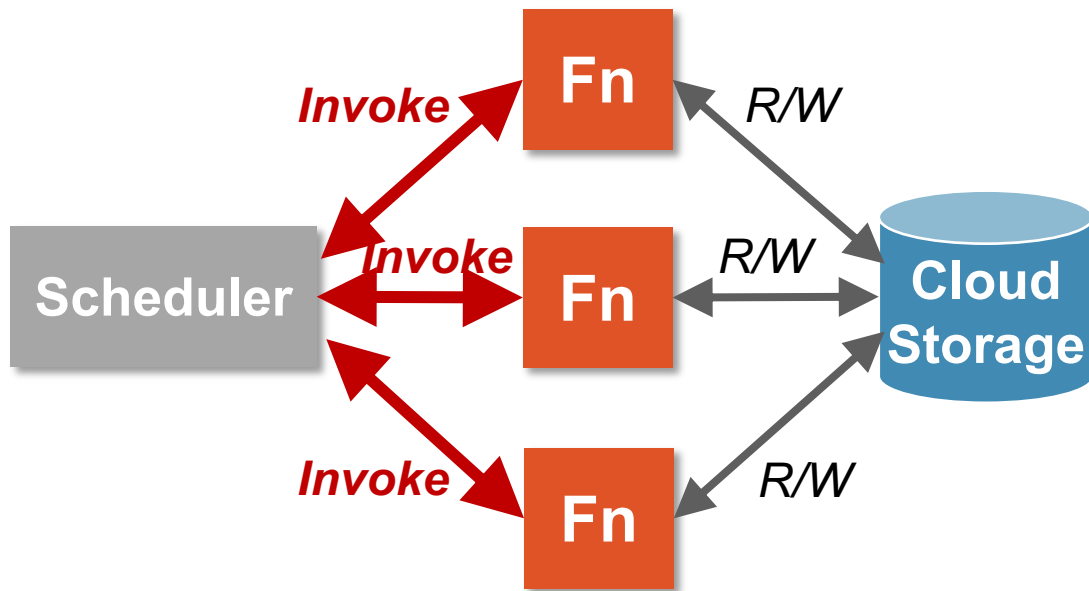


PyWren and numpywren

\* [PyWren] Occupy the Cloud: Distributed Computing for the 99%. In ACM SoCC'17.

\* [numpywren] Serverless linear algebra. In ACM SoCC'20.

# Python analytics on FaaS is slow!




State-of-the-art FaaS frameworks pay huge amounts of FaaS taxes

- **Task scheduling bottleneck:** Too slow to scale to thousands of functions

\* [PyWren] Occupy the Cloud: Distributed Computing for the 99%. In ACM SoCC'17.


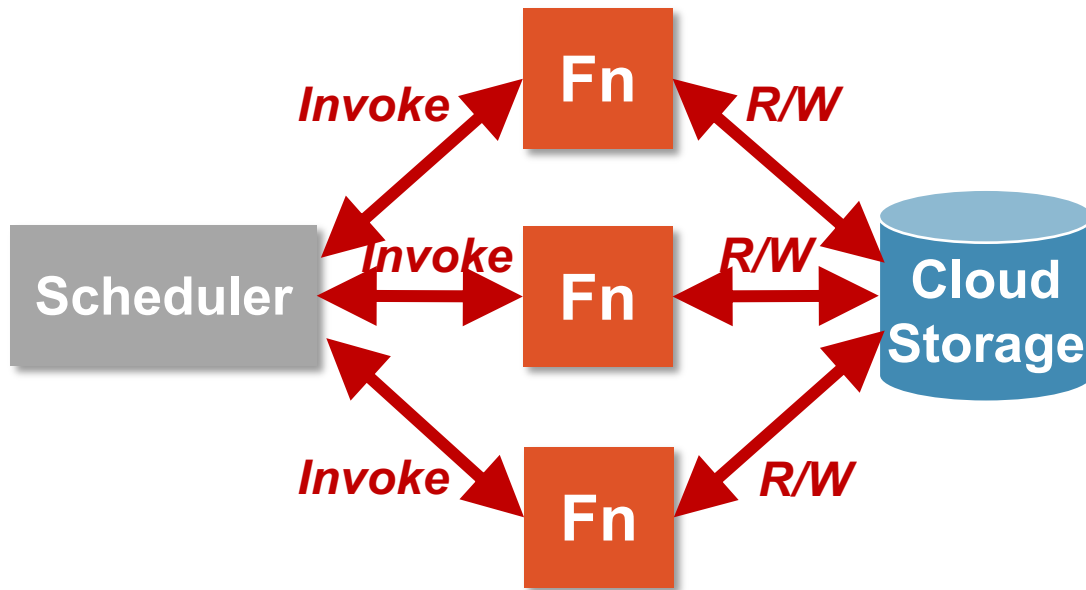
\* [numpywren] Serverless linear algebra. In ACM SoCC'20.

# High HTTP invocation cost for AWS Lambda





# Python analytics on FaaS is slow!



State-of-the-art FaaS frameworks pay huge amounts of FaaS taxes

- **Task scheduling bottleneck:** Too slow to scale to thousands of functions
- **I/O bottleneck:** Excessive data movement cost due to FaaS constraint

\* [PyWren] Occupy the Cloud: Distributed Computing for the 99%. In ACM SoCC'17.

\* [numpywren] Serverless linear algebra. In ACM SoCC'20.

**Naively porting a stateful cluster computing application to FaaS won't work!**

**Need a FaaS-centric approach**

**Insight:** A FaaS framework may not care about traditional metrics (load balancing, cluster util.)

# Wukong

**WUKONG** 悟空  
SERVERLESS DAG ENGINE




Wukong is a **FaaS-centric** parallel computing framework

<https://github.com/ds2-lab/Wukong>

**Key idea:** Partitions the work of a centralized scheduler across many functions to take advantage of FaaS elasticity

- Functions schedule tasks by **invoking** functions
- Functions execute multiple tasks to **reduce data movement cost**
- Functions scale out / in **autonomously**

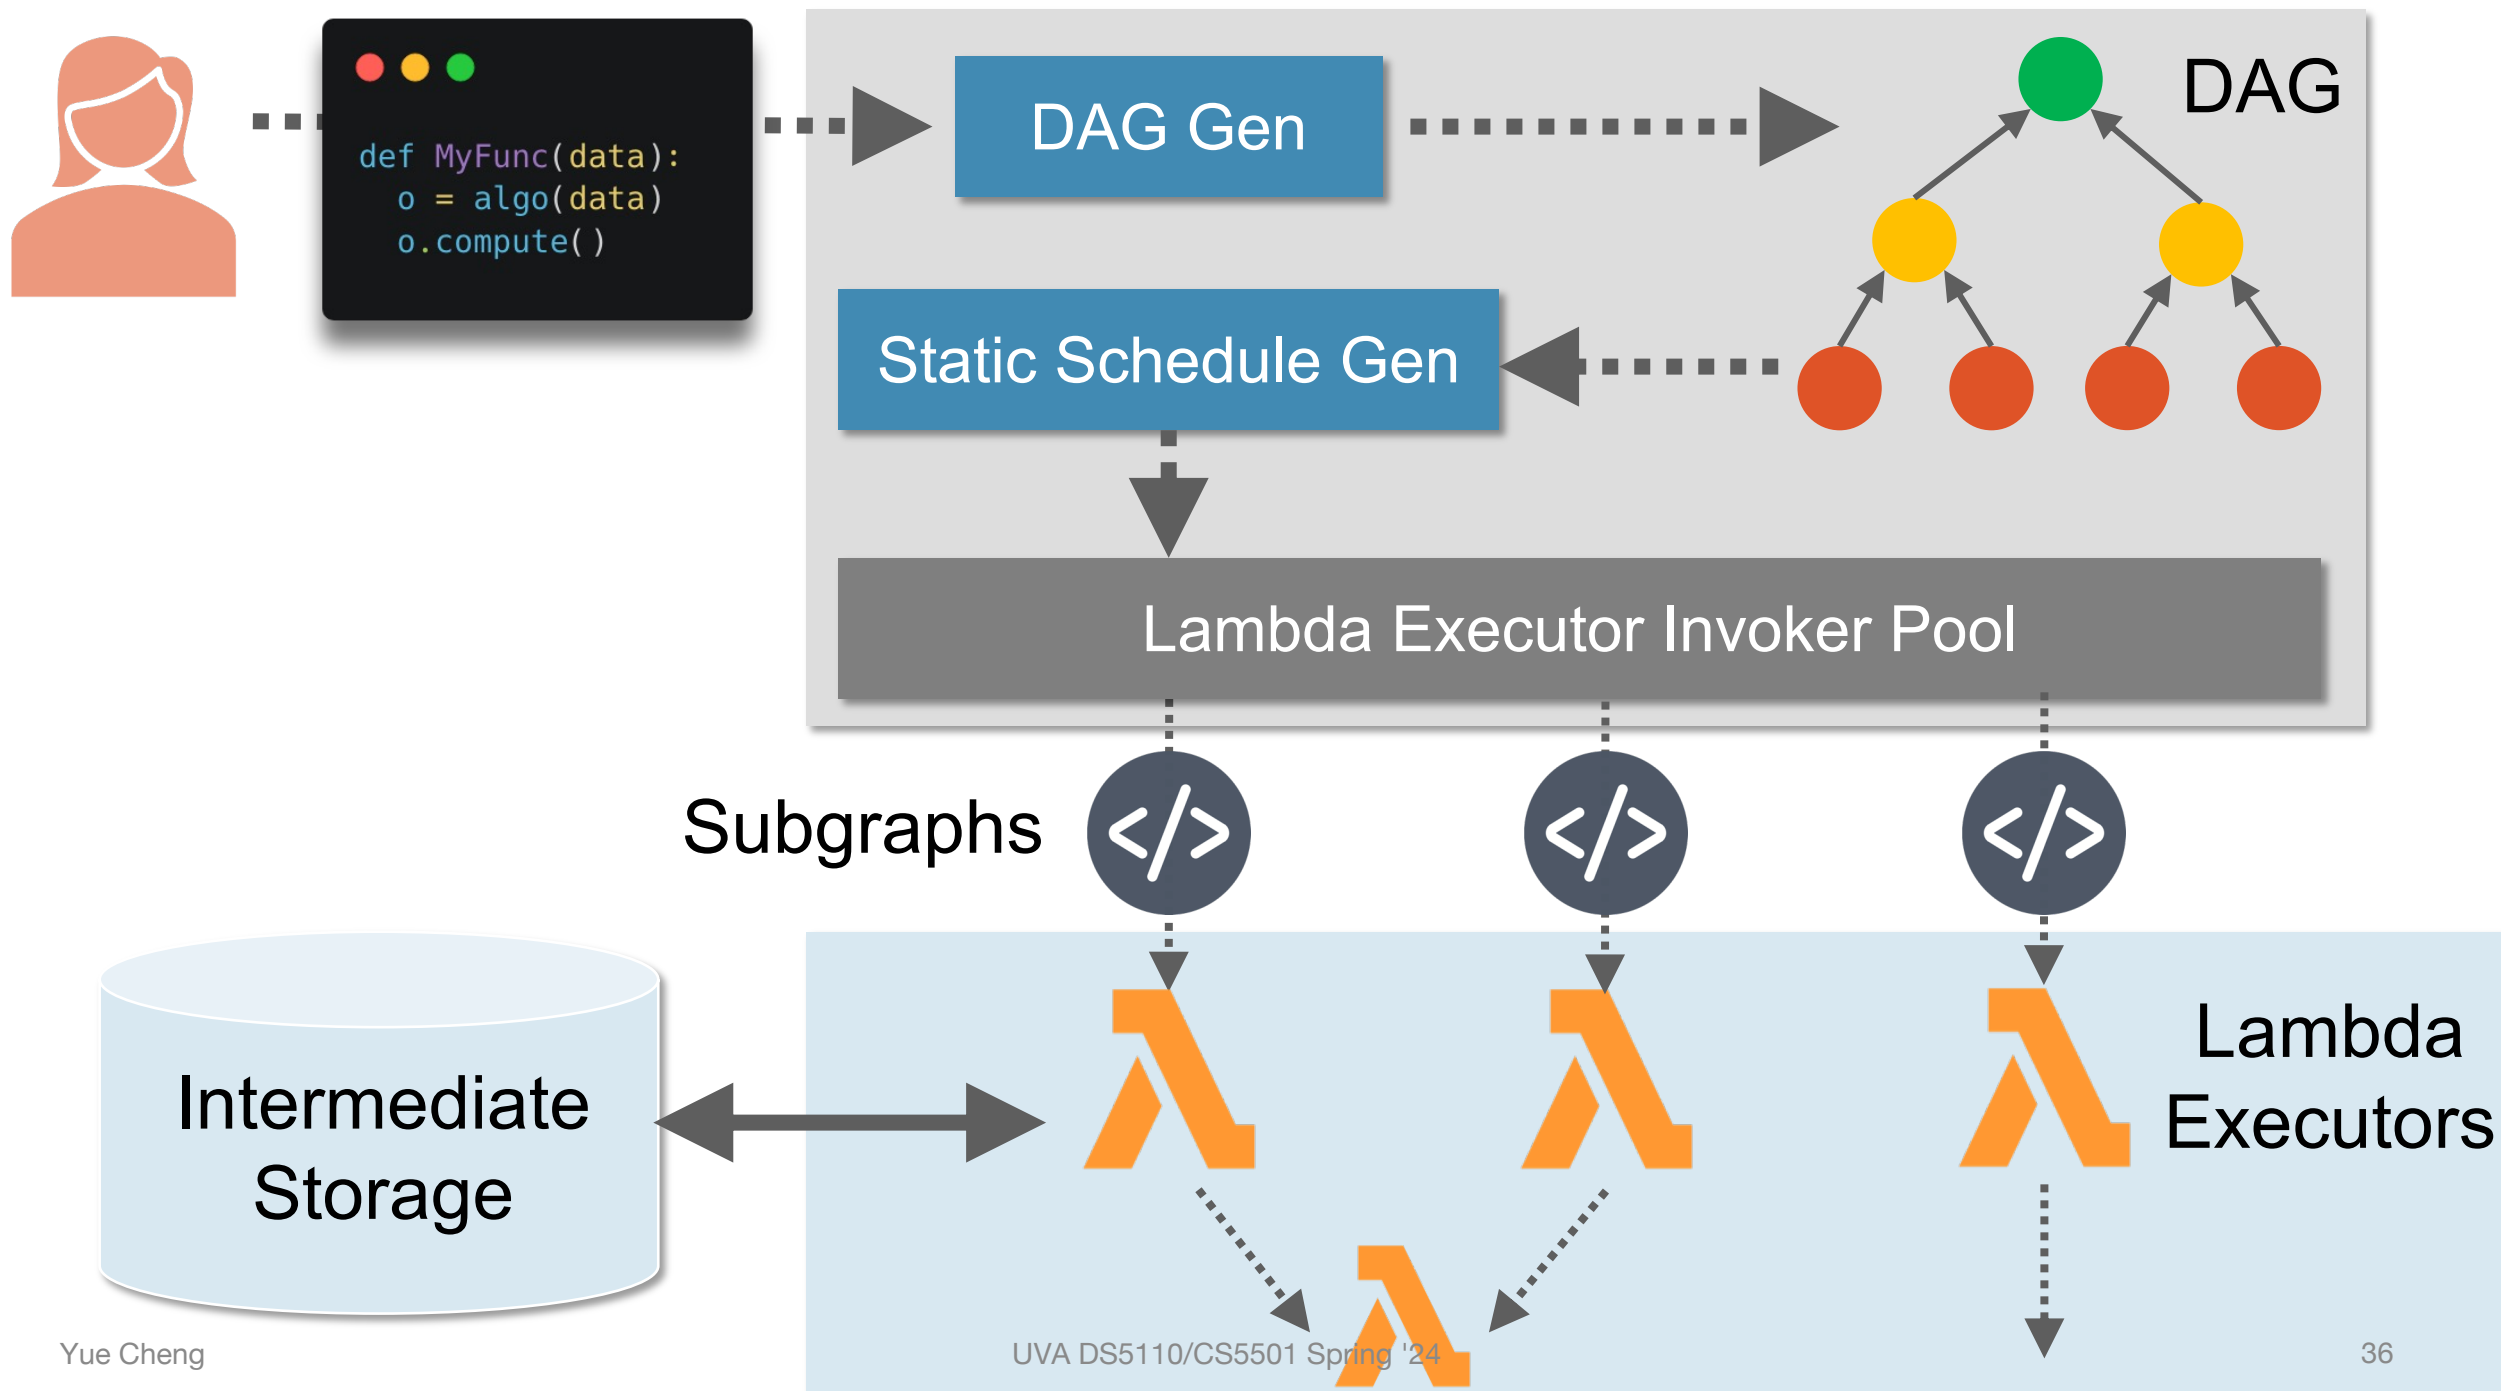
Naturally enables multiple benefits



Exploits autoscaling for scalability

Improved data locality

No tedious cluster configuration






```
def MyFunc(data):  
    o = algo(data)  
    o.compute()
```



DAG Gen

Static Schedule Gen

Lambda Executor Invoker Pool




```
def MyFunc(data):  
    o = algo(data)  
    o.compute()
```



DAG Gen

Static Schedule Gen

Lambda Executor Invoker Pool




```
def MyFunc(data):  
    o = algo(data)  
    o.compute()
```

DAG Gen

Static Schedule Gen

Lambda Executor Invoker Pool






```
def MyFunc(data):  
    o = algo(data)  
    o.compute()
```

DAG Gen

Static Schedule Gen

DAG

Lambda Executor Invoker Pool








```
def MyFunc(data):  
    o = algo(data)  
    o.compute()
```

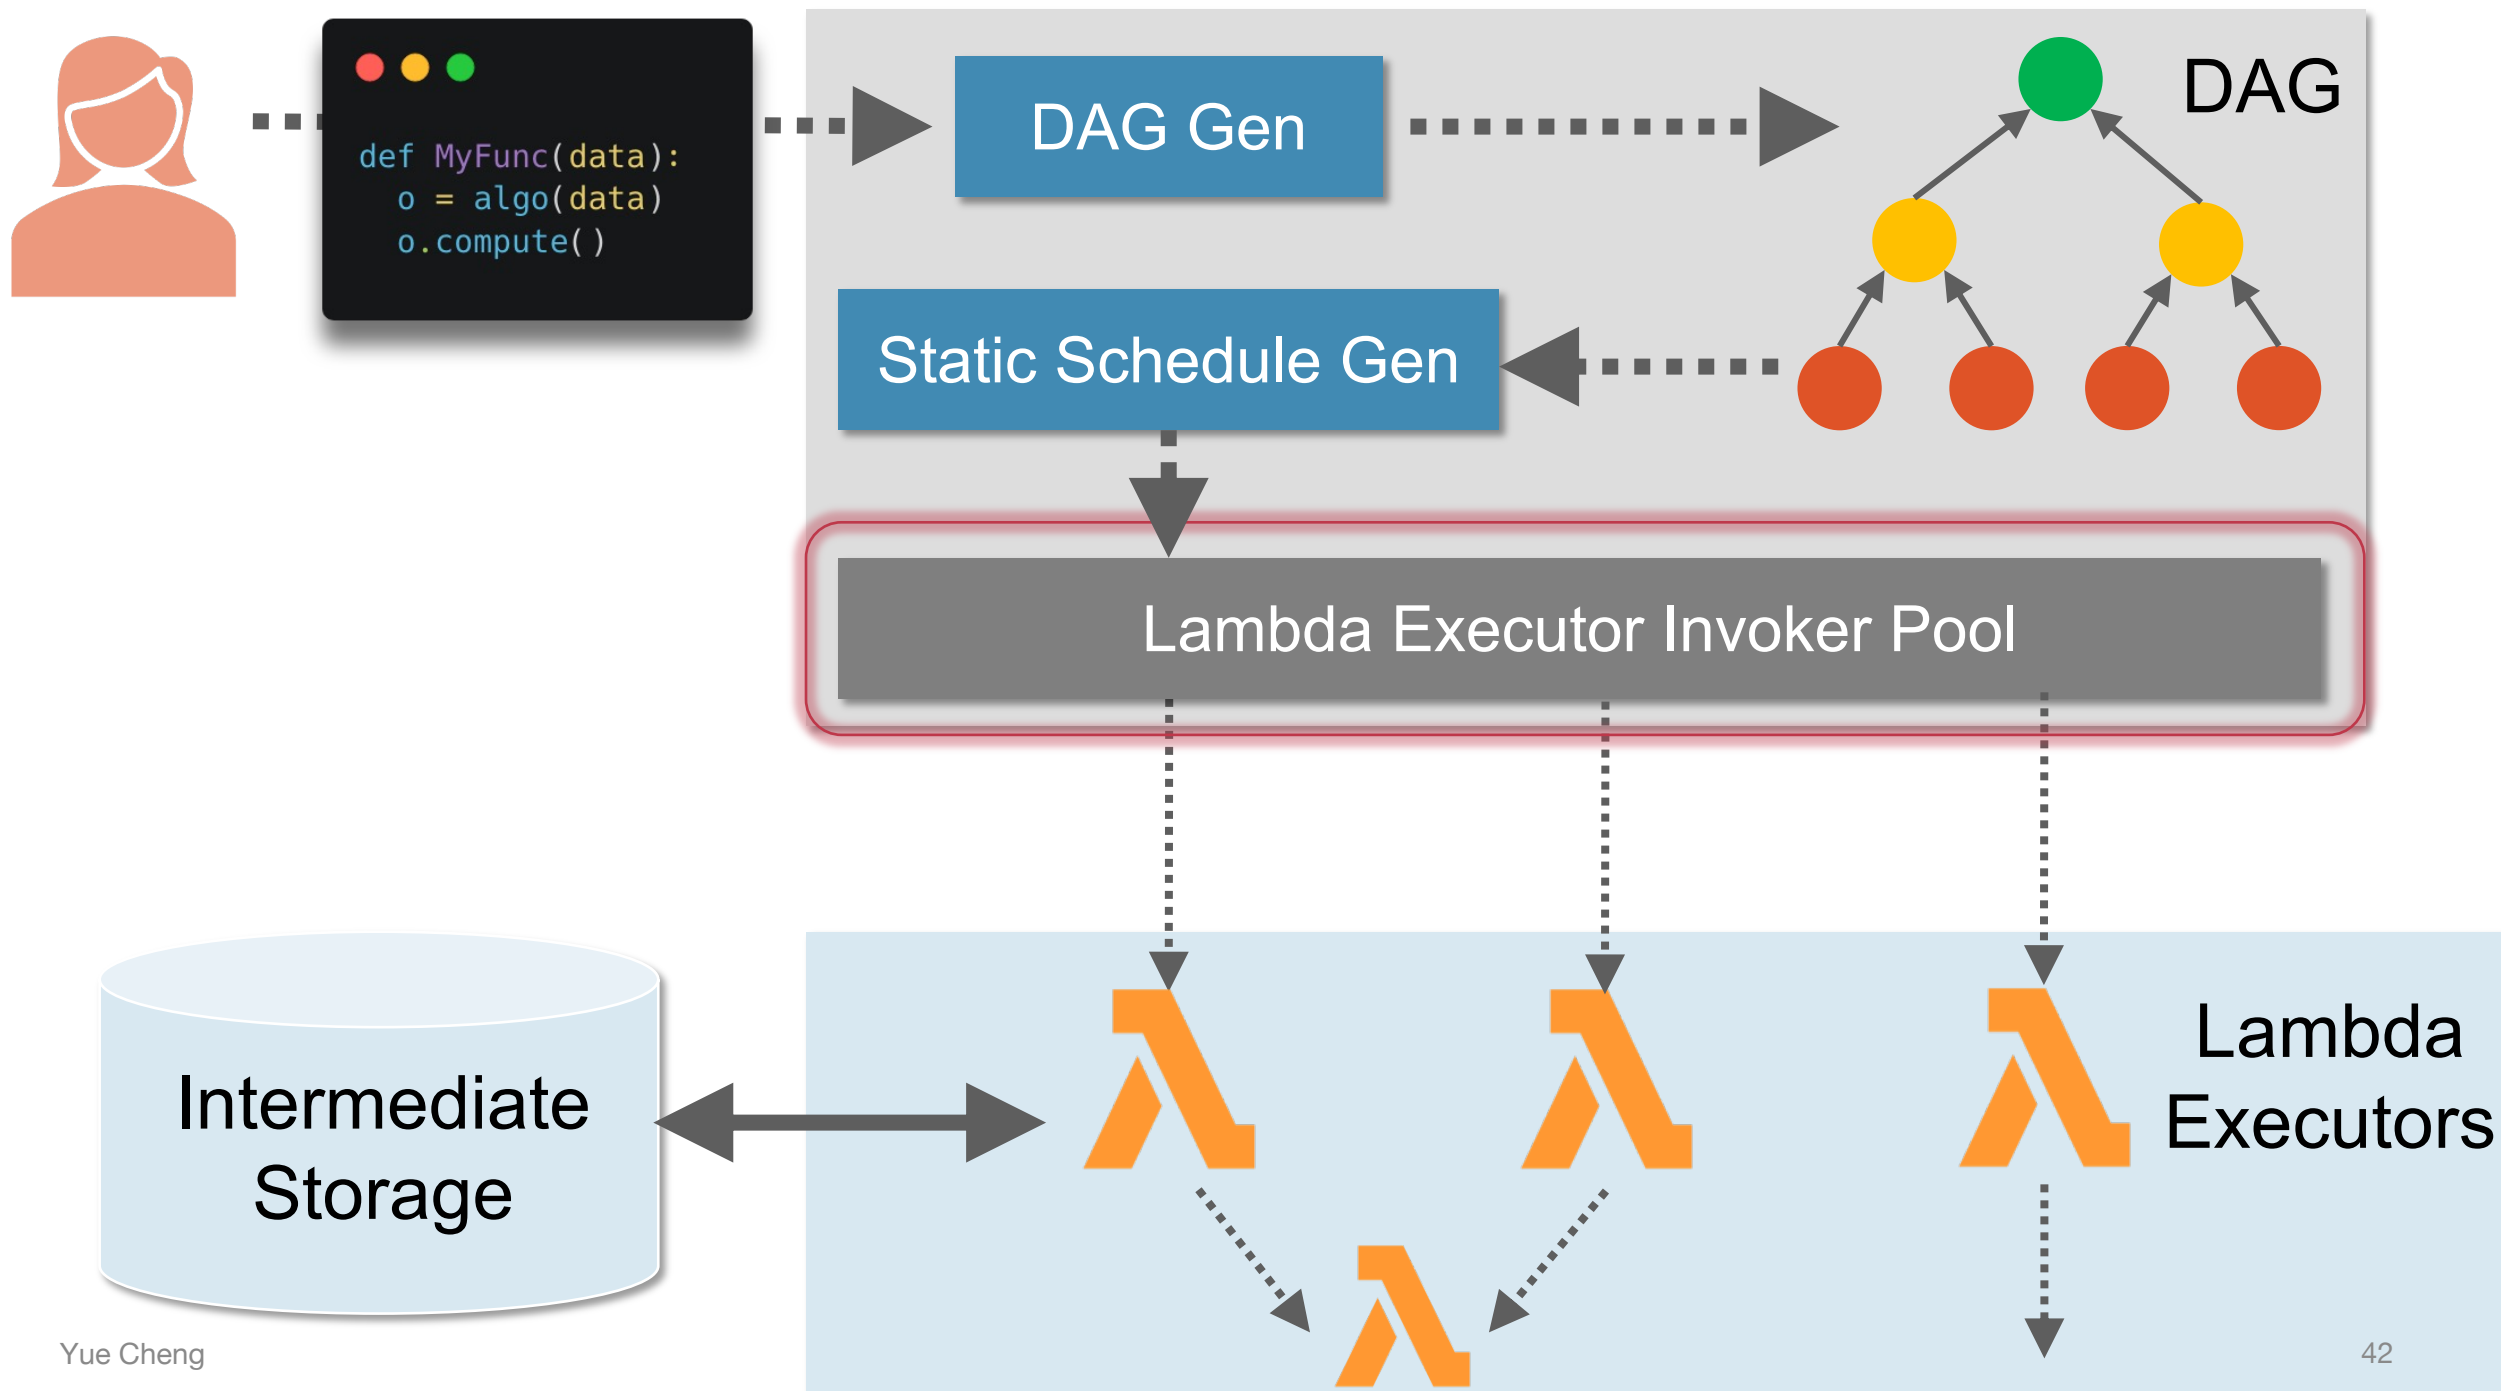
DAG Gen

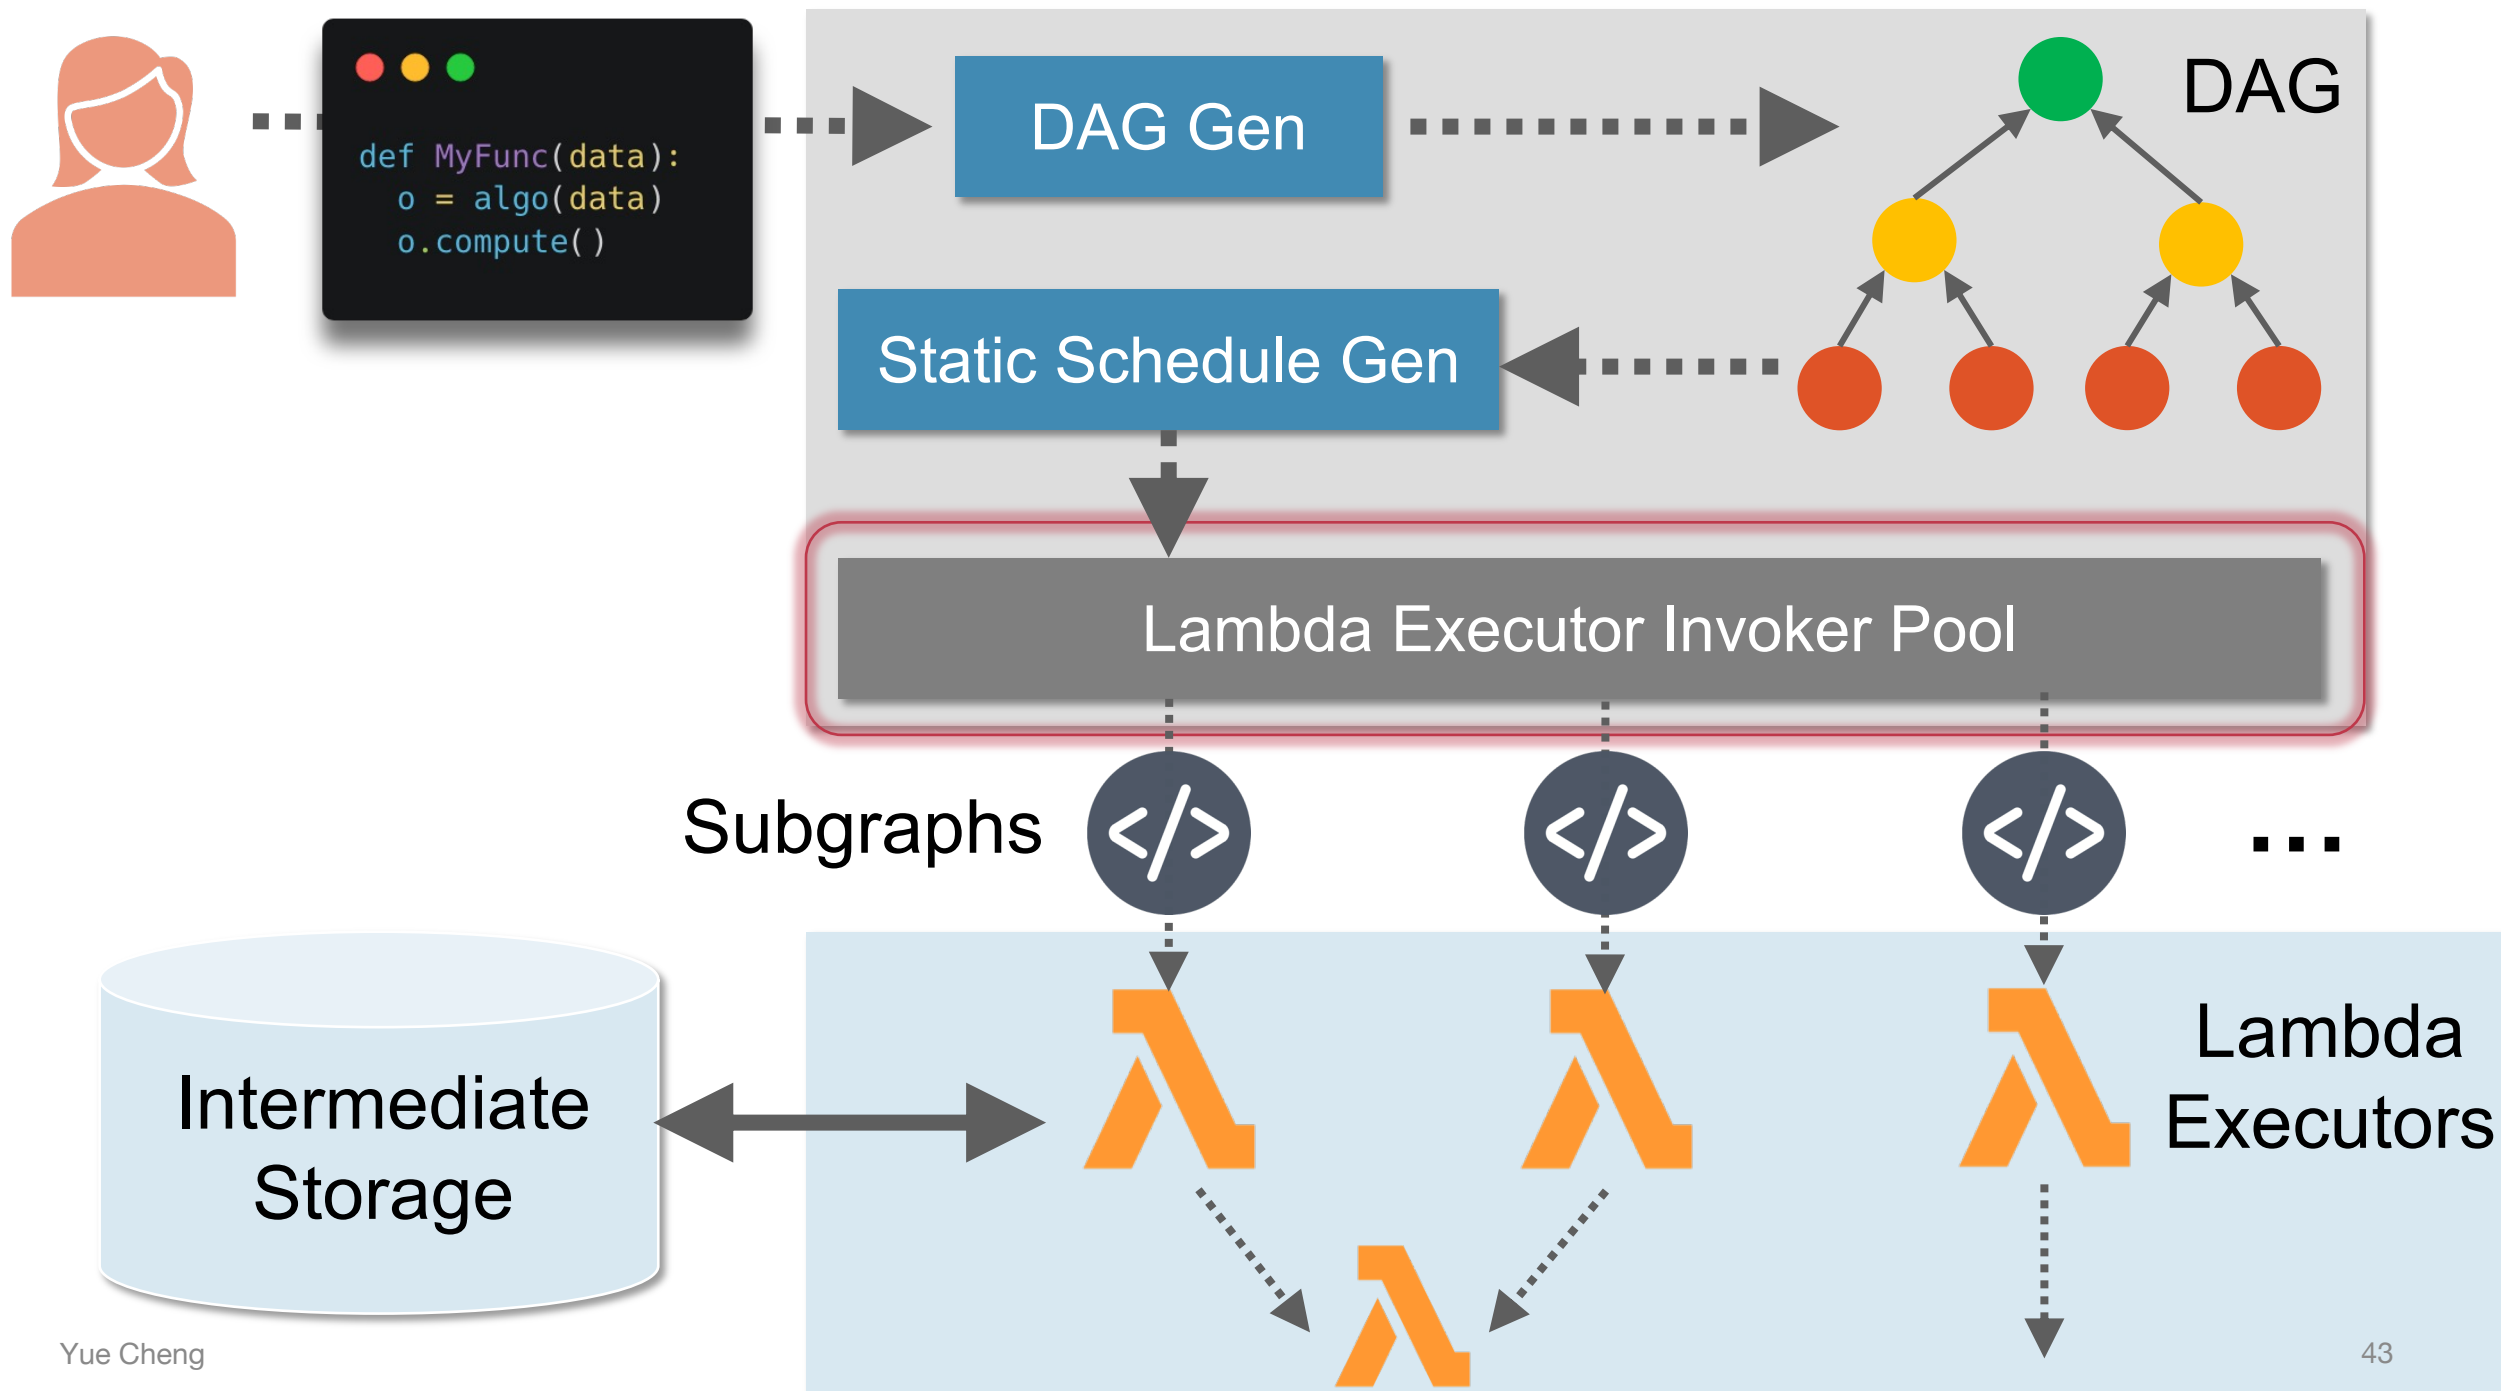
Static Schedule Gen


Lambda Executor Invoker Pool

DAG










```
def MyFunc(data):  
    o = algo(data)  
    o.compute()
```

DAG Gen

Static Schedule Gen

Lambda Executor Invoker Pool

DAG




Intermediate  
Storage




Lambda  
Executors

# Scheduling in Wukong


- Combination of **static** and **dynamic** scheduling
- Input DAG partitioned into **static schedules**, or subgraphs of the original DAG
- Serverless executors are assigned a **static schedule**
- Executors use **dynamic scheduling** to enforce data dependencies and **cooperatively** schedule tasks found in multiple static schedules






```
func MyFunc(data):  
    o = algo(data)  
    o.compute()
```

# Static scheduling




```
func MyFunc(data):  
    o = algo(data)  
    o.compute()
```



# Static scheduling

## Input DAG

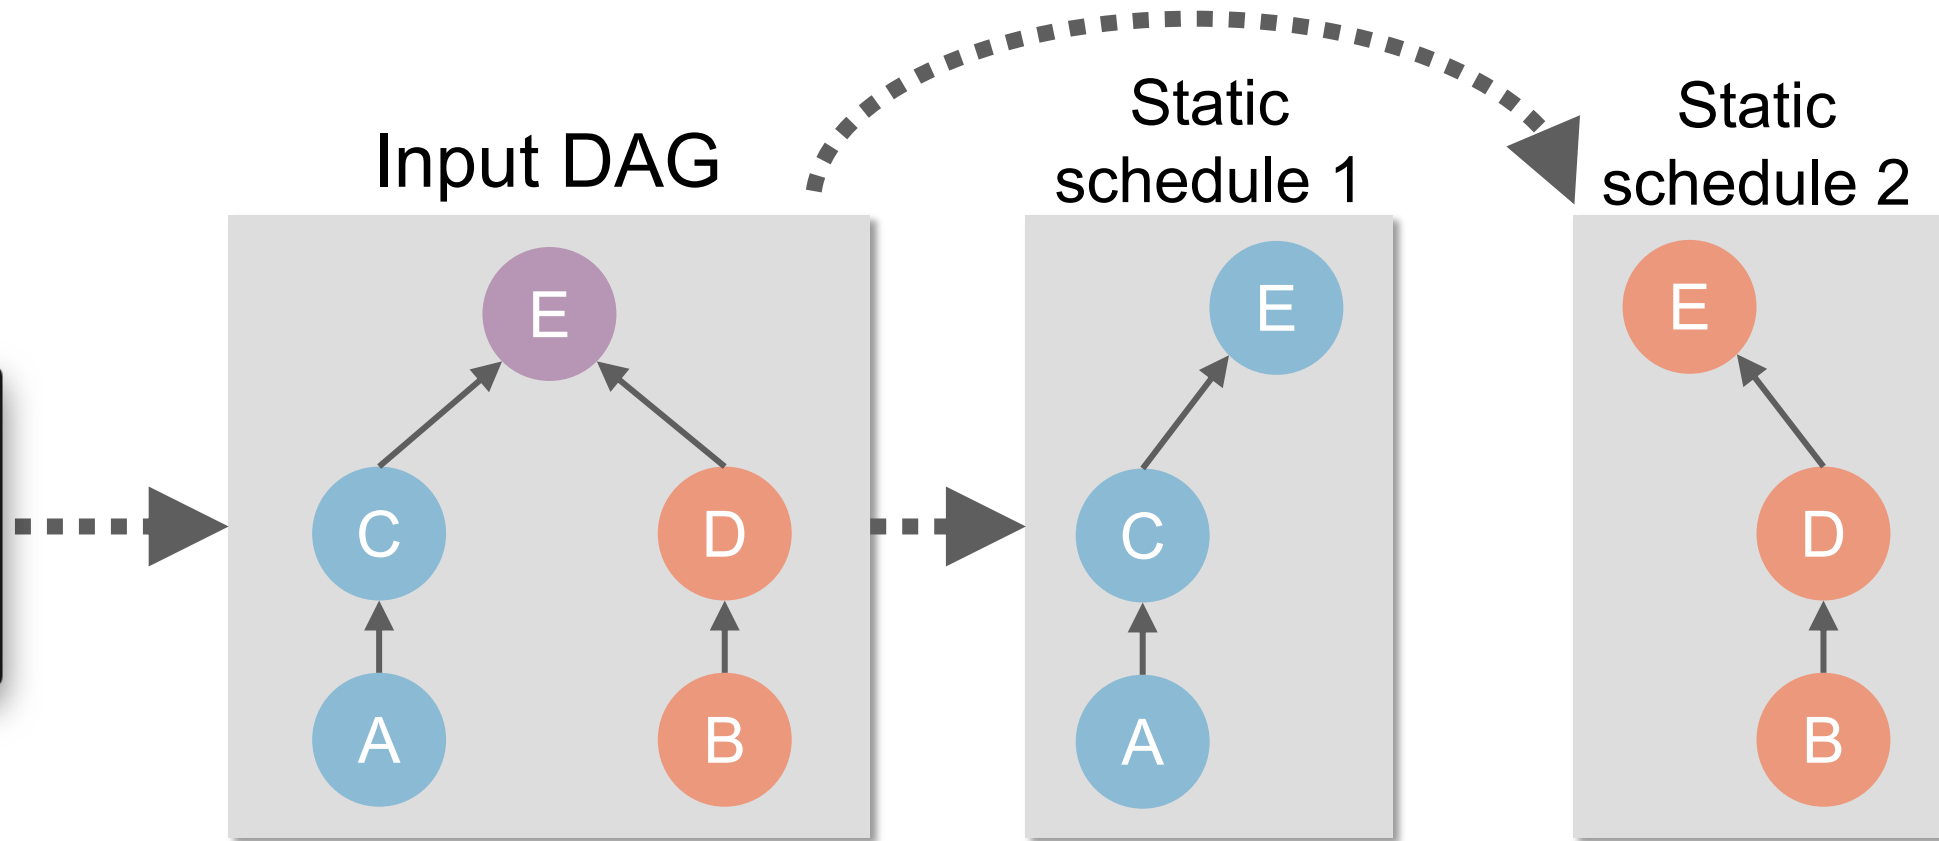
```
func MyFunc(data):  
    o = algo(data)  
    o.compute()
```



# Static scheduling




```
func MyFunc(data):  
  o = algo(data)  
  o.compute()
```




# Static scheduling

```
func MyFunc(data):  
  o = algo(data)  
  o.compute()
```


Input DAG



Static  
schedule 1




Static  
schedule 2




# Static scheduling




Input DAG



# Dynamic scheduling




Executor 1




Executor 2


Input DAG



# Dynamic scheduling




Executor 1




Executor 2

Input DAG



# Dynamic scheduling




Executor 1




Executor 2

Input DAG




# Dynamic scheduling



Executor 1


Executor 2

Input DAG



# Dynamic scheduling

Cooperate




Executor 1




Executor 2

# Application DAGs

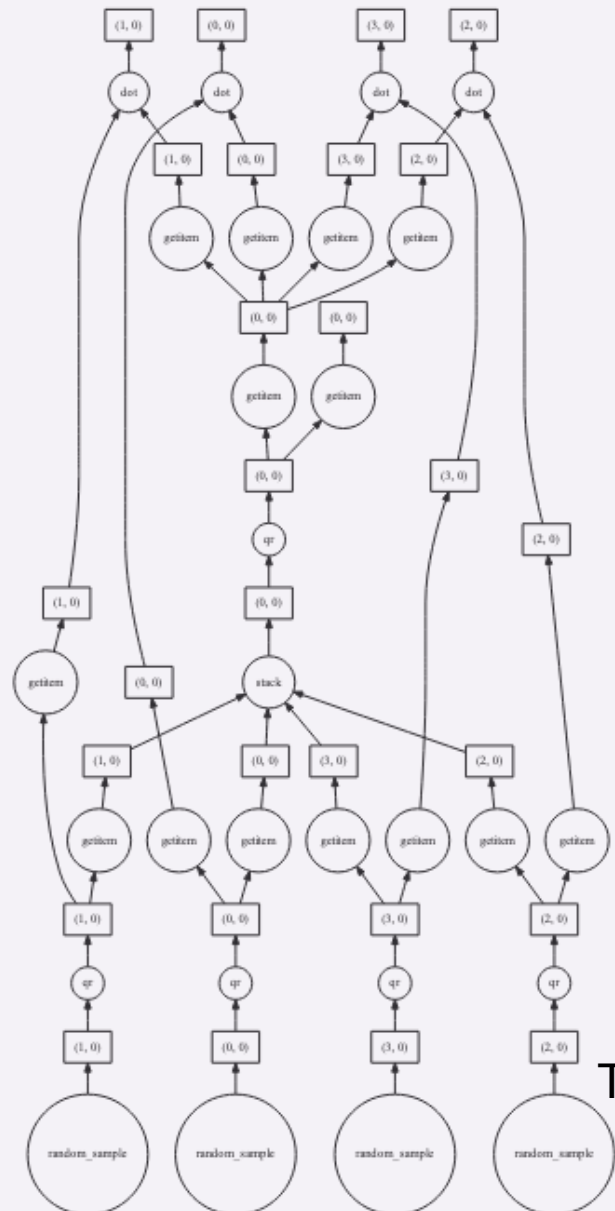
SVD1




SVD2




TSQR



GEMM




Tree reduction






# Application performance: TSQR



Wukong and numpywren ran on AWS Lambda w/ 3GB memory


Dask distributed ran on 125 c5.4xlarge EC2 VMs w/ 2,000 vCPU cores

# Application performance: TSQR




**Wukong outperforms numpywren considerably for all problem sizes.**

# Data movement cost: TSQR




# Data movement cost: TSQR



**Wukong reads and writes considerably less data than numpywren.**

# WUKONG 悟空

## SERVERLESS DAG ENGINE



### Parallelizing Prediction (sklearn.svm.SVC)

```
import pandas as pd
import seaborn as sns
import sklearn.datasets
from sklearn.svm import SVC

import dask_ml.datasets
from dask_ml.wrappers import ParallelPostFit
from distributed import LocalCluster, Client
local_cluster = LocalCluster(host='0.0.0.0:8786',
                             proxy_address = '3.83.198.204',
                             num_fargate_nodes = 10)
client = Client(local_cluster)

X, y = sklearn.datasets.make_classification(n_samples=1000)
clf = ParallelPostFit(SVC(gamma='scale'))
clf.fit(X, y)

X, y = dask_ml.datasets.make_classification(n_samples=800000,
                                             random_state=800000,
                                             chunks=800000 // 20)

# Start the computation.
clf.predict(X).compute()
```

### GEMM (Matrix Multiplication)

```
import dask.array as da
from distributed import LocalCluster, Client
local_cluster = LocalCluster(host='0.0.0.0:8786',
                             proxy_address = '3.83.198.204',
                             num_fargate_nodes = 10)
client = Client(local_cluster)

x = da.random.random((10000, 10000), chunks = (1000, 1000))
y = da.random.random((10000, 10000), chunks = (1000, 1000))
z = da.matmul(x, y)

# Start the computation.
z.compute()
```

<https://github.com/ds2-lab/Wukong>